

Concurrent Systems — Exam  
June 2017

Name: \_\_\_\_\_

Duration: 120 minutes — No document authorized

1.

a) Explain informally what is a *deadlock*. Describe a simple strategy to make sure that multiple threads that share a set of locks never deadlock when acquiring the locks.

---

---

---

---

---

---

---

---

---

---

b) Explain what is the principle of a *future* and describe briefly how one can use it (e.g., using pseudo-code).

---

---

---

---

---

---

---

---

---

---

c) The `atomicAdd(int n)` operation adds value `n` to an atomic integer. Can we implement this operation using `compareAndSet()`? If yes, provide a (pseudo-code) implementation.

---

---

---

---

---

---

---

---

---

---

d) Explain informally what is the issue with the simple “test-and-set” algorithm for implementing a spinlock, and why “test-and-test-and-set” typically performs better.

---

---

---

---

---

---

---

---

---

---

Assume you have two threads, **T<sub>1</sub>** and **T<sub>2</sub>**, executing the following code:

Implement a *rendezvous* point that guarantees that **inst1<sub>1</sub>** happens before **inst2<sub>2</sub>** and **inst1<sub>2</sub>** happens before **inst2<sub>1</sub>**. There is no order constraint between **inst1<sub>1</sub>** and **inst1<sub>2</sub>**. Use only semaphores in your implementation.

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins or other markings on the paper.

[illegible]

Consider the following barrier implementation pseudocode:


```
Lock lock = new ReentrantLock();
Semaphore sem = new Semaphore(0); // Semaphore initially 0
int count = 0, n = NUM_THREADS;

// Reaching the barrier
lock.lock();
count = count + 1;
lock.unlock();
if (count == n) {
    sem.release(); // Increment semaphore
}
sem.acquire(); // Decrement semaphore
// Critical section follows
```

The variable `count` is used to keep track of how many threads reached the barrier. It is protected by a lock. `NUM_THREADS-1` threads will block when trying to acquire the semaphore. The last thread releases the semaphore and all threads can continue to the critical section.

Is this a correct barrier implementation? If not, explain what can go wrong and suggest a solution.

[illegible]



4.

a) Consider the following history for two threads **T1** and **T2** and two FIFO queues **p** and **q**:

```
T1 p.enq(x)
T2 q.enq(y)
T1 p:void
T1 q.enq(x)
T2 q:void
T2 p.enq(y)
T1 q:void
T1 p.deq()
T2 p:void
T2 q.deq()
T1 p:y
T2 q:x
```

Is it linearizable?

---

Is it sequentially consistent?

---

b) If we change the last line of the history to:

```
...
T2 q:y
```

Is it linearizable?

---

Is it sequentially consistent?

---

c) Consider again the original history of a). Prepend one operation (method invocation and response) in front of the history so that it becomes both linearizable and sequentially consistent.

---

---

---

---

5.

Consider the following code:

```
class MyStream {  
  
    static long n0 = 1;  
    static long n1 = 1;  
  
    static Stream<Long> get() {  
        return Stream.generate(  
            () -> {  
                long r = n0 + n1;  
                n0 = n1;  
                n1 = r;  
                return r;  
            });  
    }  
}
```

a) Explain what this stream does.

---

---

---

---

---

---

b) What will be the output of the following code?

```
MyStream.get().limit(8).forEach(System.out::println);
```

---

---

c) Can we modify the code as follows? If not, explain why and propose a possible solution.

```
MyStream.get().parallel().limit(8).forEach(System.out::println);
```

---

---

---

---

---

---



A **CountDownLatch** is a synchronization aid that allows one or more threads to wait until a set of operations being performed in other threads completes. A **CountDownLatch** is initialized with a given count. The **await()** methods block until the current count reaches zero due to invocations of the **countDown()** method, after which all waiting threads are released and any subsequent invocations of **await** return immediately. This is a one-shot phenomenon; the count cannot be reset. Note that the **CountDownLatch** does not require that threads calling **countDown()** to wait for the count to reach zero before proceeding, it simply prevents any thread from proceeding past an **await()** until all threads could pass.

[illegible]