

## Solution for exercise 9

### 9.1 Flooding Consensus (3pt)

Consider *Flooding Uniform Consensus* [CGR11, Algorithm 5.3], in which every process always uses  $N$  rounds before deciding.

- a) Can one reduce the number of rounds for some process? Consider a system with two processes  $p$  and  $q$ .
- b) Is this algorithm correct if the failure detector is not perfect ( $\mathcal{P}$ ) but is only the eventually perfect one ( $\diamond\mathcal{P}$ )? Justify your answer.

#### Solution.

- a) The answer is no. In the case of two processes, the algorithm needs two communication steps because a decision cannot be reached by all correct processes after one step. Consider a system with two processes  $p$  and  $q$  and assume that every process must decide after executing only one round in the algorithm. We describe an execution where the processes do not reach uniform agreement; thus, the algorithm needs at least two rounds. Suppose that  $p$  and  $q$  propose two different values  $v$  and  $w$ , respectively. Without loss of generality, assume that  $v < w$ . In the following execution, process  $p$  is faulty. During the first round,  $p$  and  $q$  send their values to each other. Process  $p$  receives its own value and  $q$ 's value and decides at the end of the round by our assumption. It decides  $v$ , the smaller value, and then crashes. Now, assume that the message from  $p$  to  $q$  in round one is delayed arbitrarily. There is a time after which  $q$  detects  $p$  to have crashed because of the strong completeness property of the perfect failure detector. As  $q$  cannot know that  $p$  actually did send a message,  $q$  reaches the end of the first round and must decide. Process  $q$  decides its own value  $w$ , which violates *uniform agreement*. Note that in the original algorithm, where the processes decide only after two rounds, the above scenario cannot occur. This is because  $p$  crashes before it can decide (in fact, it never decides); later on,  $q$  decides  $w$ .
- b) A violation of *strong completeness* property of the perfect failure detector could lead to the violation of the *termination* property of consensus as follows. There is an execution in which a process  $p$  waits to deliver a message from a process  $q$  or to detect the crash of process  $q$ . Should  $q$  crash and  $p$  never detect the crash of  $q$ ,  $p$  would remain blocked forever and never decide. Consider now *strong accuracy*. If it does not hold, it could violate the *agreement* property. It can happen that if process  $q$  crashes after deciding  $x$ , and  $p$  is falsely suspected to have crashed by processes  $r$  and  $q$ , then  $r$  and  $q$  will decide  $y$ .

## 9.2 Leader-Driven Consensus (4pt)

Study the *Leader-Driven Consensus* [CGR11, Chap. 5.3], where epoch-change and epoch consensus are implemented as described.

- Why does this algorithm require a majority of correct processes?
- Which property is violated if there is no majority of correct processes?
- Draw or describe an execution with four processes  $p, q, r, s$  that justifies your previous answers.

**Solution.** We explain this for the case of a system of four processes  $p, q, r$ , and  $s$ . Assume by contradiction that the algorithm tolerates the crash of two processes. Assume that  $p$  and  $q$  propose a value  $v$ , whereas  $r$  and  $s$  propose a different value  $v'$ . Consider an execution (Figure 1) where  $p$  and  $q$  crash initially: in this execution,  $r$  and  $s$  decide  $v'$  to respect the *validity* property of consensus. However, processes  $r$  and  $s$  cannot distinguish this execution from one where  $p$  and  $q$  continue to run and eventually decide, but never receive any message from  $r$  and  $s$  (this execution must terminate to satisfy consensus). Then,  $p$  and  $q$  decide  $v$  by the assumption on the protocol and violate the *agreement* property. With a majority of correct processes, the issue does not occur.




Figure 1. Example of an execution of *Leader-Driven Consensus* algorithm.

## 9.3 Leader-Driven Consensus, optimized (3pt)

In practice, *Leader-Driven Consensus* [CGR11, Chap. 5.3] protocol is often implemented with slightly fewer messages. This optimization assumes that all processes start with the same process  $\ell_0$  as the first leader and that  $\ell_0$  is correct. Describe why the first epoch consensus instance may omit the first round of message exchange (of the READ and STATE messages) between  $\ell_0$  and the other processes.

**Solution.** It may omit the first round of message exchanges for reading because in the initial epoch, process  $\ell_0$  knows that no decision could have been made in a previous epoch (as there is no previous epoch). This first round in every epoch consensus instance is actually only needed to make sure that the leader will propose a value that might have been decided (more

precisely, to ensure the lock-in property of epoch consensus). The algorithm, therefore, saves one communication phase by directly having  $\ell_0$  write  $v_\ell$  and all correct processes decide after three communication steps.

## References

- [CGR11] C. Cachin, R. Guerraoui, and L. Rodrigues, *Introduction to reliable and secure distributed programming (Second Edition)*, Springer, 2011.