



1 Introduction - February 19, 2020

1.1 Defining Dependable Systems

QUOTES:

A distributed system is a system where a computer of which you did not know it exists can prevent you from getting your job done. - Leslie LAMPORT

There is perhaps a market for maybe five computers in the world. - TJ WATSON

FAULT → ERROR → FAILURE

- Train delayed because of tree has fallen on the tracks
- Travelers reach destination too late
- Alice misses her exam

	<u>FAULT</u>	<u>ERROR</u>	<u>FAILURE</u>
Train:	Tree fallen	no train	delay for passengers
Journey:	Train delay	delay	reached destination 2h after intention
Exam:	arrival 2h late	missed time-slot	repeat exam

FAULT: cause of failure

ERROR: internal state of system, not according to specification

FAILURE: observable deviation of specification

FAULT examples:

- timing
- cables
- power supply
- messages lost
- data loss (solved with RAIDs)

1.1.1 How to make systems tolerate faults

- PREVENTION
- TOLERANCE
 - Replication/Redundancy
 - Recovery
- REMOVAL
- FORECASTING/PREDICTION

SAFETY ≠ SECURITY

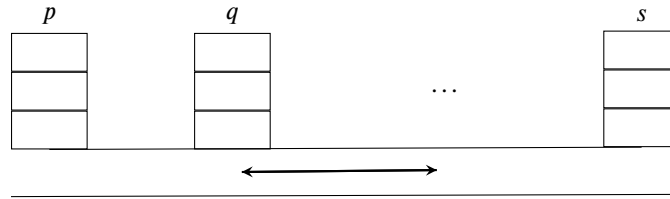
SAFETY is connected to loss of live/material due to accidents

SECURITY is connected to malicious intent

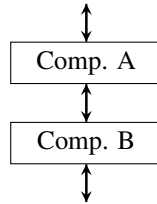
1.1.2 Defining distributed computation

Processes $\Pi = \{p, q, r, s \dots\}$

$|\Pi| = N$



COMPONENTS



EVENTS for Component c :

$\langle c, event \mid param_1, param_2 \dots \rangle$

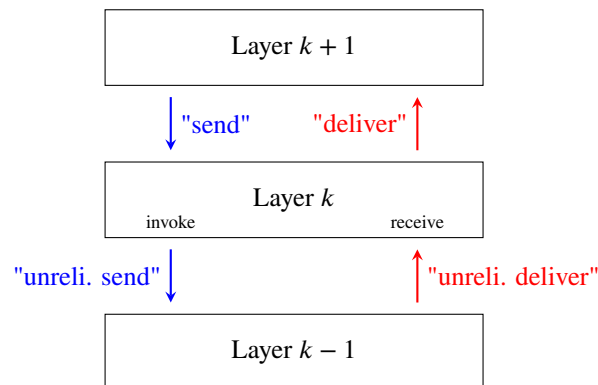
upon $\langle c, ev_1 \mid param_1 \rangle$ do

do something

trigger $\langle b, domore \mid p \rangle$

upon $\langle b, domore \mid p \rangle$ do

1.1.3 Layered modules



Events either travel:

- upwards (red): indication
- downwards (blue): request

Events on a given layer may be:

- input events (IN)
- output events (OUT)



1.1.4 Module Jobhandler

Events:

Request: $\langle jh, handle \mid job \rangle$

Indication: $\langle jh, confirm \mid job \rangle$

Properties:

Every job submitted for handling is eventually confirmed.

Implementation (synchronized) JOBHANDLER

State

...

upon $\langle jh, handle \mid job \rangle$ do

"process job"

trigger $\langle jh, confirm \mid job \rangle$

upon ...

upon ...

Implementation (asynchronized) JOBHANDLER

State

$buf \leftarrow \emptyset$

upon $\langle jh, handle \mid job \rangle$ do

$buf \leftarrow buf \cup \{job\}$

trigger $\langle jh, confirm \mid job \rangle$

upon $buf \neq \emptyset$ do

$job \leftarrow$ some element of buf

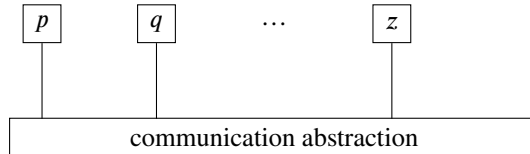
"process job"

$buf \leftarrow buf \setminus \{job\}$

1.2 Concurrency and Replication in Distributed Systems

2 Models and Abstractions - February 26, 2020

2.1 Processes and Protocols



- Set of Processes Π
 $|\Pi| = N$
- A process is an automaton
- A protocol is a set of processes

2.1.1 Execution

- Each computation step and every step of sending a message or receiving a message is an event
- An execution (history) is a sequence of all events of the processes as seen by a (hypothetical) global observer
- trace = execution

2.1.2 Properties

Used for specifying the abstractions:

- **Safety properties** (*something "bad" has not happened*)
If a property P has been violated in some execution E , then there exists a prefix E' of E such that in every extension of E' , property P is violated
- **Liveness properties** (*something "good" will happen in the future [EVENTUALLY]*)
Property P can be satisfied by some extension \tilde{E} of a given execution E

Safety or Liveness alone is not very useful. Only combination of both properties.

2.1.3 Process Failures

A process consists of different modules - if one of them fails the entire thing fails at once.

★ Crashes

- *Omission failures* (message sending and receiving events are omitted)
- *Crash-Recovery Failure*
 - store(-) operation to write to stable storage
 - upon recovery, one can restore(-) data from this stable storage
- *Eavesdropping Fault*

★ Arbitrary Fault (Byzantine Fault)



2.2 Cryptographic Abstraction

- **Hash functions** (SHA-256)
 $H : 0, 1^* \rightarrow \{0, 1\}^k$
 - collision-free: difficult to find x, x' with $x \neq x'$ and $H(x) = H(x')$
- **Message-Authentication-Code (MAC)** (HMAC-SHA256)
 - $\text{authentication}(p, q, m) \rightarrow a$
 - $\text{verifyAuth}(p, q, m, a) \rightarrow \text{YES/NO}$
- **Digital Signatures** (RSA, (EC)DSA)
 - $\text{sign}(p, m) \rightarrow s$
 - $\text{verifySign}(p, m, s) \rightarrow \text{YES/NO}$
 - ★ Correctness:
 $\forall m, p : \text{verifySign}(p, m, \text{sign}(p, m)) = \text{YES}$
 - ★ Security:
 $\forall m, p, s : \text{verifySign}(p, m, s) = \text{NO}$, unless p has executed $\text{sign}(p, m) \rightarrow s$

2.3 Communication Abstraction

Every process can send messages to every other process.

2.3.1 Stubborn point-to-point links

Events:

$\langle \text{sl.send} \mid q, m \rangle$ { send message m to process q

$\langle \text{sl.deliver} \mid p, m \rangle$ { deliver a received message m from process p

Properties:

Stubborn delivery:

If a process sends a message m to process q , then m is infinitely often delivered at q .

No creation:

If some process q delivers some message m from p then process p has previously sent m to q .

2.3.2 Perfect point-to-point links

Events:

$\langle \text{sl.send} \mid q, m \rangle$

$\langle \text{sl.deliver} \mid p, m \rangle$

Properties:

Reliable delivery:

If a correct process sends a message m to a correct process q then q eventually delivers m

No creation:

If process q delivers some m from process p then p has sent m to q

At-most-once delivery:

Every message m is delivered at most once from p to q .

2.3.3 Alg. impl. perfect links (pl) from stubborn links (sl)

```

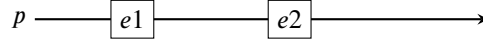
INIT:
 $\mathbb{D} \leftarrow \emptyset$ 
upon  $\langle pl.send \mid q, m \rangle$  do
    trigger  $\langle sl.send \mid q, m \rangle$ 
upon  $\langle sl.deliver \mid p, m \rangle$  do
    if  $(p, m) \notin \mathbb{D}$  then
         $\mathbb{D} \leftarrow \mathbb{D} \cup \{(p, m)\}$ 
        trigger  $\langle pl.deliver \mid p, m \rangle$ 
...

```

2.4 Timing Assumptions


- Asynchronous model (*Logical Timing*)

- **One Process**




If $e2$ happened after $e1$ in one process, we know the sequence of events.

- **Two Processes**



If we know that $e1$ caused $e2$, we know that $e2$ happened after $e1$.

- **Three processes**



Transitivity holds across processes, so if $e1$ caused \tilde{e} which cause $e2$, $e2$ happened after $e1$.

- Other time models exist

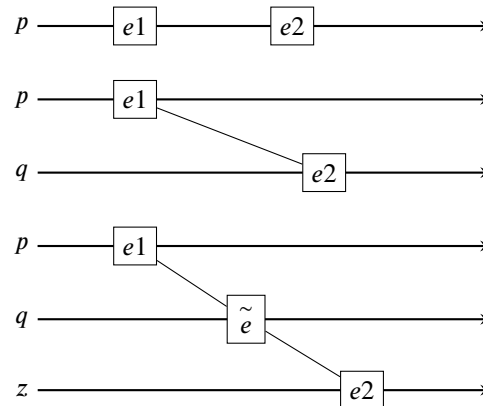
3 Timing Assumptions - March 3, 2020

3.1 Asynchronous System

Logical clock creates a logical time

- Each process p keeps a logical clock lp (initially 0)
- When an event e on p occurs, then $lp \leftarrow lp + 1$
- When p sends a message m to q , then p attaches a timestamp $ts(m) = lp$ to m
- When p receives a message m' with $ts(m')$, then p sets $lp \leftarrow \max\{lp, ts(m')\} + 1$

3.1.1 Happens-before relation



In each of these we can say that $e1$ happens before $e2$

3.1.2 Lemma

- $e1$ occurs at p at lp
- $e2$ occurs at q at lq
- $\Rightarrow e1 \rightarrow e2$, then $lp < lq$, but not the other way round!

3.2 Synchronous System

EITHER:

- Assume every process has access to a real-time clock (**RTC**)

OR:

- Synchronous computation (bounds on computation time)
- Synchronous communication (bounds on message-transmission time)

CAREFUL! when synchrony, assumptions are needed for safety properties



3.3 Partially Synchronous Model

- Synchronous most of the time
- When asynchronous, must not violate safety
Formally captured by abstraction of an eventually synchronous system.
- Initial period of asynchrony
- After some point in time (unknown to algorithm), system is synchronous

NOTE: Abstract model will remain synchronous forever after sync-point. In practice, periods of synchrony and asynchrony alternate.

3.4 Abstracting Time

DEFINITION: Perfect Failure Detecture \mathbb{P}

EVENT: $\langle \mathbb{P}.Crash \mid p \rangle$ denotes that process p has crashed.

PROPERTIES:

STRONG COMPLETENESS:

Eventually every process that has crashed is detected by all correct processes.

STRONG ACCURACY:

For any process p , if p detects that q crashed, then q has crashed.

Formally, all processes are either alive forever or they crash and stop.

Suppose a notion of time in \mathbb{N} :

$C : \mathbb{N} \rightarrow \Pi$, $C(t)$ denotes the processes that are live at time t .

$F : \mathbb{N} \rightarrow \Pi$, $F(t)$ denotes the proceses that are faulty (crashed) at time t .

$p \in F(t)$, then $\forall t' \geq t : p \in F(t')$ (crashes are irreversible)

$\mathbb{F} = \bigcup_{t \geq 0} F(t)$, set of all faulty processes

$\mathbb{C} = \Pi \setminus \mathbb{F}$, set of all correct processes

Strong Completeness:

$\exists t : \forall p \in \mathbb{F}, \forall q \in \mathbb{C} : \exists t' \geq t : \langle \mathbb{P}.Crash \mid p \rangle$ occurs on process q at time t' .

Strong Accuracy:

$\forall q \in \mathbb{C}$ if $\langle \mathbb{P}.Crash \mid p \rangle$ occurs on process q at time t then $p \in F(t)$.



3.4.1 Implementing \mathbb{P}

Initialization:

start timer Δ
alive $\leftarrow \Pi$
detected $\leftarrow \emptyset$

upon timeout do for all $p \in \Pi$ do
 if $p \notin \text{alive} \wedge p \notin \text{detected}$ then trigger $\langle \mathbb{P}.Crash \mid p \rangle$
 detected $\leftarrow \text{detected} \cup \{p\}$
 start timer with Δ
 alive $\leftarrow \emptyset$
 send msg [PING] to all $p \in \Pi$

upon receive msg. [PING] from p do
 send msg [PONG] to p

upon receiving [PONG] from p do
 alive $\leftarrow \text{alive} \cup \{p\}$

DEFINITION: Leader Election

EVENT: $\langle le.leader \mid p \rangle$, elects p to be leader

PROPERTIES (Eventual Leadership):

Eventually, some process l is elected leader by every correct process

ACCURACY:

If a process is elected leader then all previously elected leaders have crashed.

DEFINITION: Eventually Perfect Failure Detector

EVENTS:

$\langle \diamond \mathbb{P}.Suspect \mid p \rangle$, process p is suspected.

$\langle \diamond \mathbb{P}.Restore \mid p \rangle$, process p is thought to be alive.

PROPERTIES

STRONG COMPLETENESS:

Eventually, every process that has crashed is suspected by every correct process

EVENTUAL STRONG ACCURACY:

Eventually, every process that has crashed is suspected permanently by every correct process.

Model	Processes	Timing	
fail-stop	crash-stop	synchronous	$\langle \mathbb{P} \rangle$
fail-noisy	crash-stop	partially synchronous	$\langle \diamond \mathbb{P} \rangle, N > 2F$
fail-silent	crash-stop	asynchronous	$N > 2F$



4 4th Lecture - March 4, 2020

4.1 sub