

u^{*b*}

b

**UNIVERSITÄT
BERN**

Advanced Networking and Future Internet

II. Queue Management

Prof. Dr. Torsten Braun, Institut für Informatik

Bern, 21.09.2020

Advanced Networking and Future Internet: Introduction

Table of Contents

1. Introduction

1. Router Architecture
2. Congestion in the Internet

2. Queue Management

1. Why Queue Management in IP Routers ?
2. Requirements for Queue Management
3. Max-Min Fair Share
4. FIFO Scheduling
5. Priority Scheduling
6. Conservation Law
7. (Non-)Work-Conserving Disciplines
8. Token Buckets

3. Scheduling Disciplines


1. Generalized Processor Sharing
2. (Weighted) Round Robin
3. Deficit Round Robin
4. Weighted Fair Queuing
5. Class-Based Queuing
6. Earliest Due Date
7. Rate-Controlled Scheduling

4. Packet Dropping

1. Issues
2. Late and Early Dropping
3. Random Early Detection
4. RED with In and Out

1. Introduction


1. Router Architecture



A packet scheduler at the output port must select one packet among those queued for transmission.

1. Introduction

2. Congestion in the Internet




- Congestion in the Internet is still common.
 - Backbones, e.g., transatlantic links
 - Edge routers aggregating traffic
 - Access networks, e.g., modem lines, wireless links
- Output interfaces might be overloaded.
- Congestion (not transmission errors) is the main reason for packet loss in the Internet.

2. Queue Management

Components

- **Scheduling disciplines**
 - Which packet to serve next, and when
- **Dropping mechanism**
 - How and when to drop packets under varying load conditions
- **Buffer management**
 - How many queues and how much buffer space to use



2. Queue Management

1. Why Queue Management in IP Routers?

Best-Effort (BE) applications

- e.g., elastic applications such as e-mail or file transfer that can adapt to available resources
- Resources (bandwidth, buffers etc.) should be shared in a fair way among users / applications.
- Protection of connections from misbehaving connections

Non-BE applications

(some form of QoS guarantee)

- e.g., real-time applications, e.g., audio conferencing, video retrieval, transactions, interactive games, cloud computing (partially), that might require performance bounds (bandwidth, delay, jitter)
- Performance guarantees for applications by resource reservations

2. Queue Management

2. Requirements for Queue Management

- Ease of implementation
- Fairness and Protection (for BE)
 - Fairness: The allocation of a link capacity satisfies the max-min fair share allocation criterion.
 - Protection: Misbehavior of one connection does not affect performance of other connections.
- Performance bounds (for non-BE)
 - Arbitrary connection performance bounds, limited by the conservation law only.
- Ease and efficiency of admission control (for non-BE)
 - Limitation of number of connections

2. Queue Management


3. Max-Min Fair Share

Problem

- Dividing resources among a set of users
- Some users demand fewer resources than others.
- How to divide the resources left by these users ?


Max-Min Fair Share Allocation

- Resources are allocated in the order of increasing demand.
- No source gets a resource share larger than its demand.
- Sources with unsatisfied demands get an equal share of the resource.
- Allocation algorithm
 - Divide capacity C by n :
 C/n resources for each connection
 - Connection 1 needs x_1 resources ($x_1 < C/n$)
 - Distribute exceeding resources $C/n - x_1$ equally among other connections, so that each connection gets $C/n + (C/n - x_1)/(n-1)$ resources allocated
 - Continue process if resource allocation is larger than x_2



2. Queue Management


4. FIFO Scheduling



- First In First Out (FIFO) / First Come First Serve (FCFS)
- Packet transmission in the same order as packet arrival
- Packets are dropped if insufficient buffer space.
- + Simple implementation
- + Low packet processing delay
- No service differentiation
 - High-priority and/or delay-sensitive flows may be 'trapped' behind generic flows.
- Small packets can undergo large delays if they get behind big packets.
 - In general, flows with larger packets get better service.
- Greedy flows can overload the buffer.
 - Congestion-reactive flows get less and less bandwidth.

2. Queue Management

5. Priority Scheduling



- Priority levels with buffers
- Priority may depend on protocol type, application, IP addresses etc.
- Scheduler selects packet with highest priority.
- Packets with lower priority are selected only if there are no packets with higher priority. → starvation of low priority flows

2. Queue Management

6.1 Conservation Law

- A work-conserving scheduler is idle only if its queue is empty.
- FCFS is work-conserving
- Conservation Law:
$$\sum_{i=1}^N \rho_i \cdot q_i = \text{const}$$
- ρ_i : mean utilization of a link due to connection i
- q_i : connection i 's mean waiting time at the queue
- If a particular connection receives a lower delay than with FCFS, it must be at the expense of another connection.

2. Queue Management

6.2 Example: Conservation Law

2 connections A (10 Mbps) and B (25 Mbps) share a 155 Mbps link and each having a mean queuing delay with FCFS: 0.5 ms

New scheduling mechanism reduces A's queuing delay to 0.1 ms

$$\begin{aligned} & - 10 / 155 \cdot 0.5 + 25 / 155 \cdot 0.5 = \\ & \quad 10 / 155 \cdot 0.1 + 25 / 155 \cdot B \\ & - B = 0.66 \text{ ms} \end{aligned}$$

2. Queue Management


7. (Non-)Work-Conserving Scheduling Disciplines

- Non-work-conserving scheduling disciplines may be idle, even if there are packets to serve.
- A packet is only sent if it is eligible. Otherwise it is delayed until it becomes eligible.
- Are non-work-conserving scheduling disciplines useful ?
 - Downstream traffic can be made more predictable reducing buffer sizes (important in routers) and jitter.
 - Bursts are eliminated and traffic becomes smoother.
But jitter can also be compensated by playback buffers in end systems!
- Today's routers implement work-conserving schedulers mainly.

2. Queue Management

8. Token Bucket

- Mechanism for traffic shaping
- Maximum number of packets entering network during time interval t : $r \cdot t + b$



3. Scheduling Disciplines

For Best-Effort

- Generalized Processor Sharing (GPS)
- Round-Robin (RR)
- Weighted Round-Robin (WRR)
- Deficit Round-Robin (DRR)
- Weighted Fair Queuing (WFQ)

For Guaranteed Service

- Weighted Fair Queuing (WFQ)
- Delay Earliest Due Date (EDD)
- Jitter Earliest Due Date (J-EDD)
- Rate-Controlled Scheduling
(non-work-conserving)

3. Scheduling Disciplines

1. Generalized Processor Sharing

- serves packets as if they are in separate logical queues.
- visits each non-empty queue in turn.
- serves an infinitesimally small amount of data in each queue so that it can visit each queue at least once in any finite time interval.


Characteristics

- + Fair
- + Ideal
- Not implementable
- Work-conserving

3. Scheduling Disciplines

2. (Weighted) Round Robin

- Several queues per output, e.g., 1 queue per sender or flow
- **RR**: Packets are taken from each queue in turn.
- **WRR**: Queues have different weights.
- Characteristics
 - + Aggressive flows get punished.
 - Mean packet sizes must be known in advance.
 - Fairness over a time interval much longer than a round time



3. Scheduling Disciplines

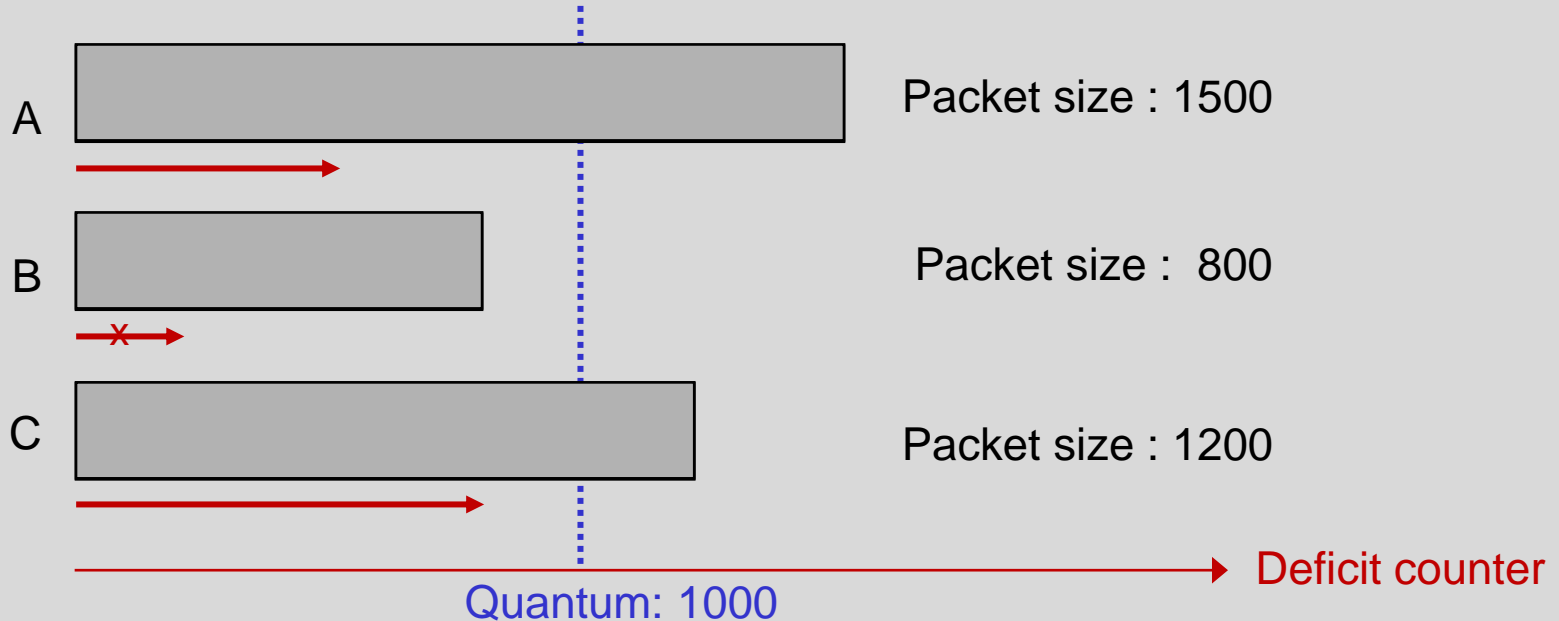
3.1 Deficit Round Robin

- DRR supports variable packet sizes.
- Each queue has a deficit counter initialized by 0.
- Scheduler visits each queue in turn and tries to serve 1 quantum of bits.

```
if (there is a packet to be served)
    if (packet_length <
        (quantum + deficit_counter)) {
        serve packet;
        deficit_counter +=
            (quantum - packet_length);
    }
    else
        deficit_counter += quantum;
else
    deficit_counter := 0;
```

3. Scheduling Disciplines

3.2 Example: DRR



3. Scheduling Disciplines

4. Weighted Fair Queuing

Approximation of GPS

- emulates a weighted bit-by-bit round-robin discipline, i.e., generalised processor sharing
- Virtual starting and finishing ‘times’ are calculated as if GPS would be used.
- The next packet to be served is the one with the smallest finishing time (at the time of calculation).

Implementable GPS version providing different amounts of capacity to different flows.

- WFQ closely approximates the properties of GPS.
- WFQ provides uniform and appropriate level of service to each flow.


3. Scheduling Disciplines

4.1 WFQ: Finish Number

- Calculation of a finish number (\neq finish time !) indicating the order, in which packets are served assuming a **bit-by-bit service**
- Round number = number of rounds a bit-by-bit RR scheduler has completed at a given time (variable duration of rounds !)
- A connection is active if (largest finish number in the connection's queue or of the packet last served) > current round number
- Finish number
 - of a packet arriving at an inactive connection = current round number + packet size [bits]
 - of a packet arriving at an active connection = largest finish number in the queue or of the packet last served + packet size [bits]
 - $F(i, k, t) = \max \{F(i, k - 1, t'), R(t)\} + P(i, k, t) / \phi(i)$
 - $F(i, k, t)$: finish number of k^{th} packet of connection i arriving at time t
 - $P(i, k, t)$: packet length of k^{th} packet of connection i arriving at time t
 - $R(t)$: number of the round at time t
 - $\phi(i)$: weight of connection i
- Packets are ordered by their finish number and are served in that order.

3. Scheduling Disciplines

4.2 Token Buckets and WFQ



3. Scheduling Disciplines

4.3 WFQ: Delay Bound

Minimum throughput for flow i
by WFQ = $R \cdot \phi(i) / \sum \phi(j)$


Maximum delay for flow i ?

- Assumption: token bucket i is full.
- Then, burst of b_i packets arrives and consumes all tokens.
- b_i packets are served with rate of at least $R \cdot \phi(i) / \sum \phi(j)$.
- Last packet will experience maximum delay = $b_i / (R \cdot \phi(i) / \sum \phi(j))$.

3. Scheduling Disciplines

5. Class-Based Queuing

- Hierarchical Link Sharing
- Implementation by WFQ, DRR
- Definition of bandwidth available for each class
- Class =
 - Protocol / application
 - Traffic from / to certain sub-networks



3. Scheduling Disciplines

6. Earliest Due Date

- Assignment of a deadline (= arrival time + period of traffic stream) to each packet
- Scheduler selects packets in the order of their deadlines.
- Packets missing their deadline are discarded.

3. Scheduling Disciplines

6.1 Delay-EDD

- Negotiation of a service contract (in addition to EDD):
If a source obeys a peak rate, the delay is less than some delay bound.
- Deadline =
arrival time + delay bound
- Delay-EDD requires
 - admission control and
 - bandwidth reservation at peak rate.
- Example
 - One packet arrives every 0.2 s with a delay bound of 1 s
 - Deadline of k^{th} packet = $(0.2k+1)$ s

3. Scheduling Disciplines


6.2 Jitter-EDD

- Extension of Delay-EDD: Delay-jitter-regulator precedes EDD scheduler.
- Packets get timestamp = deadline - finishing time.
- Next server delays packet by timestamp value.
- All packets receive the same delay at every hop (except at the last hop).

3. Scheduling Disciplines

7. Rate-Controlled Scheduling

- can provide bandwidth and delay bounds to connections
- Regulator
 - determines the packet's eligibility time.
 - forwards eligible packets to scheduler, e.g., FIFO, WFQ, EDD.



4. Packet Dropping

1. Issues

Limited queue length requires packet dropping by scheduler in overload conditions.

- Degree of aggregation
 - Protection versus state information
- Marking of several drop priorities
 - Packets with high drop priorities are dropped first.
 - Either applications or routers must be able to mark packets.
 - Low dropping priority for packets that traveled very long



- Late / early dropping
 - Early drop: Packets are dropped even in case of non-full queues. Cooperative sources experience low packet loss.
- Drop position
 - Tail: easy to implement, may be unfair
 - Head: more complex, sources notice dropping earlier
 - Random: very complex, fair

4. Packet Dropping

2. Late and Early Dropping

Late Dropping

- Arriving packets are discarded, if queue is full.
- Certain flows can monopolize the queue.
- Higher delays
- Queues can not absorb bursts.
- Reaction onto congestion, but no congestion avoidance
- Global synchronisation
 - When queue is full, all packets from all flows are dropped.
 - All flows back-off simultaneously.
 - Periods of congestion followed by periods of low network utilisation

Early Dropping

- Flows (e.g., TCP connection) reduce rate before queue becomes overloaded.
- Examples
 - Early Random Drop:
Arriving packet is dropped with fixed probability, if queue level exceeds a certain limit.
 - Random Early Detection

4. Packet Dropping


3.1 Random Early Detection

```
avg := exponential average of queue length;  
if (avg < THmin)  
    store packet in queue  
else if (THmin ≤ avg < THmax) {  
    P := calculate dropping probability;  
    discard packet with probability P;  
    store packet with probability 1-P}  
else if (avg ≥ THmax)  
    discard packet;
```


4. Packet Dropping

3.2 Random Early Detection

Dropping probability increases with the exponential average queue length.

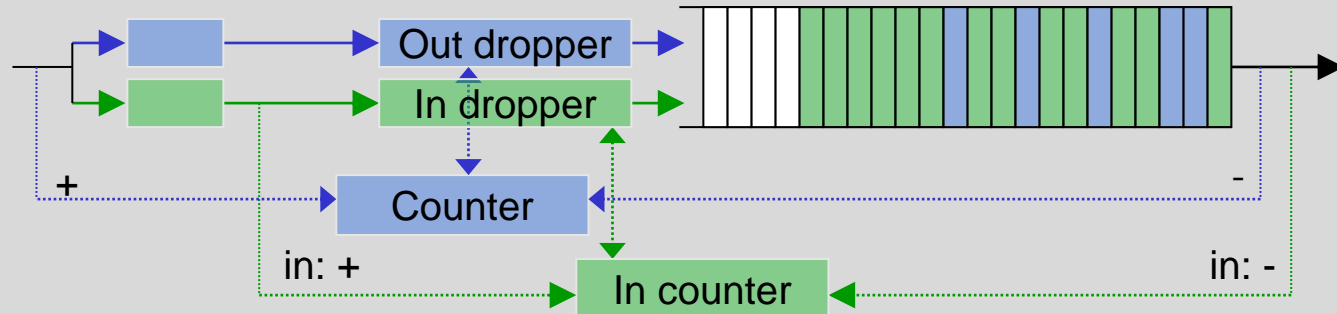


4. Packet Dropping

4.1 RED with In and Out (RIO)


Different droppers for high / low priority (in-service / out-of-service) packets

- **Out dropper** drops earlier and more aggressively.
- **In dropper** considers high-priority packets only.



4. Packet Dropping

4.2 RIO



Thanks

for Your Attention

Prof. Dr. Torsten Braun, Institut für Informatik

Bern, 21.09.2020

u^b

^b
UNIVERSITÄT
BERN

