

Docker

An Introduction to Application Deployment

University of Fribourg
Department of Informatics
Software Engineering Group



October 29, 2020

1. Introduction
2. Concepts
3. Usage
4. Example
5. Conclusion

- Two important aspects of software applications are:
 - Reproducibility: How do I make sure that my application is functional, regardless of where it is running or who is running it?
 - Robustness to change: How do I make sure that my application will still be running if I change something?



Neon Bones
@ianholmes

Follow



You can download our code from the URL supplied. Good luck downloading the only postdoc who can get it to run, though
[#overlyhonestmethods](#)

8:52 AM - 8 Jan 2013

<https://twitter.com/ianholmes/status/288689712636493824>

- Possible challenges to those aspects:
 - It may be complicated to set up the environment for the application (e.g. server setup).
 - The steps to set up the application may be unclear or difficult to reproduce (e.g. server migration, releasing to clients).
 - Different applications running on a same environment may need different versions of a library.
 - Updating server libraries may break retro-compatibility of applications.

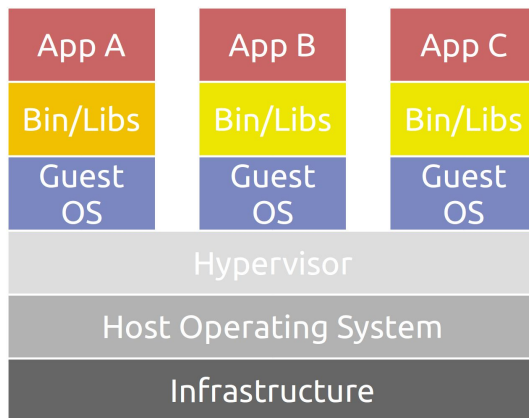
- Is there a solution?

Introducing Docker

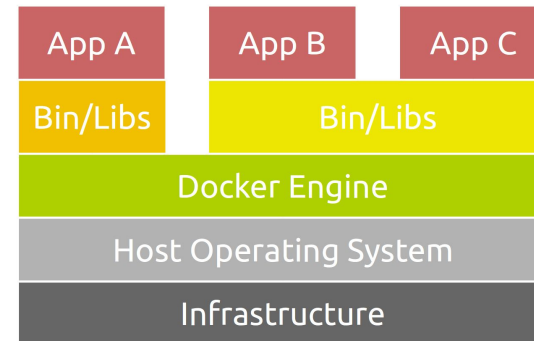


1. Introduction

- Open-source project released in 2013.
- Simplifies software deployment by guaranteeing consistency among different running environments.
 - Docker creates an abstraction of an application.
 - Applications are self-contained in (sort of) minimalistic Virtual Machines.



Virtual machines



Docker

Concepts

Images and Containers

2. Concepts

- Conceptually, Docker mainly consists of **images** and **containers**.
 - **Images:**
 - Are virtualized application
 - Are immutable
 - **Containers:**
 - Are an “instance” of an image
 - Can be stopped and restarted
 - Are given an arbitrary name when created

- Images are composed by a list of read-only filesystem layers.
- Each layer only contains the diff from the previous layer.
- Each layer is identified using a secure hash.
- A container is an extra writable layer on top of an image.
- There exist layers with no changes in filesystem (0 bytes).




Image Layers II

2. Concepts

- Images are shared between different containers.
- Image layers can also be shared by different images.
 - Great for reusability and for saving memory.

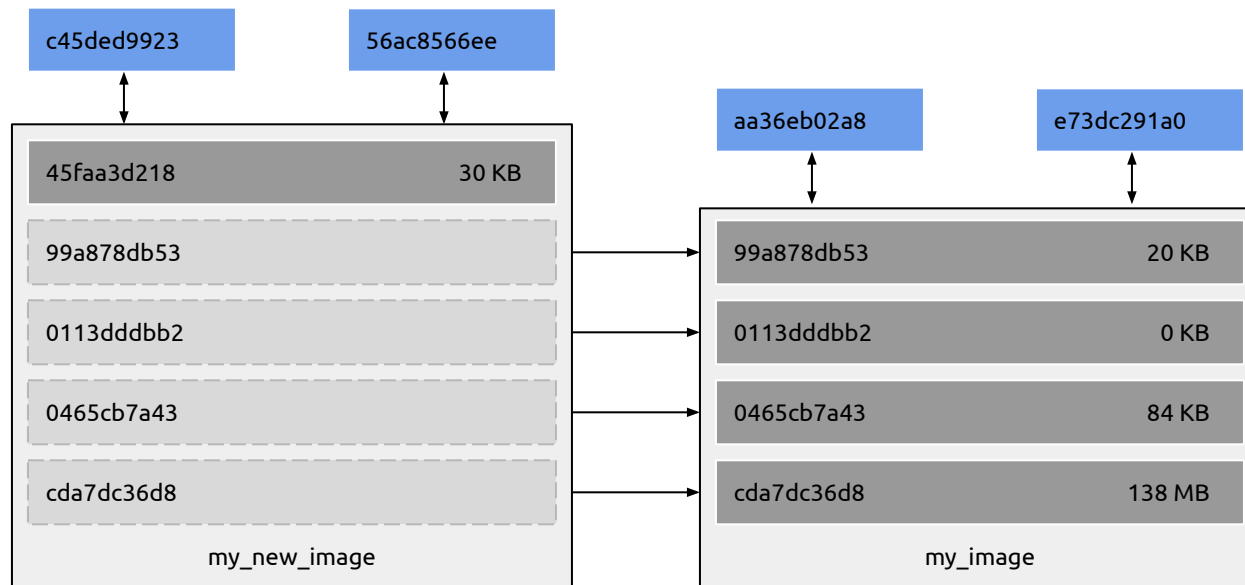



Image Layers III

2. Concepts

- When you pull an image based on an existing image on your machine, Docker uses layer hashes to only pull the missing layers.
- It is not possible for an image to use only the top-layers of another image, since they are only diffs from the lower ones.



Data Access I

R/W Filesystem

2. Concepts

- When a running container needs to access program data of its image, Docker searches from the top layers until it finds them.
 - If the container then needs to change these data, docker creates a local copy of them in the container layer.
- Containers are not supposed to be long-lived and should not be used to store actual data (i.e. databases, config files, etc).
 - Instead, data volumes are used.


Data Access II

Data Volumes

2. Concepts

- Data volumes are directories (or files) on the host system.
- They have several purposes:
 - Data persistence after a container is deleted
 - Data exchange between a container and the host
 - Data exchange between multiple containers
- Data volumes are mounted directly into one or several containers.
- Multiple data volumes can be mounted in a single container.

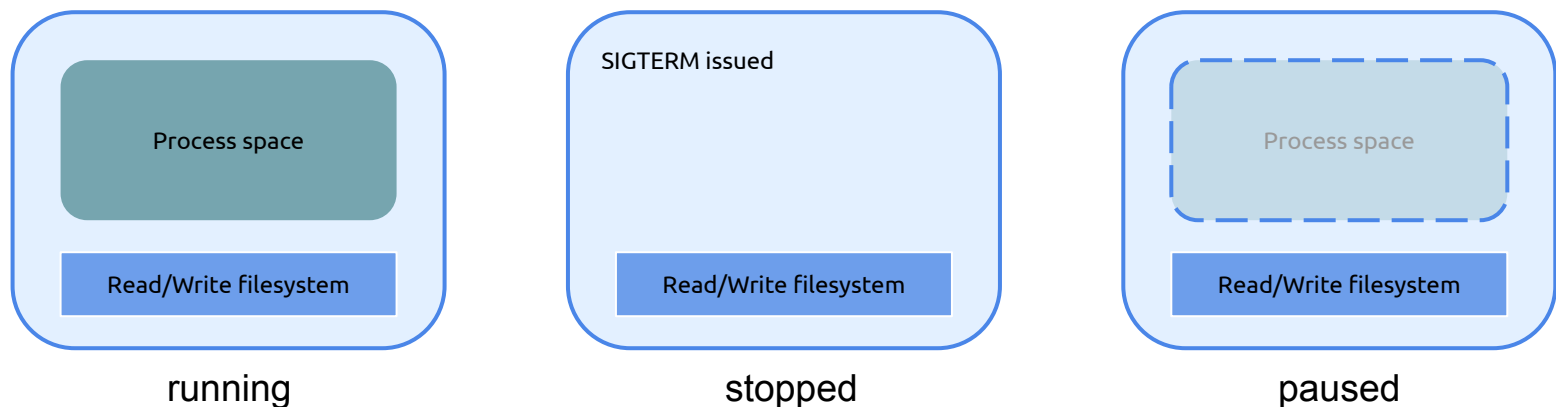
- Ports are used to communicate with a container.
 - As most applications are web-based, this method is well-suited.
- The ports defined inside the container have to be mapped to the outside.
 - This is essential to avoid conflicts.



Container States

2. Concepts

- A container can be in a number of states, mainly:
 - Running
 - Stopped (or Exited)
 - Paused
- A running container has its own process space.
- Containers can be paused, maintaining their process space.



Building Images I

Dockerfile

2. Concepts

- Images are built using a *Dockerfile*, a set of instructions.
- Each command corresponds to an image layer.
- Your *Dockerfile* should be placed alongside your code.

Base image	→	FROM node:7.0.0
Information about the author	→	MAINTAINER pascal.gremaud@unifr.ch
Environment variables	→	ENV appdir /usr/src/app/
Execute an UNIX command	→	RUN mkdir -p \$appdir
Set path for the image and dockerfile	→	WORKDIR \$appdir
Copy file or folder into image	→	COPY package.json . RUN npm install COPY . .
Expose a port outside the container	→	EXPOSE 8080
Command to execute in the container	→	CMD ["node", "."]

Building Images II


Commands and Layers

2. Concepts

- The “CMD” command is a single, default command to execute when the container starts.
- If an image has a “CMD” command, it will override any “CMD” command from the images below it.
- Even the commands with no effect on the diffs create a layer.
- Here is the result of the docker history command on node:7.0.0

IMAGE	CREATED	CREATED BY	SIZE
04c0ca2a8dad	11 hours ago	/bin/sh -c #(nop) CMD ["node"]	0 B
<missing>	11 hours ago	/bin/sh -c curl -SLO "https://nodejs.org/dist...	45.87 MB
<missing>	11 hours ago	/bin/sh -c #(nop) ENV NODE_VERSION=7.0.0	0 B
<missing>	11 hours ago	/bin/sh -c #(nop) ENV NPM_CONFIG_LOGLEVEL=in...	0 B
<missing>	11 hours ago	/bin/sh -c set -ex && for key in 9554F0...	108.3 kB
<missing>	11 hours ago	/bin/sh -c groupadd -r node && useradd -r -g ...	330.4 kB
<missing>	3 days ago	/bin/sh -c apt-get update && apt-get install ...	318.3 MB
<missing>	13 days ago	/bin/sh -c apt-get update && apt-get install ...	122.6 MB
<missing>	13 days ago	/bin/sh -c apt-get update && apt-get install ...	44.3 MB
<missing>	13 days ago	/bin/sh -c #(nop) CMD ["/bin/bash"]	0 B
<missing>	13 days ago	/bin/sh -c #(nop) ADD file:23aa4f893e3288698c...	123 MB

overridden



- Registries are used to share images.
- Different versions of an image are grouped in a repository, as with git.
- The most commonly used registry is *Dockerhub* (www.dockerhub.com)
 - It is used by default when pulling an image.
- You can also push your own built images to *Dockerhub*.


Docker Compose I

2. Concepts

- Sometimes it is better (or needed) to separate application components in different containers.
- Docker Compose is used to orchestrate these elements.
- It uses a *docker-compose.yml* file.
- If a container uses a service in another container (e.g. a database), some extra scripts may be needed to ensure that the service is already up and running before the application starts.

```
first_container:  
  build: my_custom_image  
  ports:  
    - 8002:8080  
  depends_on:  
    - second_container  
second_container:  
  image: image1  
  ports:  
    - 8001:8080  
  volumes:  
    - "/data:/data/db"
```

- Docker starts your containers one after the others, without waiting.
 - This can lead to dependency issues (e.g. a container trying to communicate via another, not set up yet container).
 - This issue needs to be solved in your code, not by Docker.



- Networks
 - It is possible to ask docker to create networks for “direct” communication between containers, even on different hosts.
- Docker Swarm
 - Grouping of several distant host machines into a single virtual host in order to scale up applications.
- Automated builds with *Github* (or *Bitbucket*)
 - Image automatically built and published on *Dockerhub* (not possible to use with Docker-Compose).

1. Introduction
2. Docker - Concepts
3. **Docker - Usage**
4. Example

Containers Manipulation I

Basic Commands

3. Usage

- Docker commands: `docker [options] [command] [arg..]`
- Starting a container (and possibly pulling it):
`docker run -d -p <host_port>:<container_port> <image>`
 - `-d` is used to run in detached mode
 - `-p` is used to map ports
- Once started, the container is associated with a hash and a name.
 - To interact with a container, use either of them.
- See the list of non-stopped containers:
`docker ps`
 - `-a` to show all containers
 - `-q` to show only the hash of each container

Containers Manipulation II

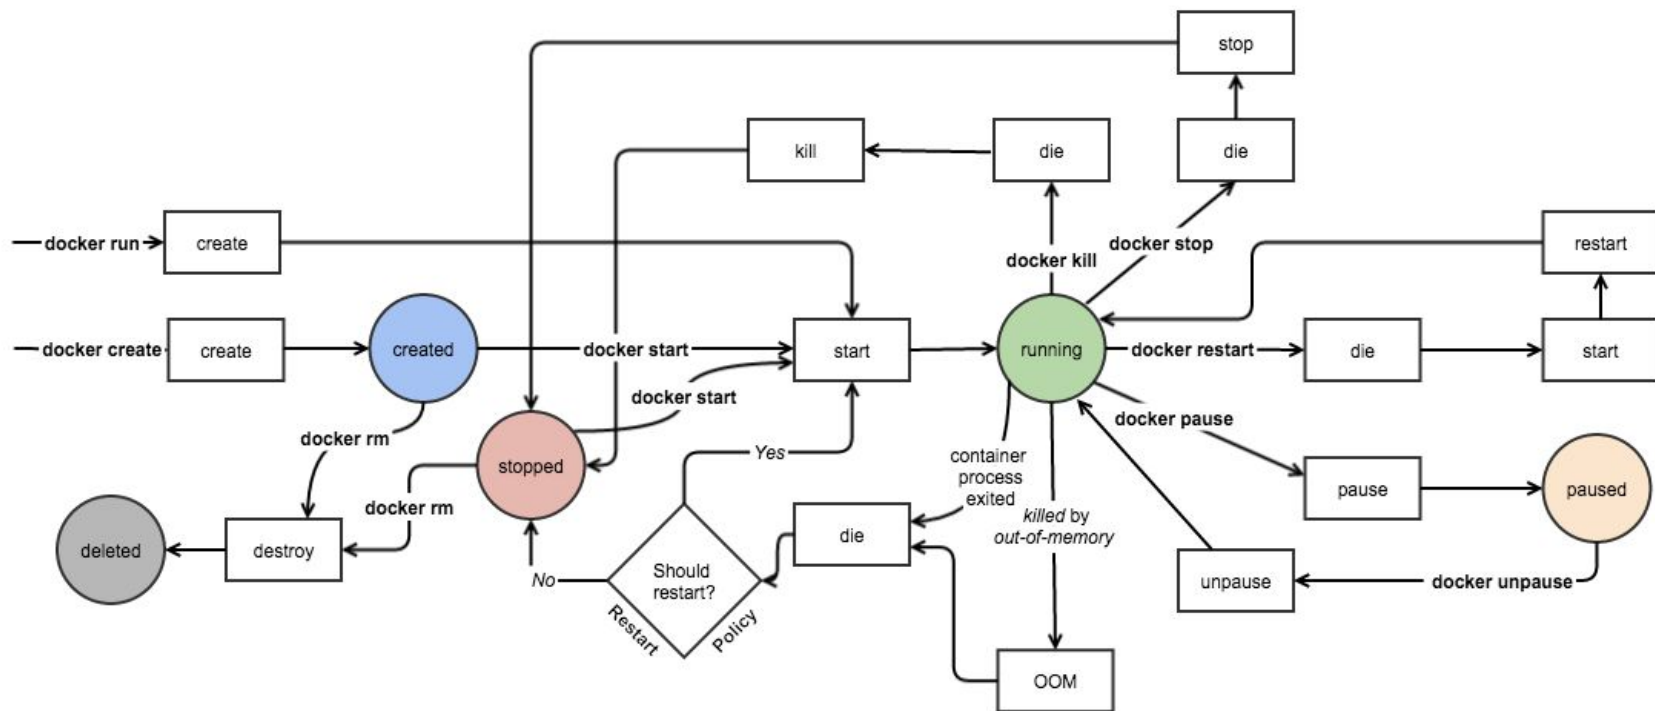
Other docker Commands

3. Usage

- create: create a new container from an image (no running)
 - If no local image is found and no registry is specified, the image is pulled from *Dockerhub*.
- run: create a new container from an image and run it
 - Same pulling mechanism
- rm: destroy a **stopped** container
- start: run a **stopped** container
- stop: stop a **running** container
- restart: stop and start a **running** container
- pause: pause a **running** container
- unpause: unpause a **paused** container

Containers Manipulation III

Container States and Transition Commands



<https://medium.com/@nagarwal/lifecycle-of-docker-container-d2da9f85959>

- Stop all containers:
 `docker stop $(docker ps -q)`
 - `$(docker ps -q)` gives the list of non-stopped containers
- Destroy all containers (stopped):
 `docker rm $(docker ps -qa)`
- Execute a command inside a container (not recommended):
 `docker exec <options> <container> <command>`
- Open a terminal inside a container (debug only!!!):
 `docker exec -ti <container> /bin/bash`

- Download an image:
docker pull <image>
- See the list of “top-level” images:
docker images
 - -a to show all images
 - -q to show only the hash of each image
- Remove an image:
docker rmi <image>
- Remove all images:
docker rmi \$(docker images -qa)
- Get the stack of an image:
docker history <image>

Images Manipulation II

Creating Images

3. Usage

- In order to build your image, run this command:
`docker build -t <image_name>[:<tag_name>] <location>`
- For instance, when inside your project at the level of your *Dockerfile*:
`docker build -t docker-node-example:latest .`
- You can then push your image to a registry (e.g. *Dockerhub*):
`docker login`
`docker tag <image>:<tag> <username>/<repo>[:<tag>]`
`docker push <username>/<repo>[:<tag>]`

1. Introduction
2. Docker - Concepts
3. Docker - Usage
4. **Example**

Step 1

Setup

4. Example

- In this example we will:
 - Build an image
 - Push it to Dockerhub
 - Create an automated build for it
- By now you should have:
 - Docker installed on your machine
 - An account on *Github*
 - An account on *Dockerhub*
 - Pulled the node:7.0.0 image

Step 2

Creating a Node App

4. Example

- Go to <https://github.com/polchky/docker-node-example> and fork the repository.
- Clone your forked repository on your machine and cd into it.
- We will not build our Nodejs app locally but directly create an image.

Step 3

Building an Image

4. Example

- Study the *Dockerfile*.
- To build your image, open a terminal and run `docker build -t docker-node-example .`
- Your image should now appear in docker.
- Run your image by executing `docker run -d -p 9000:8080 docker-node-example`
 - You can now access your app at <http://localhost:9000>
 - On OSX, you need to replace localhost by the IP of your docker machine.
 - Get it by executing `docker-machine config default`
- Play around with your container(s), exec into it.

Step 4

Pushing to *Dockerhub*

4. Example

- Login to *Dockerhub*:
docker login
- Tag your image:
docker tag docker-node-example \
<username>/docker-node-example
- Push your image to *Dockerhub*:
docker push <username>/docker-node-example

Step 5

Creating an Automated Build

4. Example

- On *Dockerhub*, go to your account settings, and then *Linked Accounts*.
 - Link your Github account, with read-write access.
- Go to your pushed docker image, *Builds*.
- Create a new automated build using Github.
- Select your github username and your docker-node-example.
- Take a look at the behavior customization.
- Save the automated build.
- Each time you push to the master branch (default behavior) on Github, a build of your image is triggered on *Dockerhub*.
 - You can also trigger a build manually.

- Docker and its tools suite can be a noticeable improvement to standard software delivery and orchestration methods.
- Docker may be still young, but is in active development and has a consequent community.
- You may consider deploying and delivering your project using Docker and Docker-Compose.
 - This will be considered as an extra “feature” of your project.
 - However, be careful to include the “dockerization” time in your planning to avoid any surprises.

External Resources

Some Useful Links

5. Conclusion

- The official docker documentation:
<https://docs.docker.com>
- Docker commands cheat sheet:
<https://github.com/wsargent/docker-cheat-sheet>
- A good explanation and visual representation of the concepts of layers and processes:
<http://merrigrove.blogspot.ch/2015/10/visualizing-docker-containers-and-images.html>