

u^b

b

**UNIVERSITÄT
BERN**

Internet of Things

III. Operating Systems

Prof. Dr. Torsten Braun, Institut für Informatik

Bern, 08.03.2021

Internet of Things: Operating Systems

Table of Contents

1. Introduction
 1. Platform Requirements
 2. Node Level Platforms
 3. Concurrency
 4. Event- and Thread-driven Execution
2. TinyOS
 1. Overview
 2. Component-based Architecture
 3. nesC
 4. Examples: Timer, Blink, FieldMonitor
 5. Program Execution
3. Contiki OS
 1. Overview
 2. Event-Driven Multitasking Kernel
 3. Protothreads
 4. uIP Stack
4. MANTIS OS
 1. Overview
 2. Memory and Thread Management
 3. Power Management and Dynamic Reprogramming
5. RIOT
 1. Overview
 2. Architecture
 3. Kernel
6. Other IoT Operating Systems
 1. Mbed
 2. Google Android Things
 3. Zephyr
 4. ReeRTOS
7. Operating System Power Management
 1. Low-Energy Earliest Deadline First
 2. Low-Energy Device Scheduling

1. Introduction

1. Platform Requirements

- Low power consumption
- Concurrency-intensive operation
- Reactivity,
at least soft-real-time capabilities
- Robustness and self-configurability
- Flexibility: programmability
and reconfiguration
- Compatibility
- Security and privacy
- Memory limitations
→ low memory footprint
- Network interoperability
- System interoperability

1. Introduction

2. Node Level Platform Options

Node-centric operating system providing hardware and network abstractions of a sensor node to programmers

- Traditional operating systems: file management, memory allocation, task scheduling, device drivers, networking etc.
- For sensor nodes: simplified versions of traditional operating systems
- Example: MANTIS OS, RIOT

Node-level programming tools

- Language platform providing library of components to programmers
- Example: TinyOS, Contiki

1. Introduction

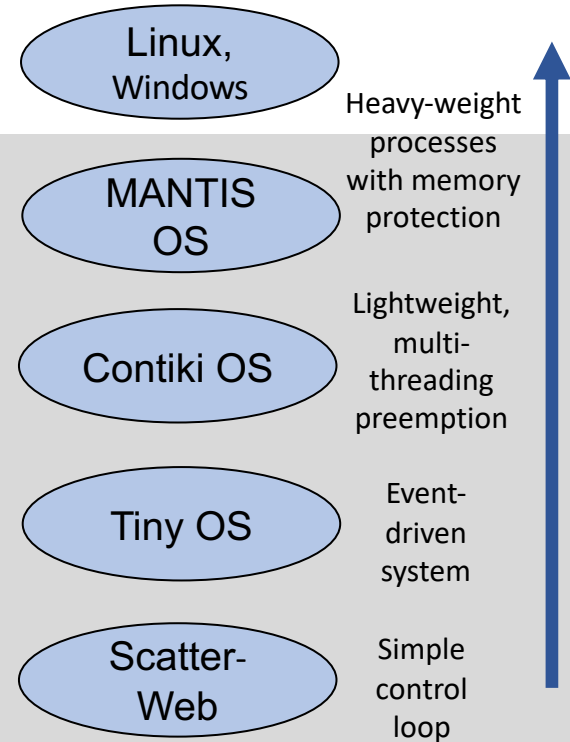
3.1 Concurrency

True Concurrency vs. Pseudo-Concurrency

- Low-cost networked embedded systems usually have only 1 CPU core.
- Concurrent activities are mapped to pseudo-concurrent sequential processes.
- True concurrency requires true parallel processing units (CPU cores).

Concurrency

- Ability to provide concurrency “costs” CPU and memory.
- First sensor node OSs had almost no concurrency capabilities.
- OSs based on simple control loops proved to be inflexible and died, e.g., ScatterWeb.

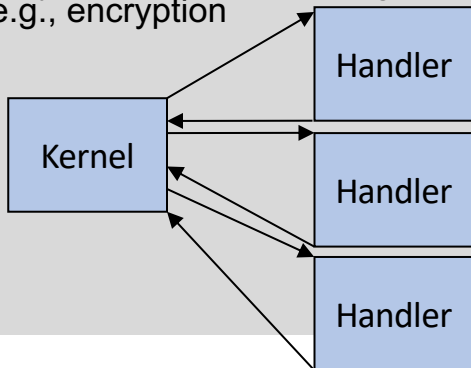


1. Introduction

3.2 Concurrency

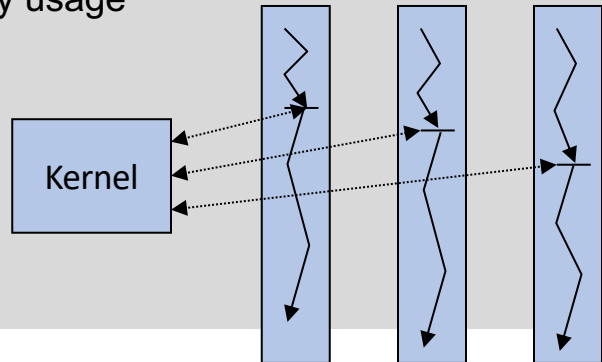
Event-driven

- Processes do not run without events.
- Kernel invokes event handler when event occurs.
- Only one event can run at a time!
- Event handler runs to completion (explicit return).
 - may be a problem for long computations, e.g., encryption



Multi-Threading

- Blocked threads are waiting for events.
- Kernel unblocks threads when event occurs.
- Thread runs until next blocking statement.
- Each thread requires own stack.
- Larger memory usage

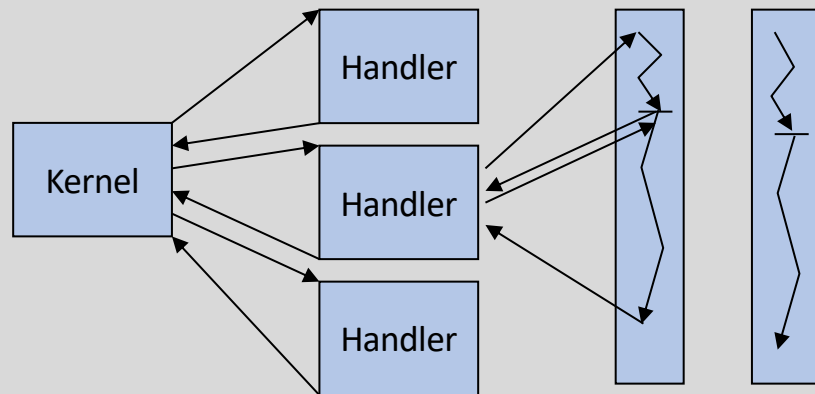


1. Introduction

3.3 Concurrency

Mix: Threads on top of event-driven kernel

- Kernel is event-based, invokes event handler when an event occurs.
- Multi-threading implemented as library
- Threads only used when needed



1. Introduction

4. Event- and Thread-Driven Execution

Event-driven Execution

- + (Pseudo) Concurrency with low resources
- + complements the way networking protocols work
- + Inexpensive scheduling technique
- + Highly portable
- Event-loop is in control.
- Program needs to be chopped to subprograms.
- Bounded buffer producer-consumer problem
- High learning curve

Thread-driven Execution

- + eliminates bounded buffer problem
- + Programmer in control of program
- + Automatic scheduling
- + Real-time performance
- + Low learning curve
- + simulates parallel execution
- Complex shared memory
- Expensive context switches
- High memory footprint
- Not portable due to stack manipulation
- performs better on multiprocessors

2. TinyOS

1. Overview



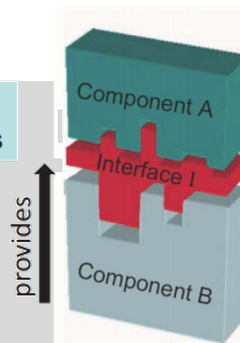
- 1st OS for sensor network applications on resource-constrained HW platforms.
- classical event-based OS
- Simplifications
 - No file system
 - Static memory allocation
 - Implementation of a simple task model
 - Thread support by application level thread library using standard synchronization mechanisms, e.g., semaphores, condition variables
 - Minimal device and network abstractions
- Language-based application development approach
 - Only necessary parts of operating system are compiled with application.
 - or: Each application is built into the operating system.
- Concurrency
 - Concurrency issues are left to the programmer by large extent.
 - Event model allows to handle concurrency in a light-weight fashion → to avoid processor idle times.
 - TinyOS avoids multi-tasking overhead.

2. TinyOS

2. Component-based Architecture

- Component-based approach
 - Components are organized into layers.
 - Component = Interface (bidirectional) + implementation
 - Components encapsulate software functionalities.
 - Some components are thin wrappers around hardware.
 - Components are typically implemented as reentrant state machines.
 - All variables are inside components (no dynamic memory allocation).
 - TinyOS provides system software components as set of libraries.
- TinyOS application = scheduler + graph of components
 - Applications are typically developed using a special language (nesC) and wire components together

can signal events,
must implement commands



can call commands,
must implement event handlers

2. TinyOS

3. nesC

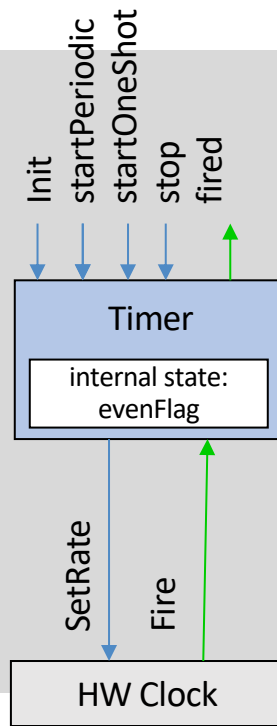
- C extension
- Static language
 - no dynamic memory allocations
 - call graph known at compile time
- Direct support of event-based TinyOS design
- nesC applications are built out of components with well-defined interfaces.
 - Component provides and uses interfaces, which are the only point of access to a component.
 - Bidirectional interfaces: commands and events
- Types of components
 - *Modules* provide application code, implementing one or more interfaces
 - *Configurations* are used to wire other components together by connecting component interfaces.

2. TinyOS

4.1 Example: Timer Component

- Timer can set rate of the clock (wrapper around hardware clock).
- Hardware clock generates periodic interrupts and toggles internal state of timer.
- Timer component fires dependent on internal state.

```
interface Timer<precision_tag> {  
  
    event void fired();  
  
    command void startPeriodic(...);  
    command void startOneShot(...);  
    command void stop();  
    ...  
}
```



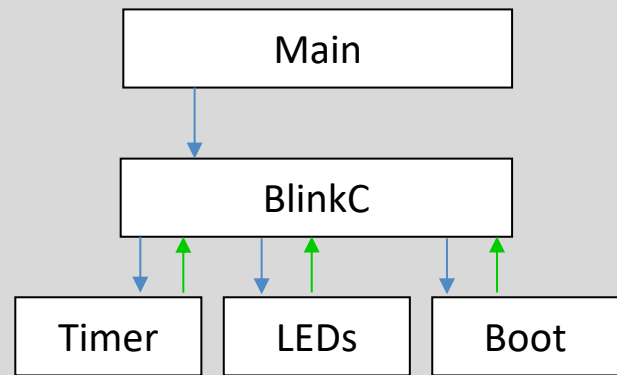
2. TinyOS

4.2 Example: Blink Application

Task: Blink the red LED every 1 second

Implementation: On boot, start a 1 second timer. On timer fire (countdown at 0): toggle state of red LED

```
module BlinkC {  
  uses interface Timer<TMilli>  
    as BlinkTimer;  
  uses interface Leds;  
  uses interface Boot;  
}  
implementation{  
  event void Boot.booted() {  
    call BlinkTimer.startPeriodic(1000);  
  }  
  event void BlinkTimer.fired() {  
    call Leds.led0Toggle();  
  }  
}
```

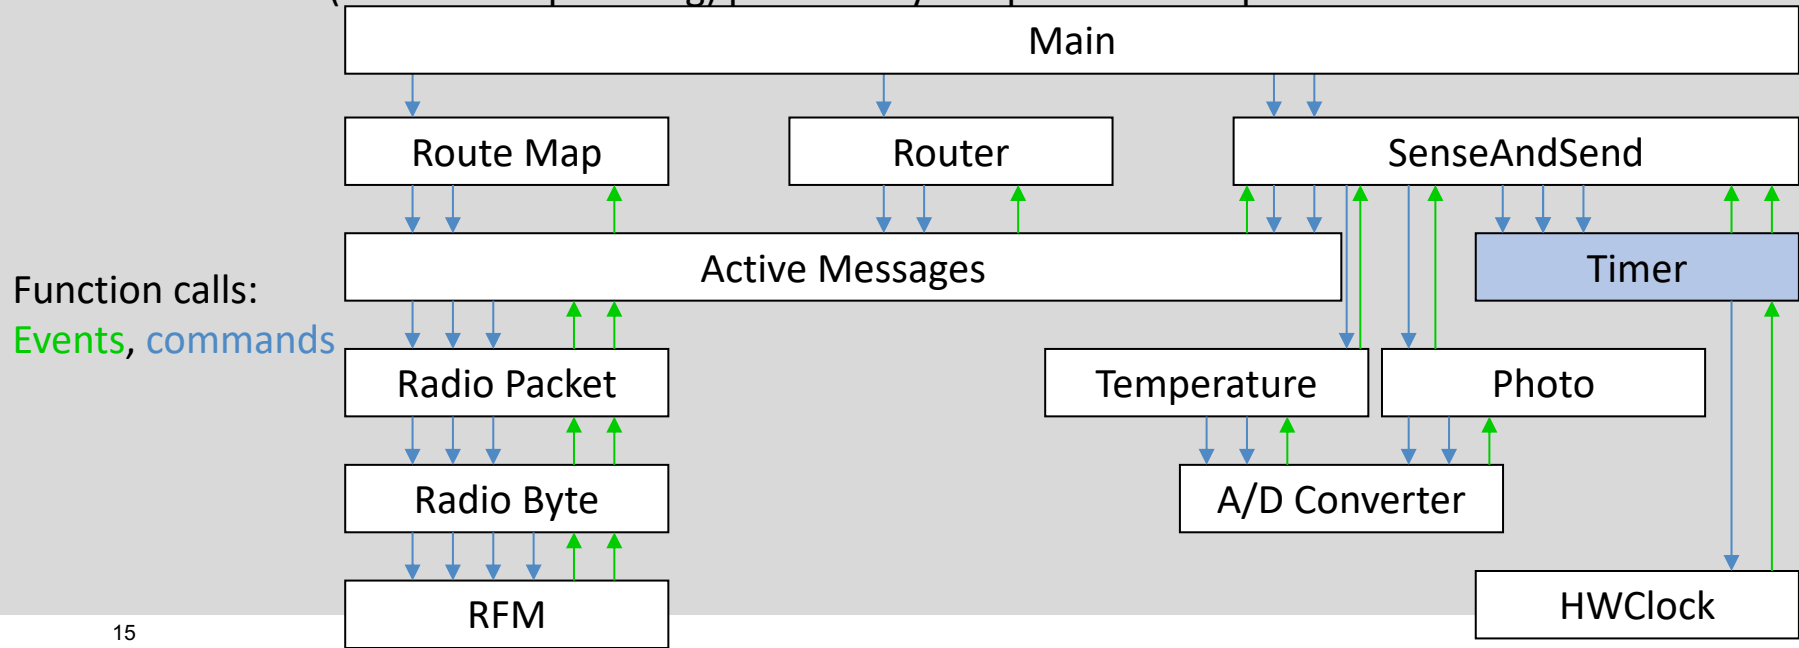


Function calls: Events, Commands

2. TinyOS

4.3 Example: FieldMonitor Application

Nodes send (via multi-hop routing) periodically temperature and photo sensor data to a base station.



2. TinyOS

5.1 Events and Tasks

A program executed in TinyOS has two contexts

1. Events

- Interrupt sources: clock, digital inputs, radio chip, etc.
- Execution of interrupt handler = event context
- Processing of events runs to completion, but preempts tasks and can be preempted by other events (last in first out, LIFO).
- Programmers are required to chop code (in particular in event contexts) into smaller execution pieces to avoid long blocking of tasks

2. Tasks

- are created by a component and posted to a task scheduler.
- are deferred computations.
- always run to completion without preempting other tasks or being preempted by other tasks.

Default TinyOS scheduler

- maintains a task queue and invokes tasks in posting order (FIFO).
- puts node into sleep state if no tasks are available in the queue.

2. TinyOS

5.2 Split-Phase Operation

Separation of **method call initiation** and **return of call** (similar to asynchronous method / function calls)

- Client call returns immediately without performing body of call.
- Server executes operation later.
- Server signals completion by calling an event handler in client component.

Example: Packet transmission (in Active Messages component) would block system for a long time.

- **send()** returns immediately.
- Message transmission: conversion of packet to bytes and bits, driving radio circuit
- Caller is notified by **sendDone()** about completion.

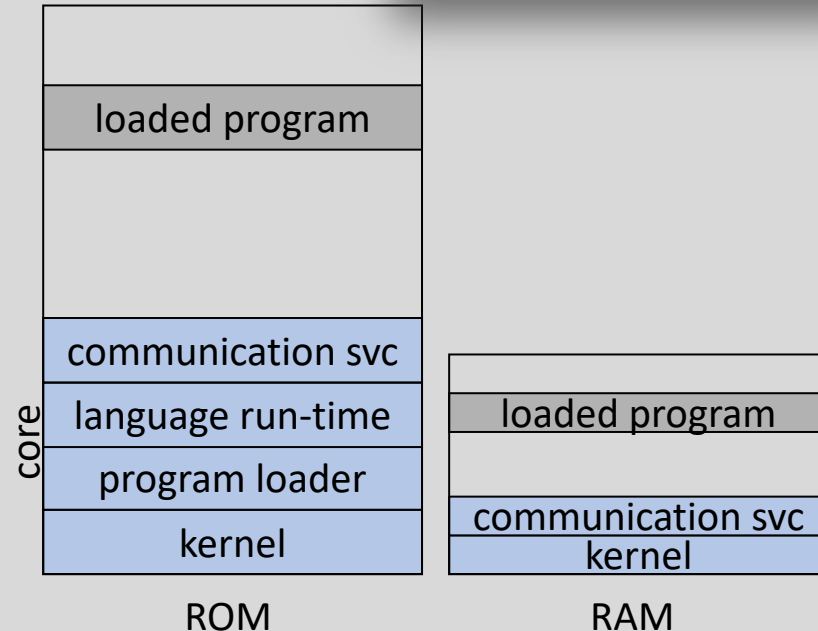
3. Contiki OS

1. Overview

- Event-driven kernel but with preemptive multitasking support
- Code size
 - Typical memory footprint: 40 KB ROM, 2 KB RAM
 - TCP/IP: < 5 KB code, < 2 KB RAM
- Ported to dozens of platforms incl. computers, video game consoles, sensors, pico-satellites
- System parts
 - Core
 - compiled into single binary image
 - Loaded programs
 - loaded by program loader
 - Program obtained via communication stack or EPROM

Contiki

The Open Source OS for the Internet of Things



3. Contiki OS

2. Event-Driven Multitasking Kernel

- Contiki System = Kernel + Libraries + set of processes
- Kernel dispatches events to running processes.
- Process is defined by
 - an event handler and
 - an optional polling handler function to poll hardware state.
- Processes are implemented as Protothreads.
- Process state is kept in its own memory.
Kernel keeps a pointer to the state.
- Processes share the same address space and do not run in different protection domains.
- Single shared stack for process execution
- Inter-process communication:
Processes post events to each other.

3. Contiki OS

2.1 Lightweight Event Scheduler

dispatches **events**
(posted by kernel or polling
mechanism) to processes (FIFO).

- Event types
 - Asynchronous:
 - Events are queued by kernel and dispatched to target process later (deferred procedure call).
 - Synchronous:
 - Immediate scheduling of target process and return of control after target has finished event processing (inter-process procedure call)

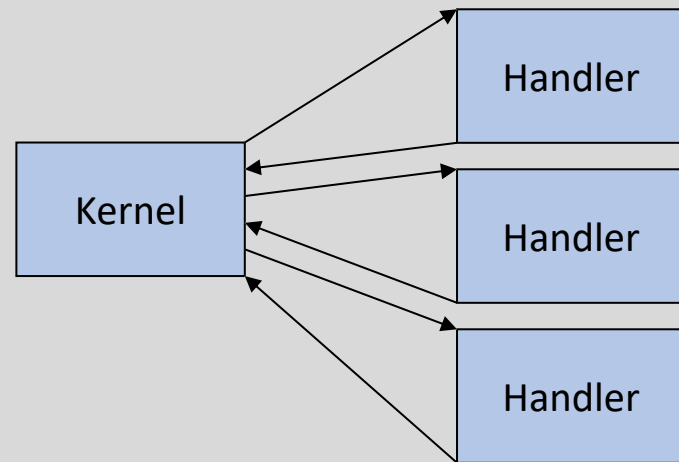
calls processes' **polling handlers**
periodically.

- Polls = high priority events scheduled between asynchronous events to check hardware device status
- For a scheduled poll, all processes implementing a poll handler are called according to their priority.

3. Contiki OS

2.2 Preemption and Interrupts

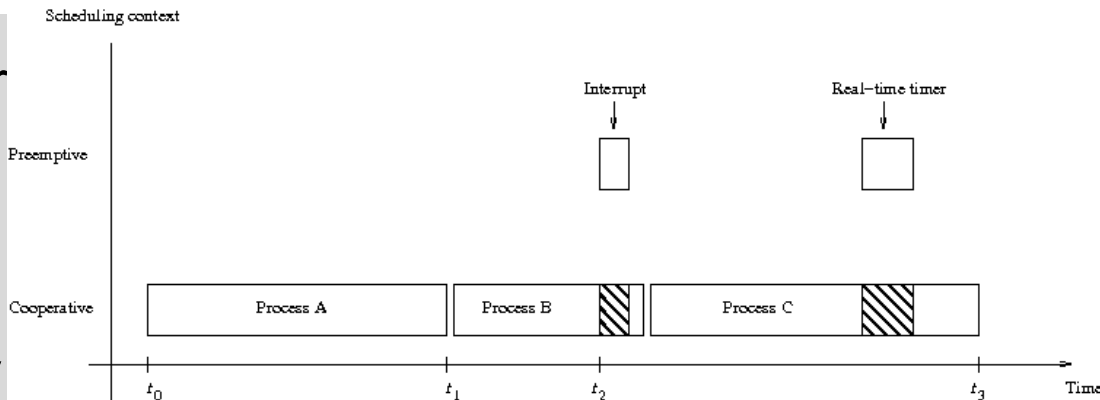
- Event handlers can not preempt each other.
- Event handlers run to completion.
- Preemption of events by interrupts only.
- Interrupts are never disabled.
- Interrupt handlers are not allowed to post events to avoid race conditions but can request immediate polling.



3. Contiki OS

2.3 Execution Contexts

- Code in Contiki runs in either of two execution contexts.
 - cooperative (Contiki processes)
 - preemptive (Interrupts, real-time timers)
- Preemptive code temporarily stops cooperative code.
- Cooperative code must run to completion before other cooperatively scheduled code can run.



3. Contiki OS

3. Protothreads

Motivation

- Event handlers cannot execute blocking wait due to run-to-completion semantics in case of a single shared stack.
- Disadvantage: State machines become complex and cannot be expressed by a single event handler
→ difficult to program and maintain

New concept of lightweight threads: **Protothreads**

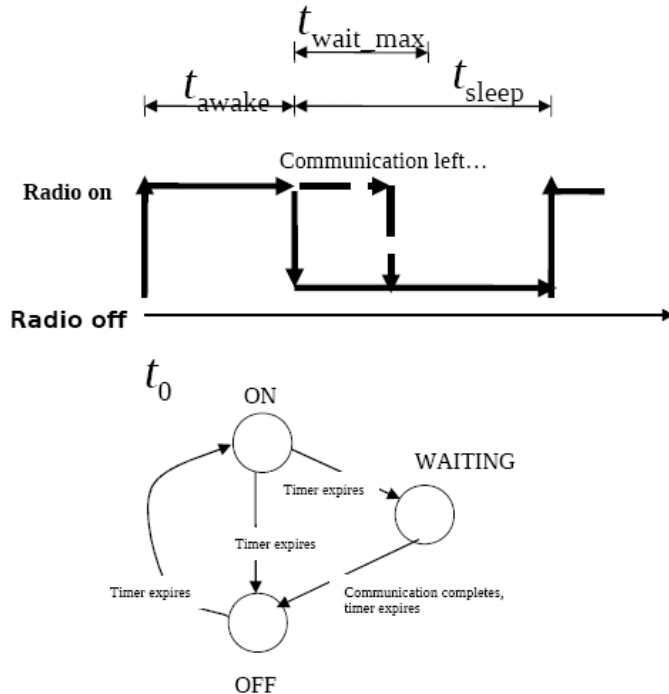
- = memory-efficient programming abstraction that shares features of both multi-threading and event-driven programming to achieve low memory overhead.
- combine events and threads (blocking event handlers)
- provide several conditional blocking wait statements: e.g., `PROCESS_WAIT_EVENT_UNTIL()` blocks until conditional statement becomes true
→ linear sequencing of statements in event-driven programs
- are stackless: All Protothreads share the same stack, which is rewound every time a Protothread blocks. Local variables are not preserved across blocking wait statements.
- are cooperatively scheduled. Protothreads must always explicitly yield control back to kernel.

3.1 Example: Simple MAC Protocol

- 24

3. Contiki OS

3.2 Simple MAC Protocol with State Machine



```
enum {ON, WAITING, OFF} state;

void eventhandler() {
    if(state == ON) {
        if(expired(timer)) {
            timer = t_sleep;
            if(!comm_complete()) {
                state = WAITING;
                wait_timer = t_wait_max;
            } else {
                radio_off();
                state = OFF;
            }
        }
    } else if(state == WAITING) {
        if(comm_complete() ||
           expired(wait_timer)) {
            state = OFF;
            radio_off();
        }
    } else if(state == OFF) {
        if(expired(timer)) {
            radio_on();
            state = ON;
            timer = t_awake;
        }
    }
}
```

3. Contiki OS

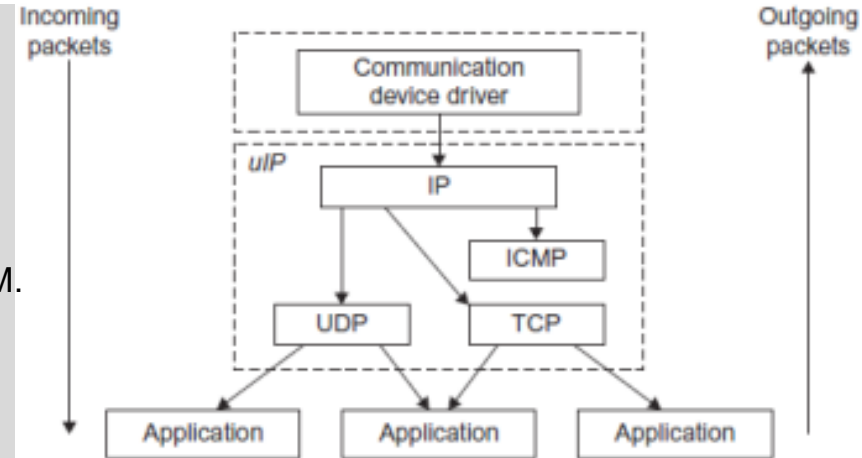
3.3 Simple MAC Protocol with Protothreads

```
static struct etimer timer;
static struct etimer wait_timer;
PROCESS_THREAD(radio_wake_process, ev, data)
{
    PROCESS_BEGIN();
    while(1) {
        NETSTACK_RADIO.on();
        etimer_set(&timer, T_AWAKE);
        PROCESS_WAIT_EVENT_UNTIL(etimer_expired(&timer));
        etimer_set(&timer, T_SLEEP);
        if(!NETSTACK_RADIO.channel_clear()) {
            etimer_set(&timer, T_WAIT_MAX);
            PROCESS_WAIT_EVENT_UNTIL(NETSTACK_RADIO.channel_clear() ||
                                     etimer_expired(&wait_timer));
        }
        NETSTACK_RADIO.off();
        PROCESS_WAIT_EVENT_UNTIL(etimer_expired(&timer));
    }
    PROCESS_END();
}
```

3. Contiki OS

4. uIP Stack

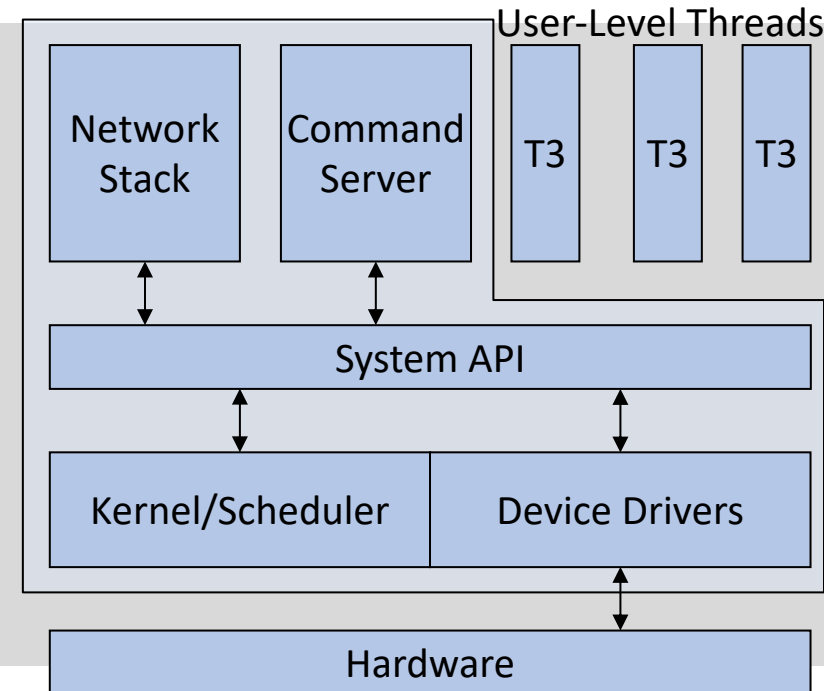
- small full TCP/IP stack supporting IPv4 and IPv6
- open source
- RFC-compliant to a large extent!
- intended for tiny microcontroller systems, where code size and RAM are severely constrained.
- uIP requires 4-5 KB code space and a few 100 bytes RAM.
- may eat up large share of memory of TelosB, only use when needed!
- limitations to minimize resource usage
 - Only one packet buffer for sending & receiving
 - Only one segment “in flight” at a time
 - congestion window = 1
 - no need for congestion control



4. MANTIS OS

1. Overview

- Classical layered multi-threaded operating system structure including
 - Multi-threading
 - Pre-emptive scheduling with time slicing
 - I/O synchronization via mutual exclusion
 - Network stack
 - Device drivers
- MANTIS code size
 - 14 KB flash memory
 - 500 bytes RAM



4. MANTIS OS

2. Memory and Thread Management

- Memory management
 - Allocation of space for global variables at compile time
 - Rest of memory managed as heap. Stack for each thread is allocated from heap.
- Thread management
 - Static, fixed size thread table
 - Semaphore support for applications
- Interrupts
 - Timer interrupts are handled by kernel.
 - Hardware interrupts are sent to device drivers.
 - No software interrupts
- Network stack options
 - Different layers (application, network, MAC, physical) can be implemented in different threads.
 - All layers can be implemented in a single thread.

4. MANTIS OS

3. Power Management and Dynamic Reprogramming

Power Management

- **sleep(period)**
enables power-save mode.
- If all threads sleep:
OS determines earliest deadline
and shuts down microcontroller.

Dynamic Reprogramming

- Reprogramming granularities
 - Re-flashing of entire OS
 - Reprogramming single threads
 - Changing of variables within a thread
- Implementation as system call library
 - Application (e.g. command server) may write a new code image through call to system call library into EEPROM.
 - Application calls **commit** operation, which writes a control block for the boot loader.
 - Control block causes installation of the new code on reset.

5. RIOT

1. Overview

- Microkernel architecture supporting multi-threading, including mutual exclusion
- Static memory allocation to guarantee run-times
- supports various C++ libraries
- TCP/IP network stack (6LoWPAN, RPL, IPv6, TCP, UDP, CoAP) and Named Data Networking
- supports several MCUs such as a 16-bit MSP430 or a 32-bit ARM7

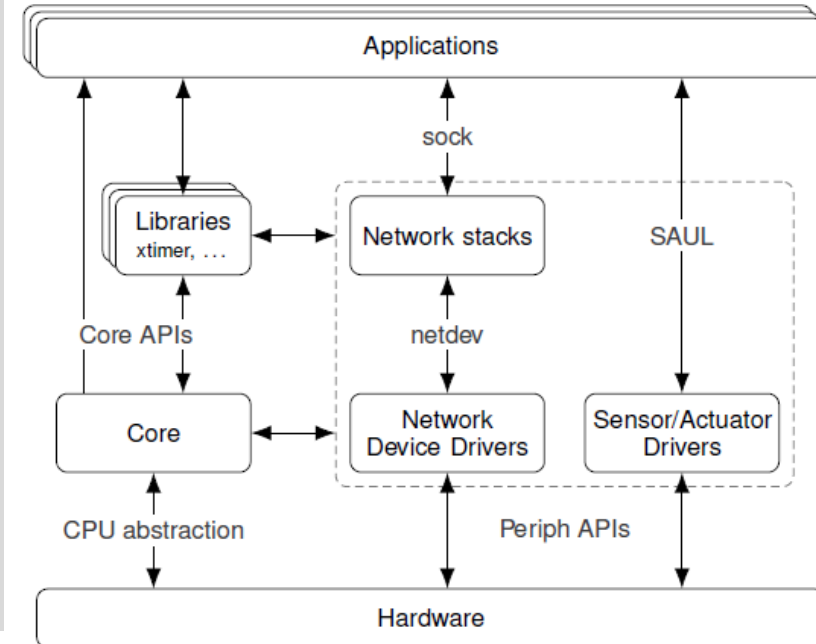
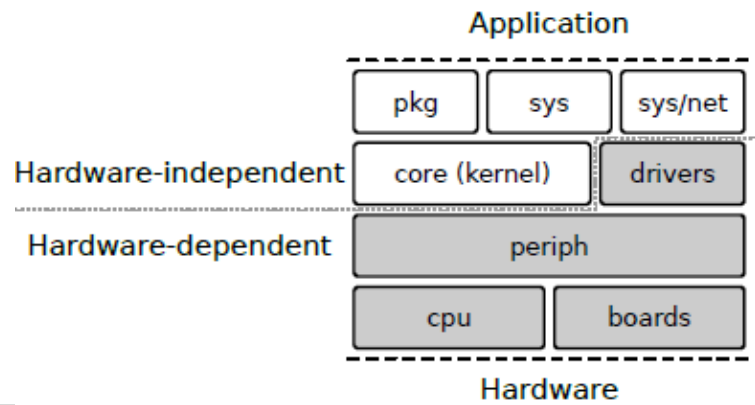
memory requirements:

RIOT Configuration	ROM	RAM
Basic RTOS	3.2 kB	2.8 kB
6LoWPAN-enabled	38.5 kB	10.0 kB
JavaScript-enabled	166.2 kB	29.1 kB
OTA-enabled	111 kB	17.5 kB

5. RIOT

2. Architecture

- core: OS kernel
- Hardware abstraction
 - i. CPU: microcontroller functionality
 - ii. boards: configures CPU and drivers
 - iii. drivers: HW independent device drivers for sensors, actuators, network transceivers, storage components with high-level APIs
 - iv. periph: unified access to microcontroller peripherals
- sys: system libraries
- pkg: 3rd party components



5. RIOT

3. Kernel

Multi-threading

- overhead by thread control block (36 bytes), stack, CPU context < 128 bytes
- Light-weight inter-process communication using semaphores, mutex, messaging

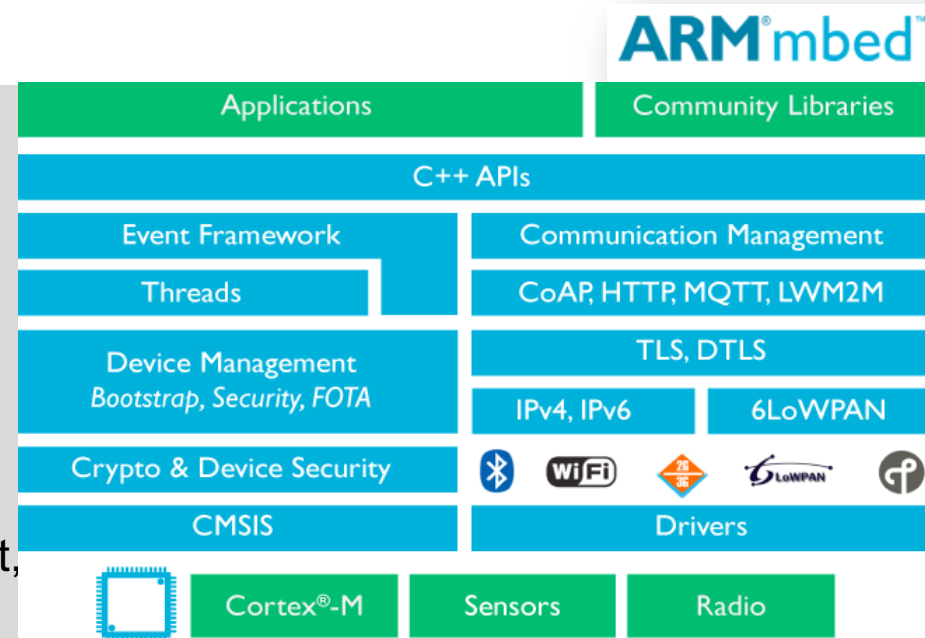
Tickless O(1) Priority Scheduler

- works without periodic events.
- Whenever there are no pending tasks, RIOT will switch to idle thread.
- The only function of the idle thread is to determine the deepest possible sleep mode, depending on the peripheral devices in use.
- Only interrupts (external or kernel-generated) wake up the system from idle state.

6. Other IoT Operating Systems

1. mbed

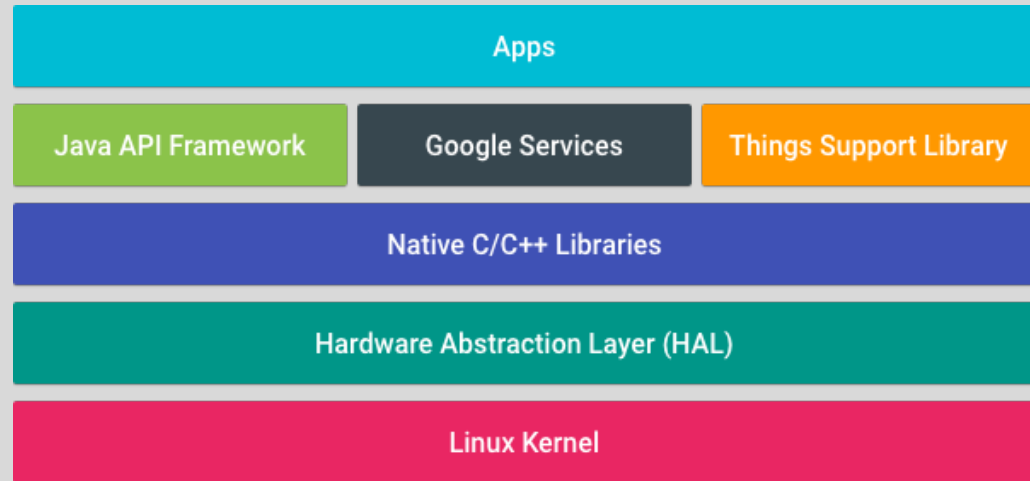
- Chip designers are developing “semi-open” OS platforms, e.g., ARM Microsystems has created mbed for 32-bit ARM Cortex-M microcontrollers.
- C/C++ based application development platform, online compilation in the cloud
- Tools for creating microcontroller firmware
- Core libraries providing peripheral drivers, networking, runtime environment, build tools, test and debug scripts.
- Components database provides driver libraries for components and services.



6. Other IoT Operating Systems

2. Android Things

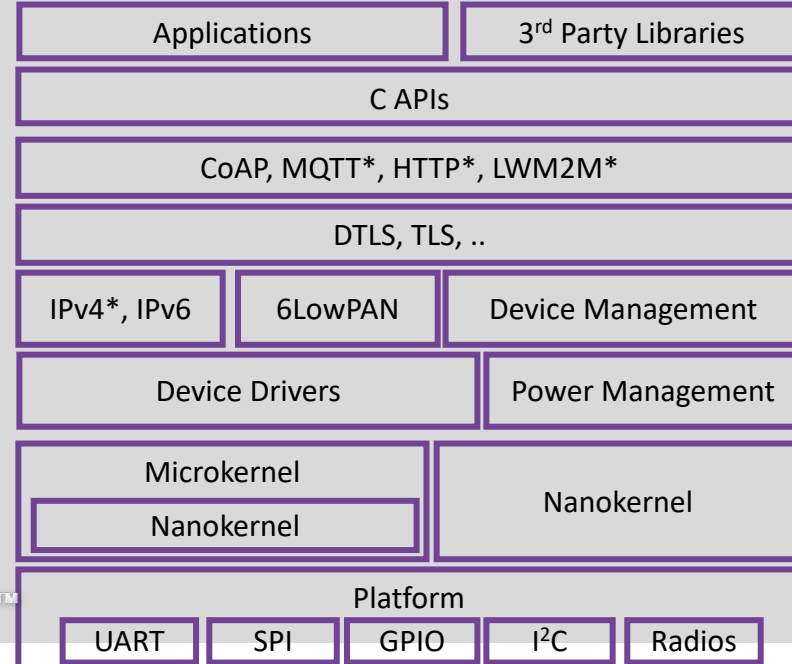
- Android based embedded OS designed to work with 32-64 MB RAM
- Supporting BLE and WiFi
- Software Development Kit for application development, similar to developing Android applications



6. Other IoT Operating Systems

3. Zephyr

- Linux Foundation hosted collaboration project.
- producing open source code
- small, scalable, library-based real-time operating system, optimized for resource constrained devices across multiple architectures
- supports hardware > 8 KB RAM
- 1 single executable executed in 1 single address space
- Multi-threading support, inter-thread communication
- Kernel mode only, no user-space, no dynamic runtimes
 - Nanokernel: Limited functionality targeting small memory footprint (below 10kB RAM)
 - Microkernel (superset of nanokernel): with additional functionality and features
- Memory and resources are typically statically allocated.



6. Other IoT Operating Systems

4. FreeRTOS

- Micro kernel OS
 - Reduced kernel size (3 C files)
 - All other functions are provided by servers running on user level.
- Preemptive priority-based round-robin scheduling
- Multi-threading support
- Idle task puts microcontroller into low power mode
- Event-based memory allocation
- Third-party network support, e.g. uIP



7. Operating System Power Management

Dynamic Power Management

- CPU-centric
 - Low-Energy Earliest Deadline First Scheduling
- I/O-centric
 - Low-Energy DEvice Scheduling

7. Operating System Power Management

1. Low-Energy Earliest Deadline First

```

 $t_c$ : current time;
 $S_h > S_{l1} > S_{l2} > \dots S_{lm}$ : Available processor speeds
schedulable = 1
1. if ready_list  $\neq$  NULL
2.   Sort task deadlines in ascending order;
3.   Select task  $\tau_i$  with earliest deadline;
4.   for  $S = S_{lm}$  to  $S_h$ 
5.     if  $t_c + \frac{l_i}{S} \leq d_i$  then
6.        $t = t_c + \frac{l_i}{S}$ 
7.       for each task  $\tau_u$  that has not completed execution
8.         if  $t + \frac{l_u}{S_h} \leq d_u$  then
9.            $t = t + \frac{l_u}{S_h}$ 
10.        else
11.          schedulable = 0
12.          break
13.        endfor
14.      if schedulable = 1 then
15.        Schedule  $\tau_i$  at  $S$ 
16.        break
17.      endif
18.    endfor

```

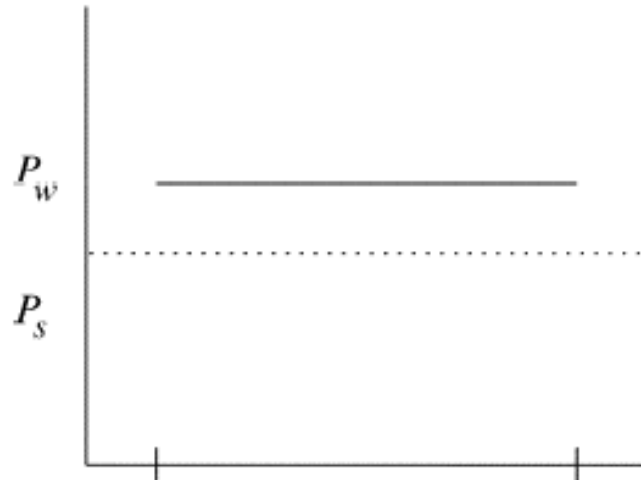
d_i : deadline for task τ_i

l_i : length of task τ_i (instruction cycles)

7. Operating System Power Management

2.1 Low-Energy DEvice Scheduling

Device stays powered up for entire interval t_{be}



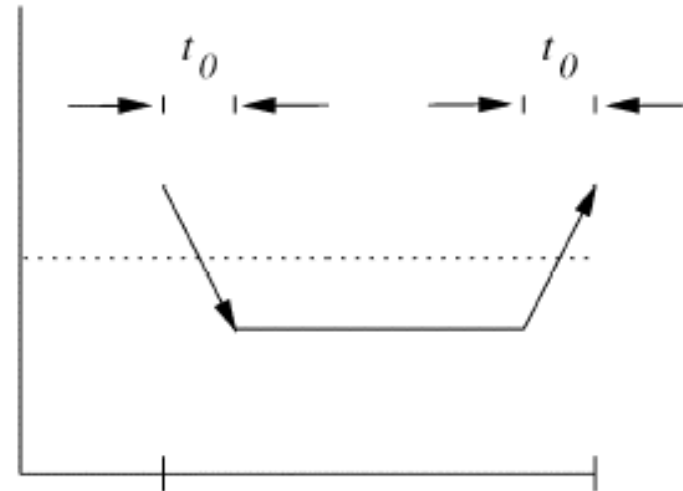
Set of tasks with

- start times s_i
- completion times c_i

$t_{0,j}$: state transition time for device k_j ,

L_i : devices needed by task i

Device performs two transitions during interval t_{be}



6. Operating System Power Management

2.2 LEDES

Algorithm LEDES(k_j, τ_i, τ_{i+1})

curr: current scheduling instant;

```

1.  if curr =  $s_i$            begin of task i
2.  if  $k_j$  is powered-up
3.  if  $k_j \notin L_i \cup L_{i+1}$  } shut down device if not used by next 2 tasks (i, i+1)
4.  shutdown  $k_j$ 
5.  if  $k_j \in L_{i+1}$  and
6.  if  $s_{i+1} - (s_i + c_i) \geq t_{0,j}$  } shut down device for task i+1 if there is enough
7.  shutdown  $k_j$                     time for starting device between tasks i and i+1
8.  if ( $k_j$  is shut down)
9.  if  $k_j \in L_{i+1}$  and  $s_{i+1} - (s_i + c_i) < t_{0,j}$  } start device for task i+1 if there is not enough time
10. wakeup  $k_j$                                 for starting device between tasks i and i+1
11. if curr =  $s_i + c_i$            end of task i
12. if  $k_j$  is powered-up
13. if  $k_j \notin L_{i+1}$  and  $s_{i+1} - \text{curr} \geq t_{0,j}$  }
14. shutdown  $k_j$ 
15. else if ( $k_j \in L_{i+1}$ )
16. wakeup  $k_j$ 

```



shut down device if not needed for task i+1 and if there remains enough time to shut down before task i+1 starts, because it may be necessary to wake the device up again then.

Thanks

for Your Attention

Prof. Dr. Torsten Braun, Institut für Informatik

Bern, 08.03.2021

u^b

^b
**UNIVERSITÄT
BERN**

