Computer Vision

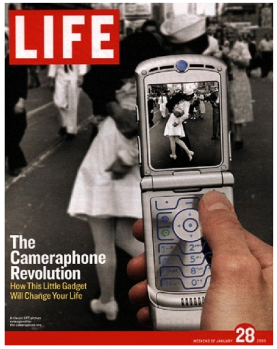

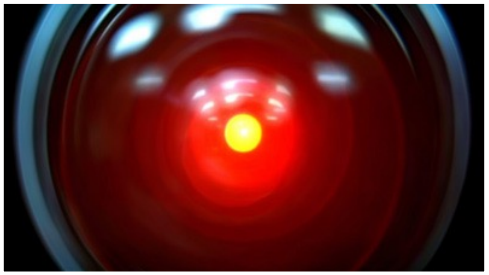

Basic Info

- Instructor: Paolo Favaro (paolo.favaro@inf.unibe.ch)
- Teaching Assistants:
  Abdelhak Lemkhenter (abdelhak.lemkhenter@inf.unibe.ch)
  Llukman Cerkezi (llukman.cerkezi@inf.unibe.ch)
  Alp Eren Sari (alp.sari@inf.unibe.ch)
- Course webpage: ILIAS
- Course material acknowledgements:
  - Svetlana Lazebnik – University of Illinois
  - Steve Seitz – University of Washington
  - Jan Koenderink – University of Leuven
  - Kristen Grauman – University of Texas at Austin
  - Andrea Vedaldi – University of Oxford
  - Srinivas Narasimhan – Carnegie Mellon University

To Do List

- On behalf of the Branch Committee for Computer Science, I would like to inform you about the **student's ToDo list for every semester**. The details can be found on our web site at http://mcs.unibnf.ch/organization.

- **JMCS students** please pay attention to these dates

1. Each student must **register for the teaching units** that they want to take during that semester until the registration deadline (see below) in Academia.

2. Each student must **register for the exams** until the registration deadline (see below) in Academia.

3. Each student must **register for the BeNeFri network** if he/she follows a teaching unit at a partner university until

- Autumn semester: September 30th
- Spring semester: February 28th

The procedure is dependent on the student's situation

- Bern: http://www.selfservice.studis.unibe.ch
- Neuchâtel:
  - JMCS student: nothing to do
  - Hosted JMCS student: https://am-institute.ch/benefri/assets/files/BeNeFri-inscript_automne%202015.pdf
- Fribourg: http://www.unifr.ch/mydata

Travel expenses are reimbursed by the student's home university and must be claimed until

- Autumn semester: January 31st
- Spring semester: June 30th

Use therefore the travel regulation (BeNeFri forms).

Dates and deadlines for the academic year 2020/2021

Teaching units

	Autumn semester 2020	Spring semester 2021
Registrations	31.08.2020 – 09.10.2020	08.02.2021 – 19.03.2021
Teaching units	14.09.2020 – 18.12.2020	22.02.2021 – 04.06.2021

Exams

	Autumn dates 2020 (autumn semester 2020)	Spring dates 2021 (spring semester 2021)	Summer dates 2021 (repetition 2021)**

Hosted JMCS students

- **Hosted JMCS students**, e.g.: bachelor students in Computer Science (major, minor, optional studies) or master students in Mathematics or Information Management have exactly to do the same tasks as JMCS students. Additionally, **every semester** they have to:
- Complete a request form for Academia access (http://mcs.unibnf.ch/node/535) at the beginning of every semester

Hosted JMCS students

Hosted JMCS students are all students other than JMCS students who follow a teaching unit of the Joint Master in Computer Science, e.g.:

- bachelor students in Computer Science (major, minor, optional studies)
- master students in Mathematics or Information Management

They have exactly to do the same three tasks as mentioned above for JMCS students and they must additionally do **every semester** the following task:

4. Each student must **complete a Request for Academia Access form** at the beginning of every semester until
   - Autumn semester: September 30th
   - Spring semester: February 28th

The student will be informed when he/she is able to login into Academia.

Textbooks

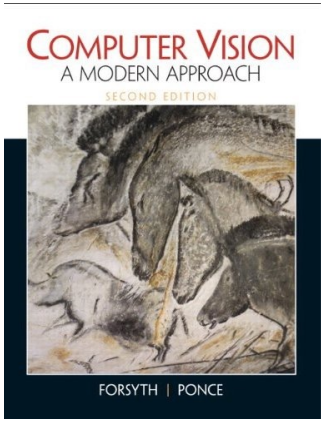

- Forsyth & Ponce, *Computer Vision: A Modern Approach*

- Kristen Grauman and Bastian Leibe, *Visual Object Recognition (pdf available online)*
- Christopher Bishop, *Pattern Recognition and Machine Learning*
- Yi Ma, Stefano Soatto, Jana Kosecka, S. Shankar Sastry, *An invitation to 3-D vision: From images to geometric models*
- Richard Szeliski, *Computer Vision: Algorithms and Applications*
- Simon J.D. Prince, Computer Vision: Models, Learning, and Inference (available online: http://computervisionmodels.com)

Course requirements

- 2 assignments
  - Deadlines specified in the next slide
  - Each assignment has a maximum score of 100 and a pass is 60
  - To register for the exam one needs at least a pass on each assignment
  - Each assignment to be made available in ILIAS
  - Assignments will require use of PYTHON (tutorial material will be provided)
- Tutorials
  - Weekly
  - Aim is exam preparation (exercises will be similar to exam problems)
- Exam
  - There will be a written final exam (duration 120 mins)
- Final Mark
  - 70% Exam and 30% Assignments

Assignments

- First assignment: Image deblurring
  - Available on October 13
  - First half is due on October 27 and the second half on November 10
- Second assignment: 3D reconstruction
  - Available on November 10
  - First half is due on November 24 and the second half on December 8
- Late submission policy
  - Solutions to the first half will be made available in ILIAS a day after the submission
  - Submission delays: No delay is allowed for the submission of the first half (as the solutions will be made available); for the second half we allow a delay of up to 1 week with a linear penalty (-0% points on the deadline until -100% on the 7^{th} day after the deadline)

Image deblurring

1 out of 4 blur cases to be solved with 3 optimization methods

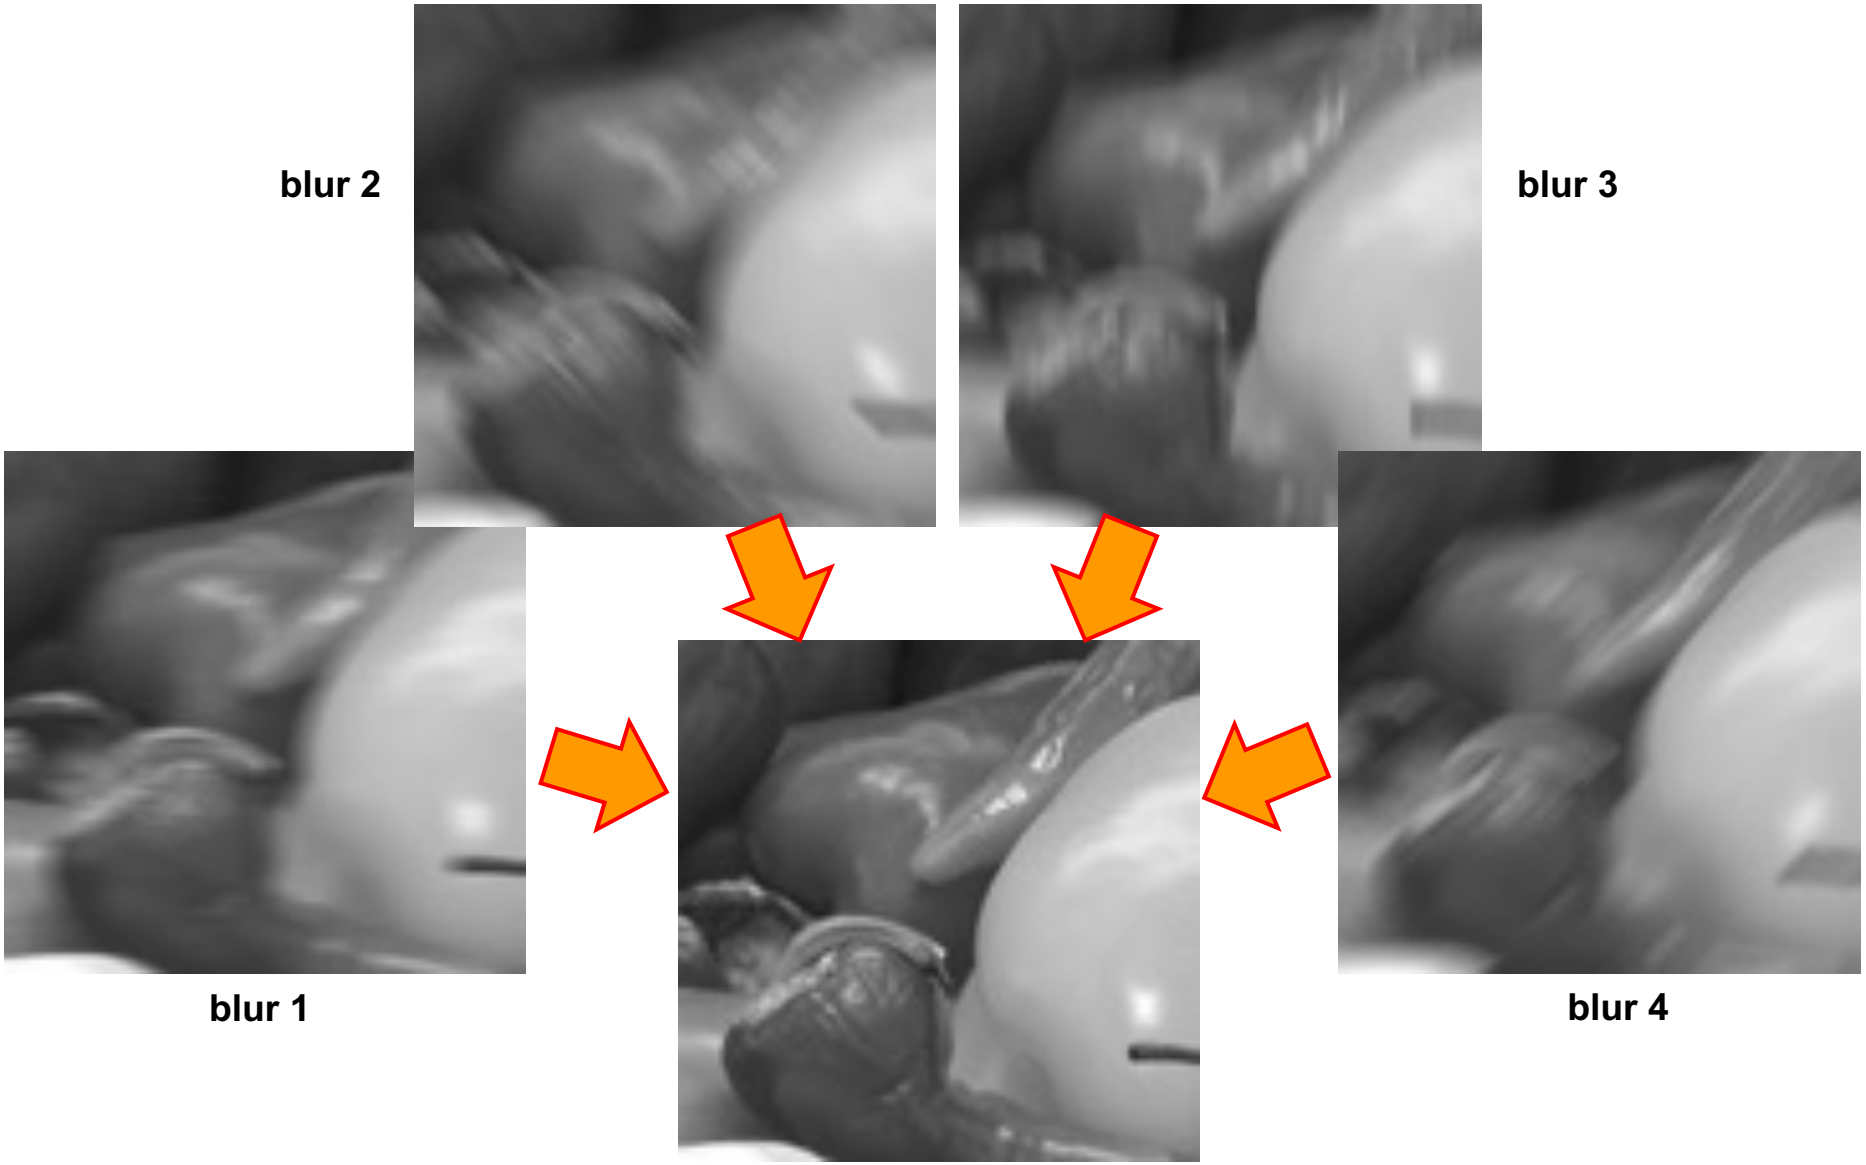

3D Reconstruction

1. Given high-res pictures of the same scene from different viewpoints

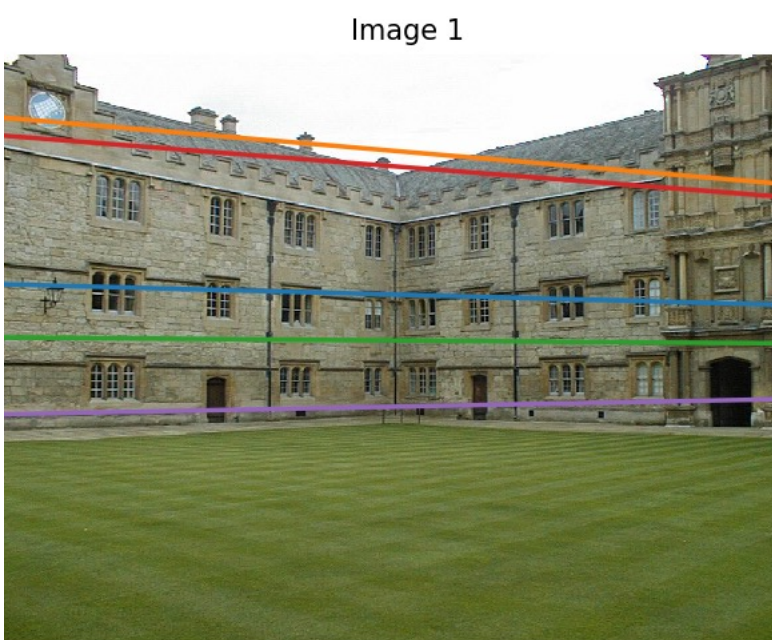

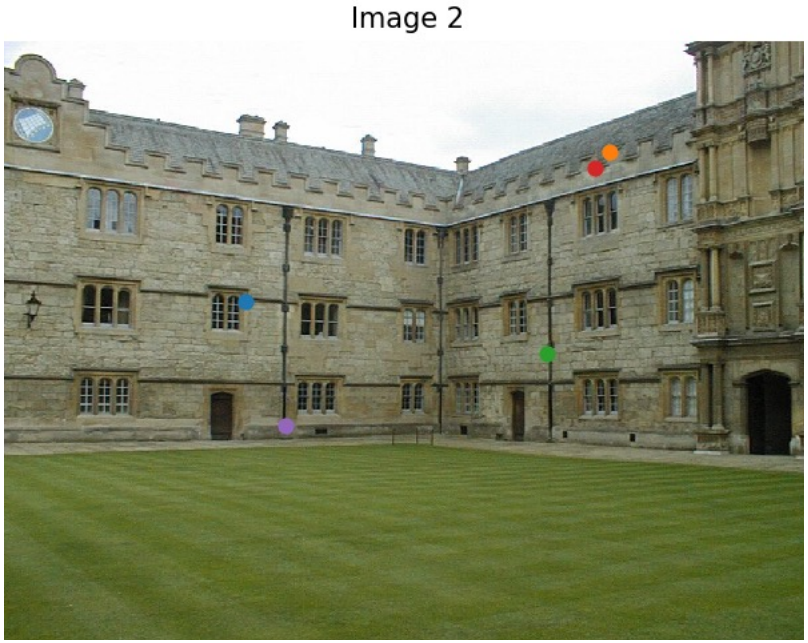

3D Reconstruction

2. Select points and apply computer vision methods (epipolar geometry) to find correspondences in the other image

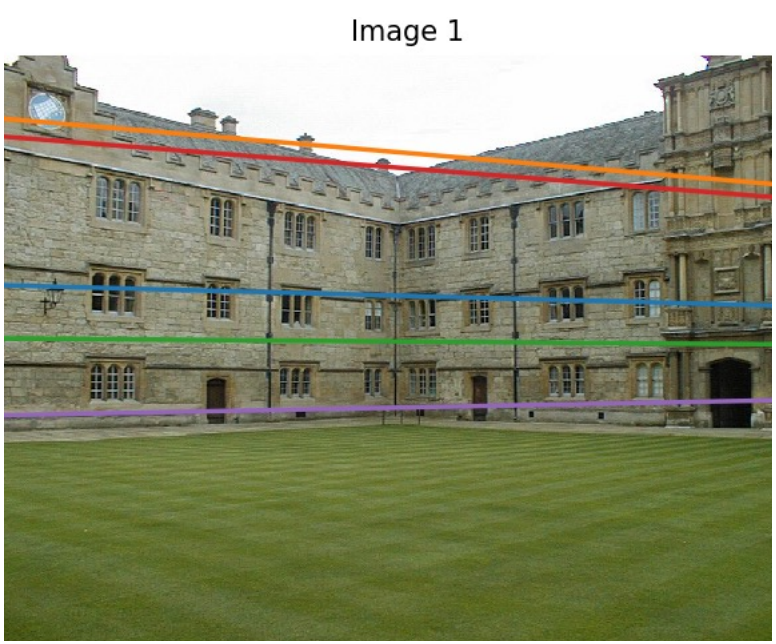

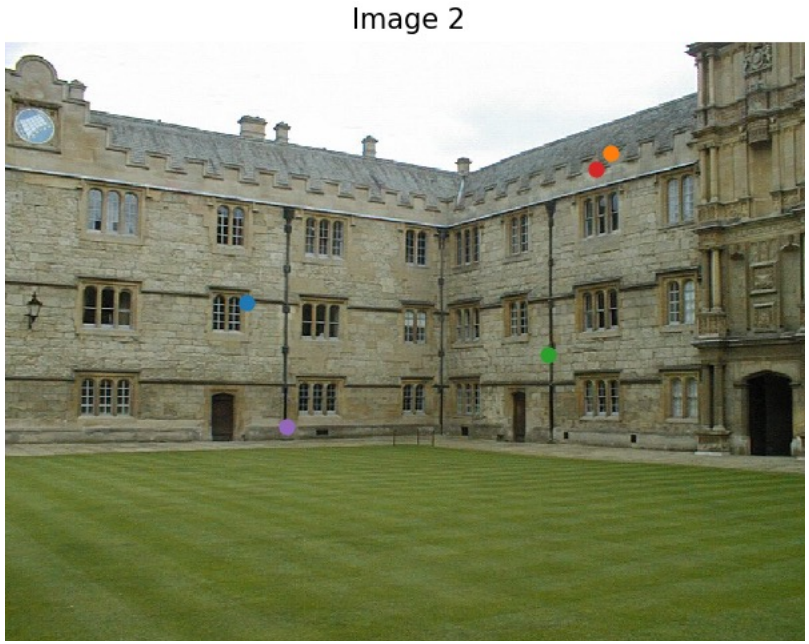

3D Reconstruction

3. Find reliable point correspondences from one image to the other

2D points from file

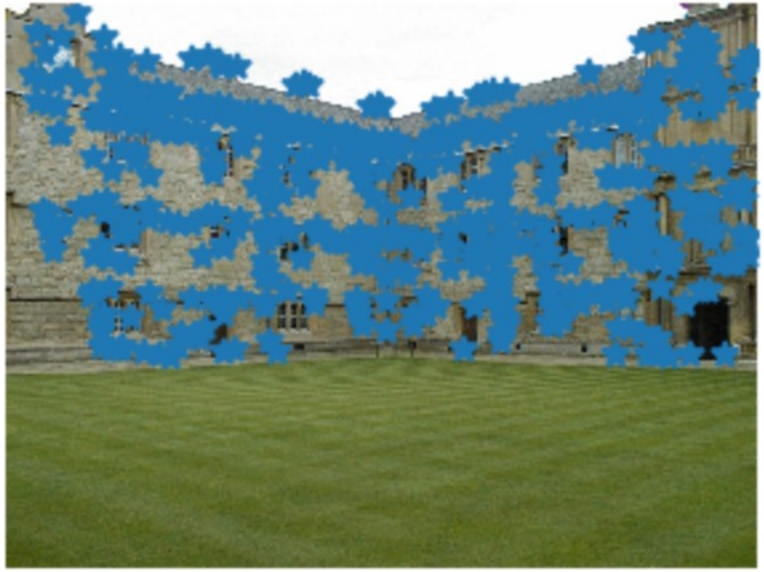

3D points projected into view 1

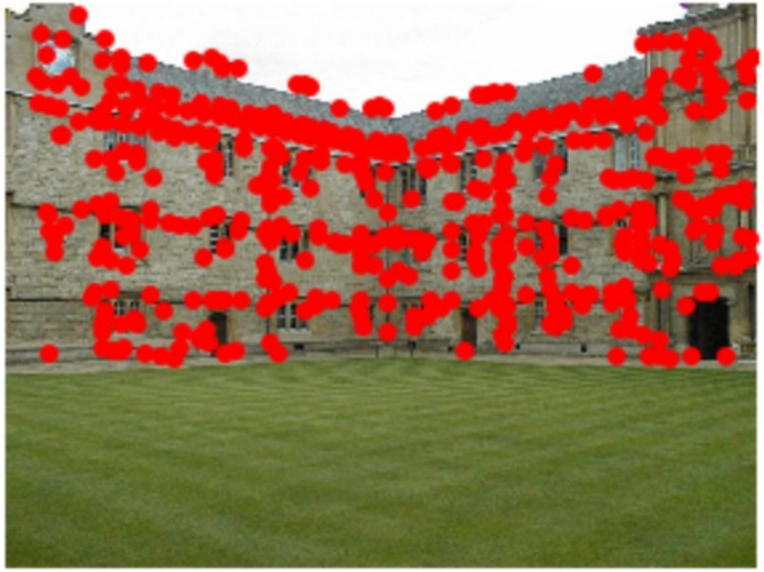

3D Reconstruction

4. Compute the 3D reconstruction of the detected points (through the 8 point algorithm and the essential matrix)

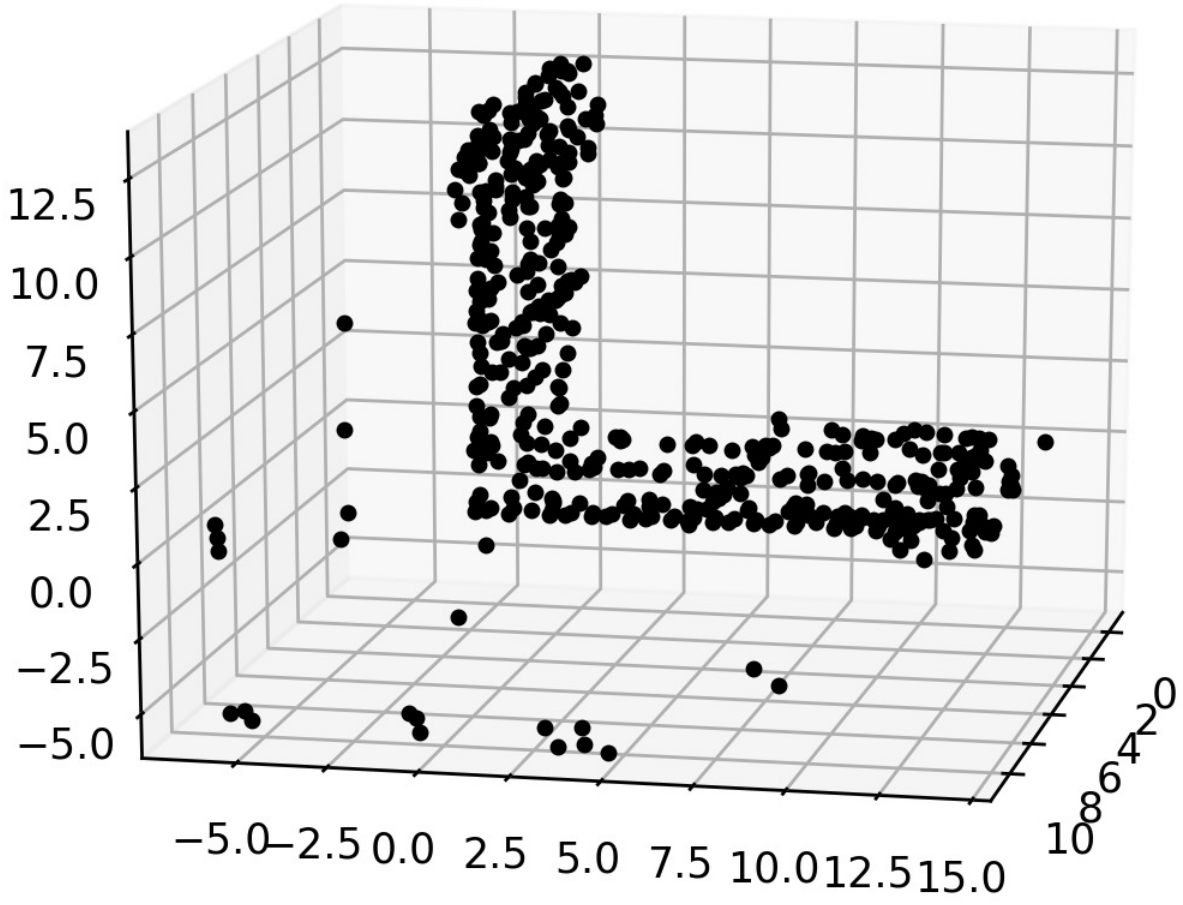

Academic integrity policy

- Feel free to discuss assignments with each other, but **report writing and coding must be done individually and without plagiarism**

- Feel free to incorporate code or tips you find on the Web, provided this does not make the assignment trivial and you **explicitly acknowledge your sources**

Computer Vision

The goal of computer vision

- To extract “meaning” from pixels

0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0

The goal of computer vision

- To extract “meaning” from pixels

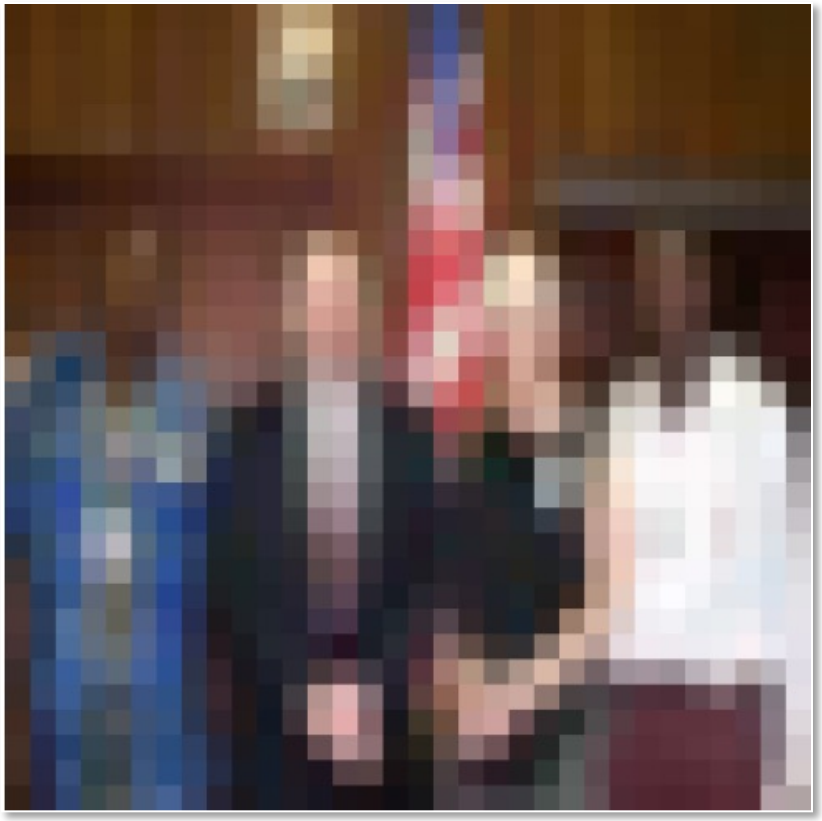

What kind of information can be extracted from an image?

Semantic information

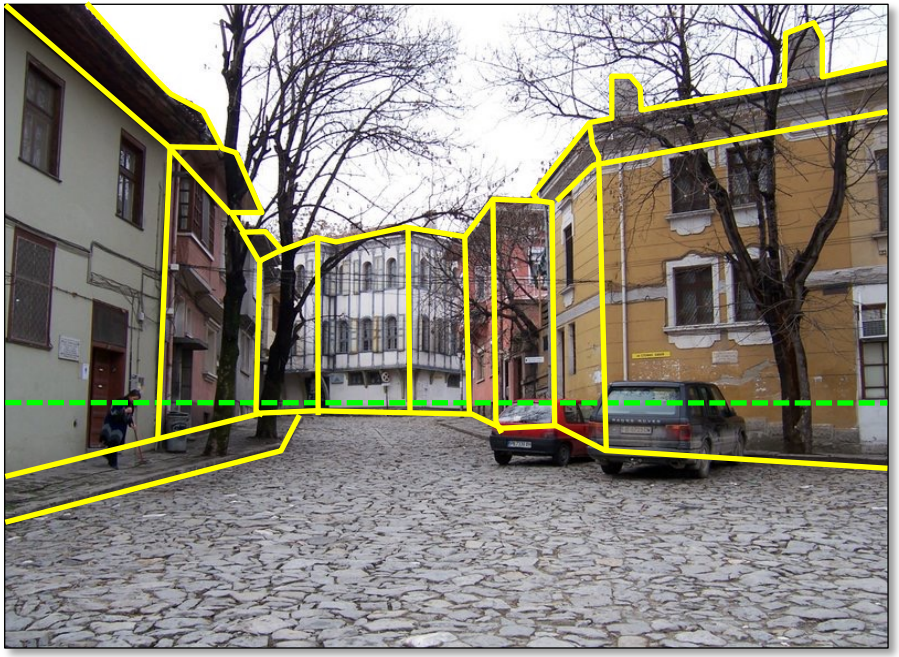

Geometric information

Why study computer vision?

- Vision is useful
- Vision is interesting
- Beware: Human vision is easy, but computer vision is difficult
  - Half of primate cerebral cortex is devoted to visual processing
  - Achieving human-level visual perception is probably "AI-complete"

Successes of computer vision to date

Optical character recognition (OCR)

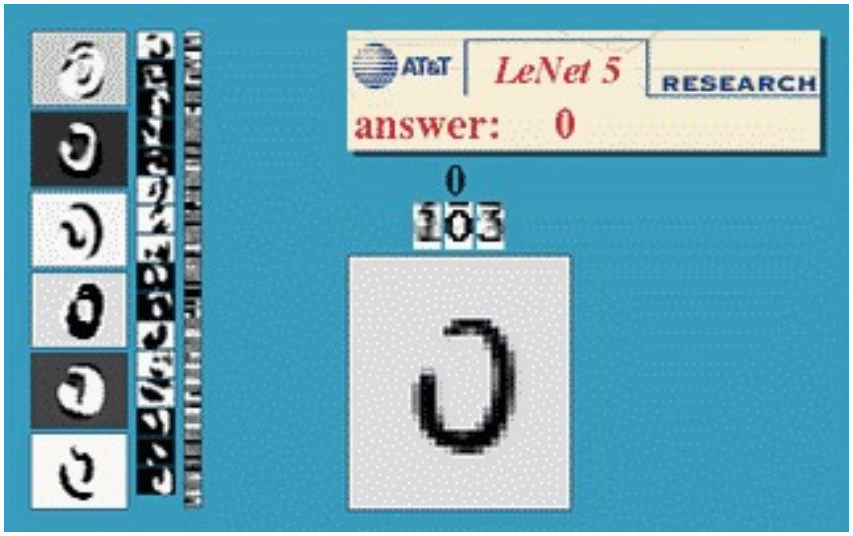

Digit recognition
yann.lecun.com

License plate readers
http://en.wikipedia.org/wiki/Automatic_number_plate_recognition

Automatic check processing

Sudoku grabber

http://sudokugrab.blogspot.com/

Source: S. Seitz, N. Snavely

Biometrics

Fingerprint scanners on many new laptops, other devices

Face recognition systems now beginning to appear more widely
http://www.sensiblevision.com/

Source: S. Seitz

Mobile visual search: Google Goggles

Google Goggles in Action

Click the icons below to see the different ways Google Goggles can be used.

Google goggles labs

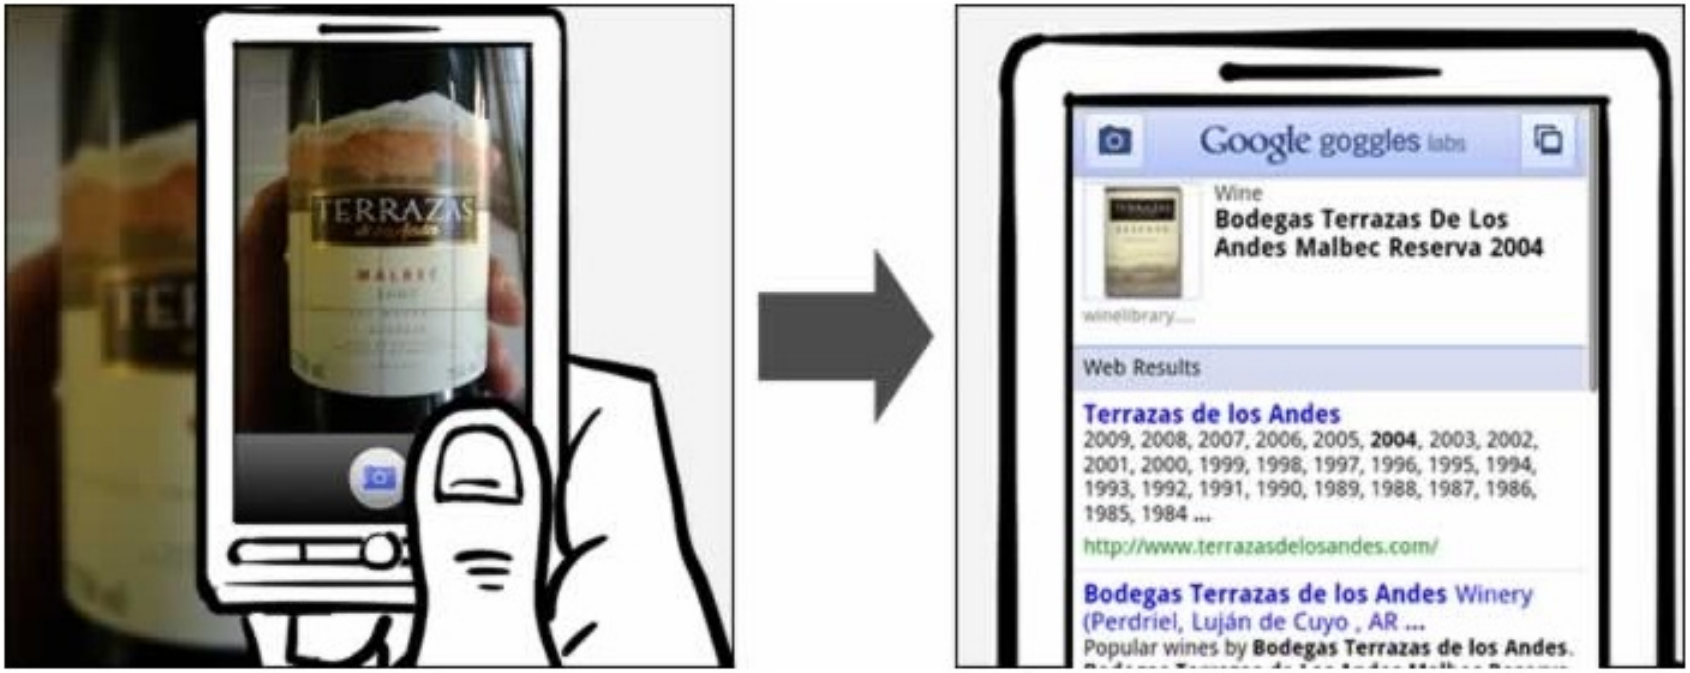

Automotive safety

Mobileye: Vision systems in high-end BMW, GM, Volvo models

- Pedestrian collision warning
- Forward collision warning
- Lane departure warning
- Headway monitoring and warning

Source: A. Shashua, S. Seitz

Google/Tesla self-driving cars

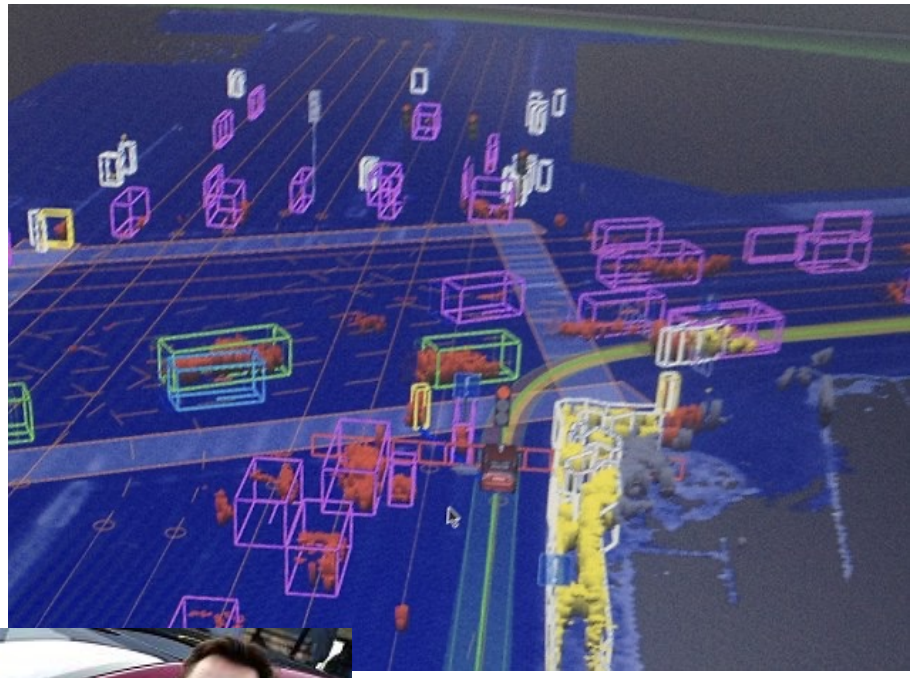

Vision-based interaction: Xbox Kinect

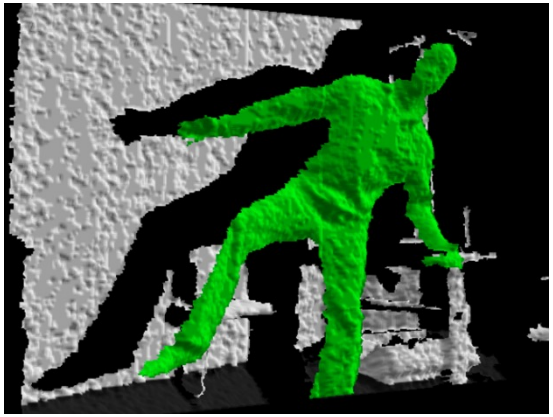

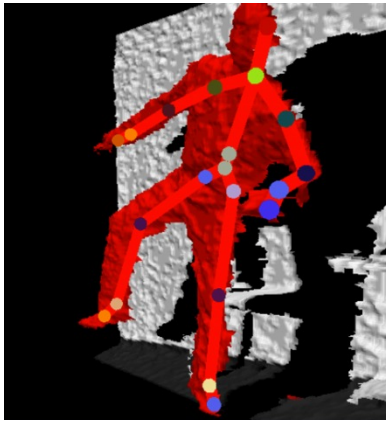

3D Reconstruction: Multi-View Stereo

YouTube Video

Special effects: shape and motion capture

Source: S. Seitz

Vision for robotics, space exploration

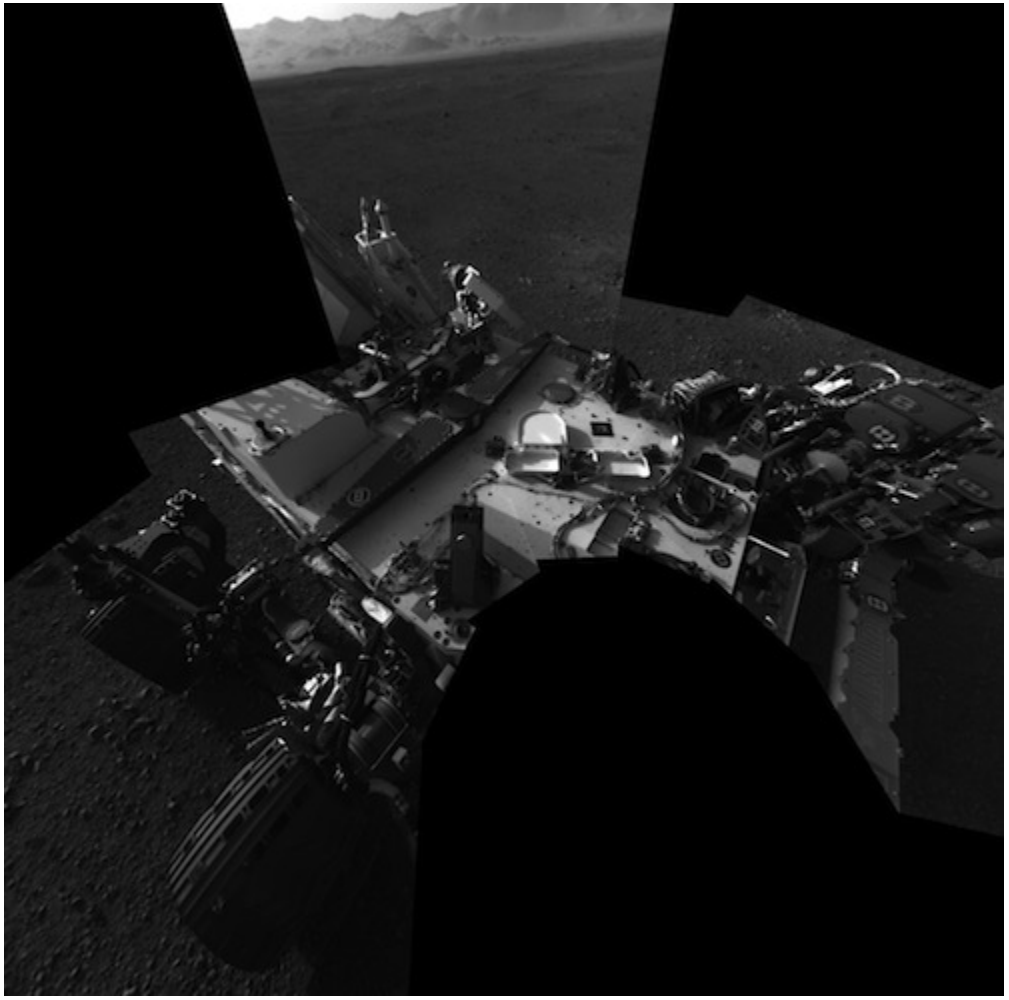

NASA'S Curiosity Rover has a system consisting of 17 cameras

Why is computer vision difficult?

Challenges: viewpoint variation

Michelangelo 1475-1564

slide credit: Fei-Fei, Fergus & Torralba

Challenges: illumination

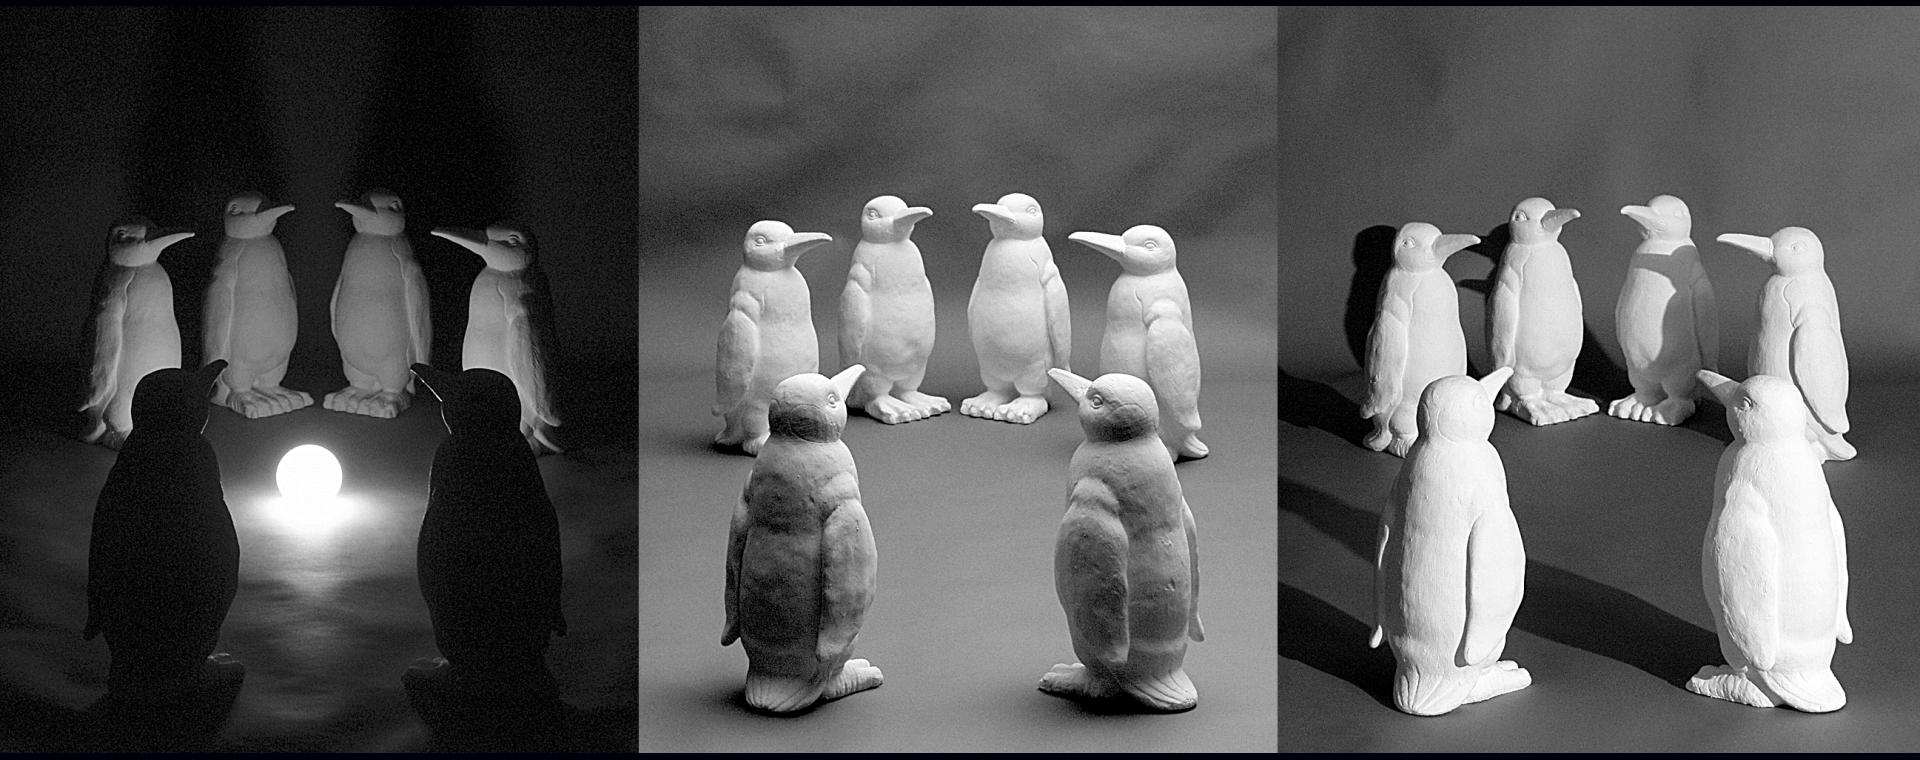

image credit: J. Koenderink

Challenges: scale

slide credit: Fei-Fei, Fergus & Torralba

Challenges: deformation

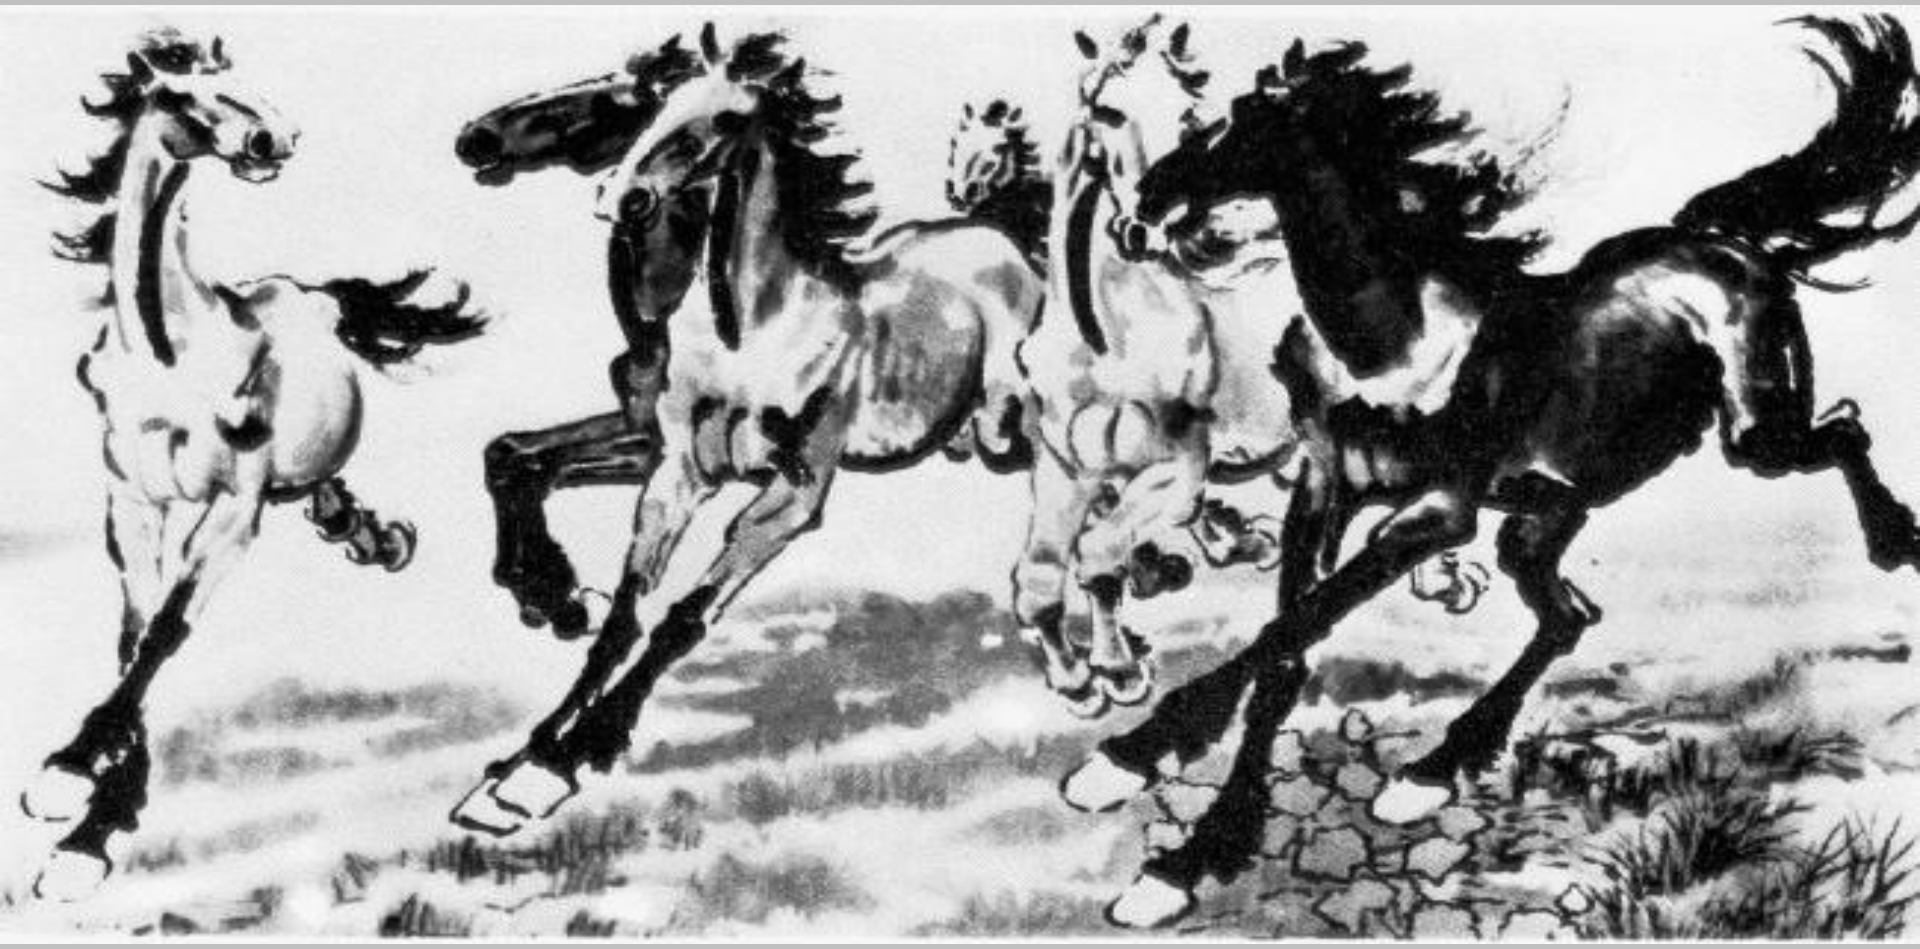

Xu, Beihong 1943

slide credit: Fei-Fei, Fergus & Torralba

Challenges: occlusion, clutter

Image source: National Geographic

Challenges: Motion

Challenges: object intra-class variation

slide credit: Fei-Fei, Fergus & Torralba

Challenges: local ambiguity

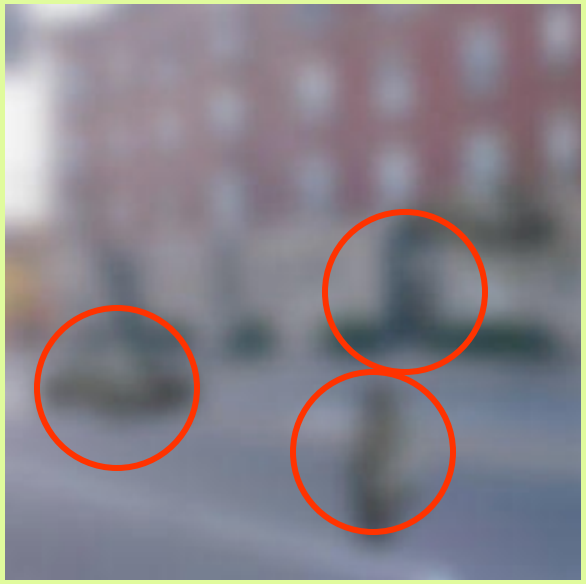

slide credit: Fei-Fei, Fergus & Torralba

Challenges: local ambiguity

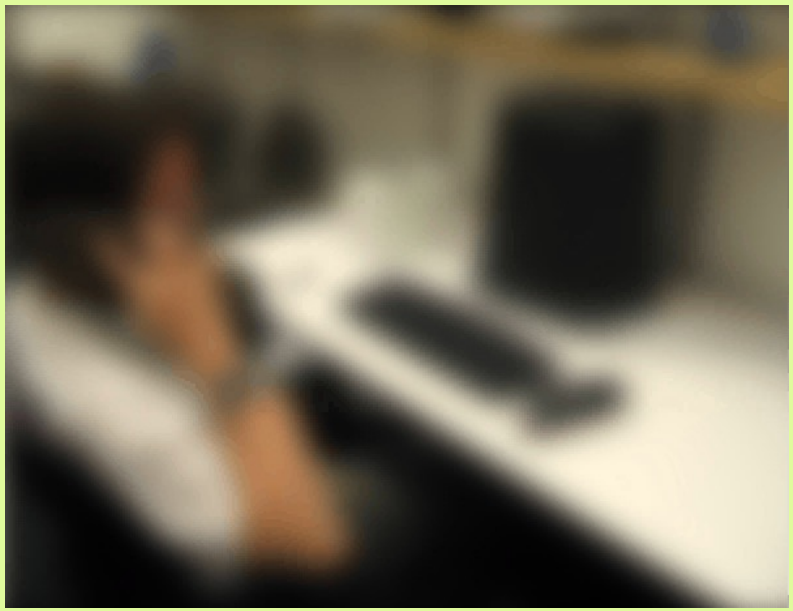

Source: Rob Fergus and Antonio Torralba

Challenges: local ambiguity

Source: Rob Fergus and Antonio Torralba

Challenges: Inherent ambiguity

- Many different 3D scenes could have given rise to a particular 2D picture

Challenges or opportunities?

- Images are confusing, but they also reveal the structure of the world through numerous cues
- Our job is to interpret the cues!

Depth cues: Linear perspective

Depth cues: Aerial perspective

Depth ordering cues: Occlusion

Shape cues: Texture gradient

Shape and lighting cues: Shading

Position and lighting cues: Cast shadows

Grouping cues: Similarity (color, texture, proximity)

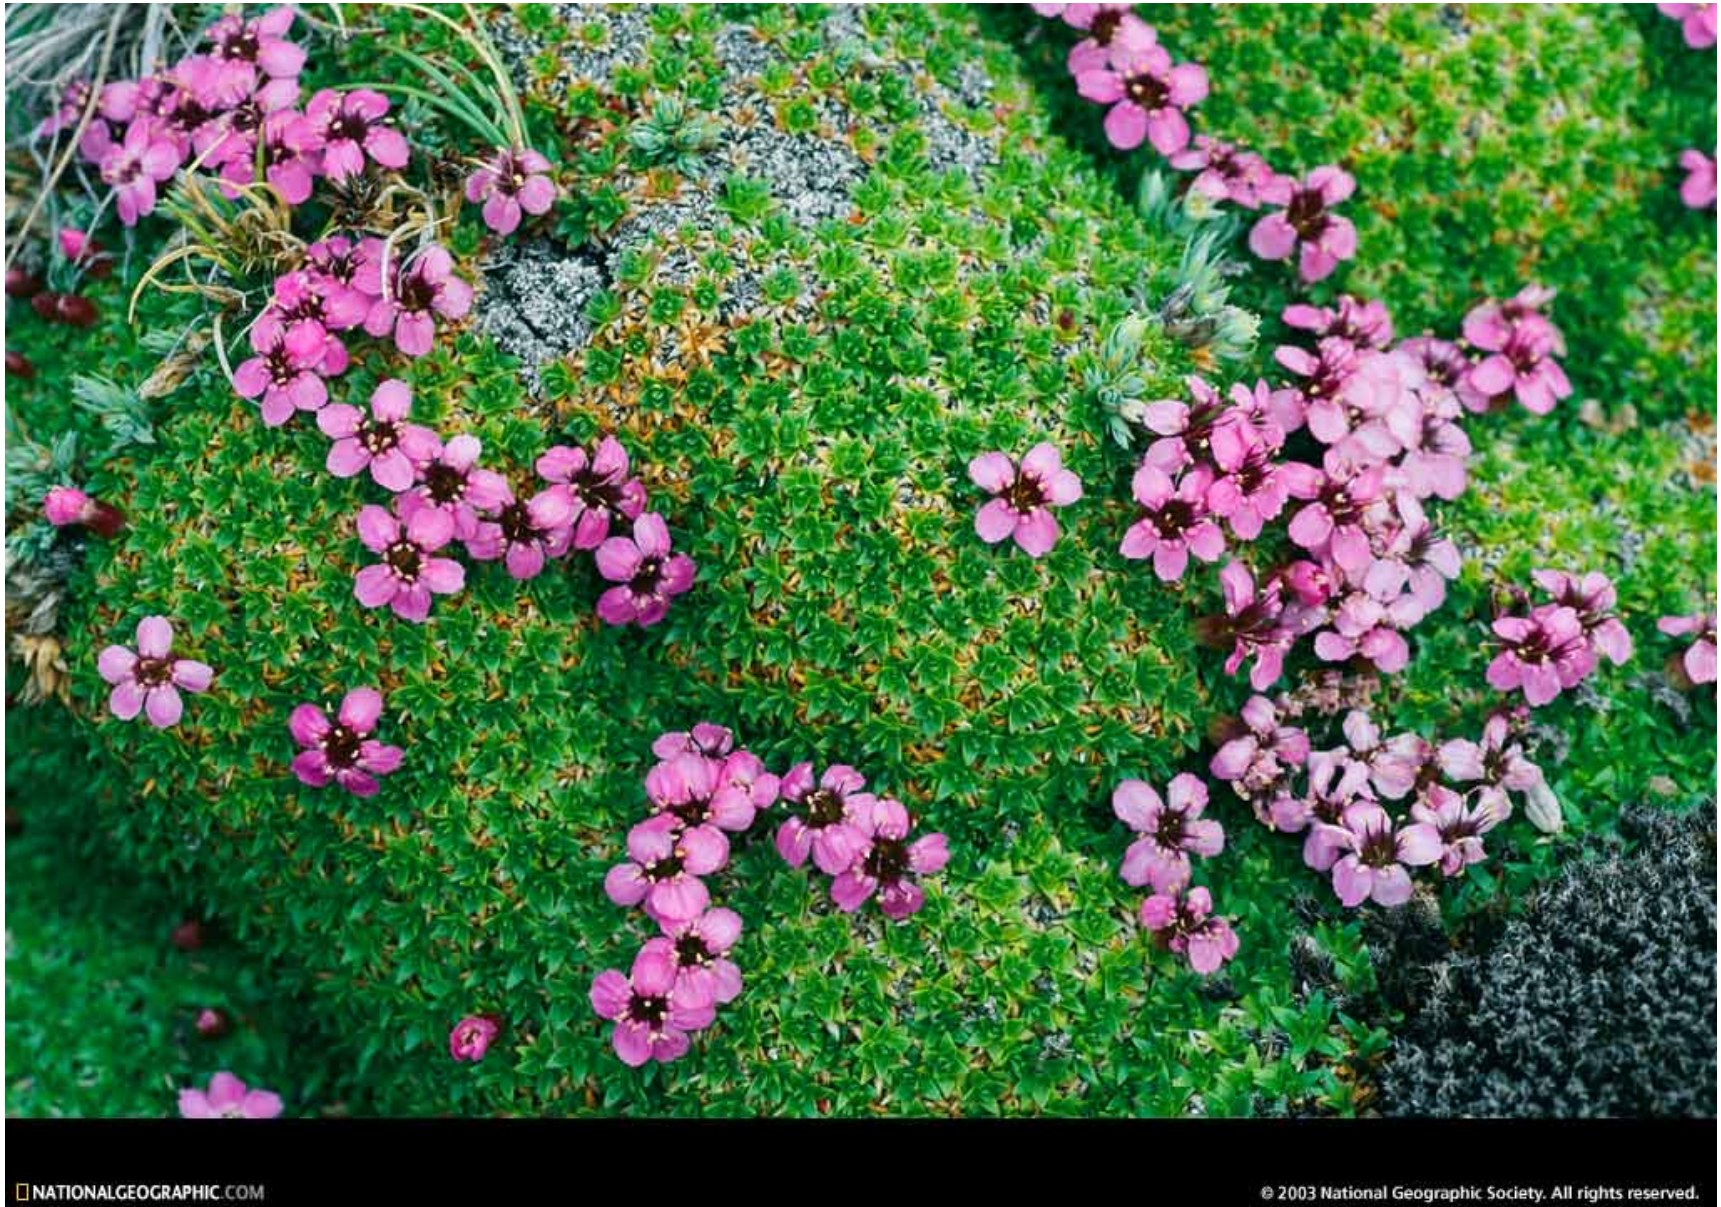

Grouping cues: “Common fate”

Perceptual illusions (limitations)

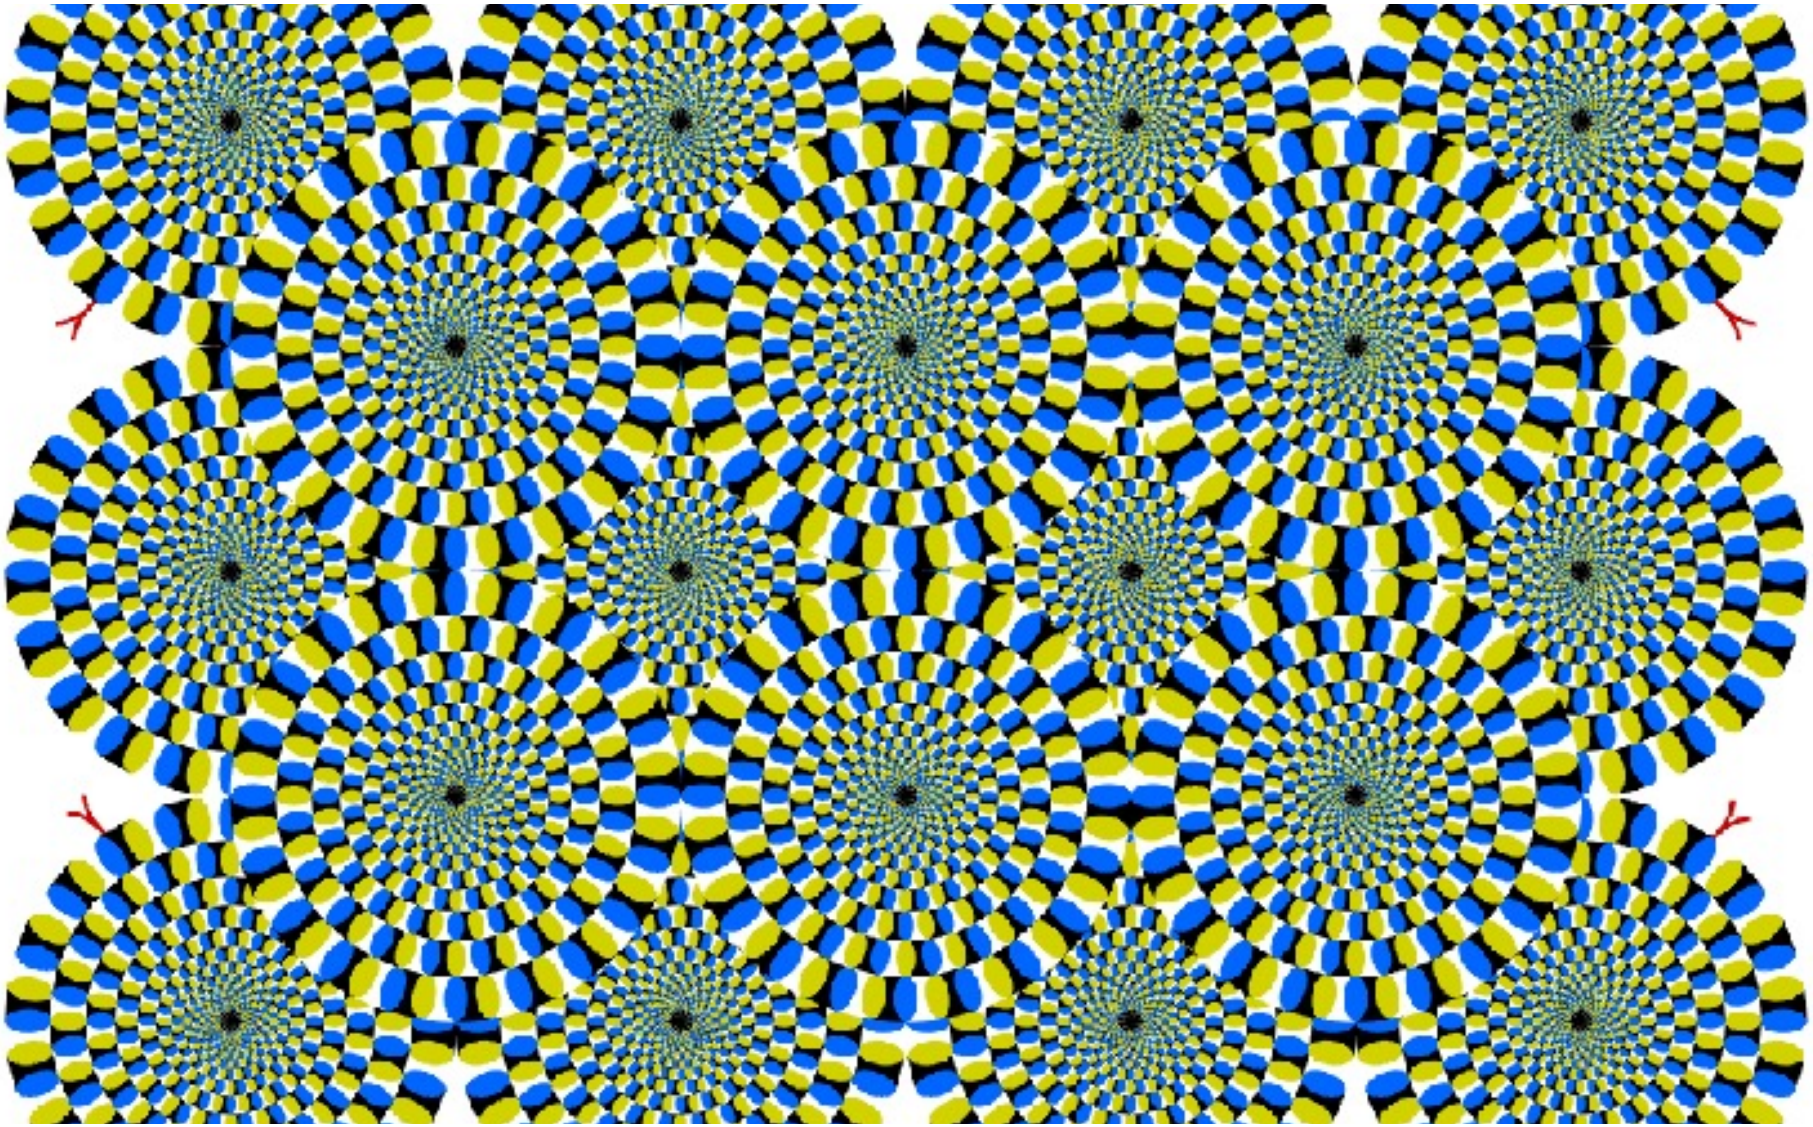

http://www.ritsumei.ac.jp/~akitaoka/index-e.html

Perceptual illusions (concave/convex)

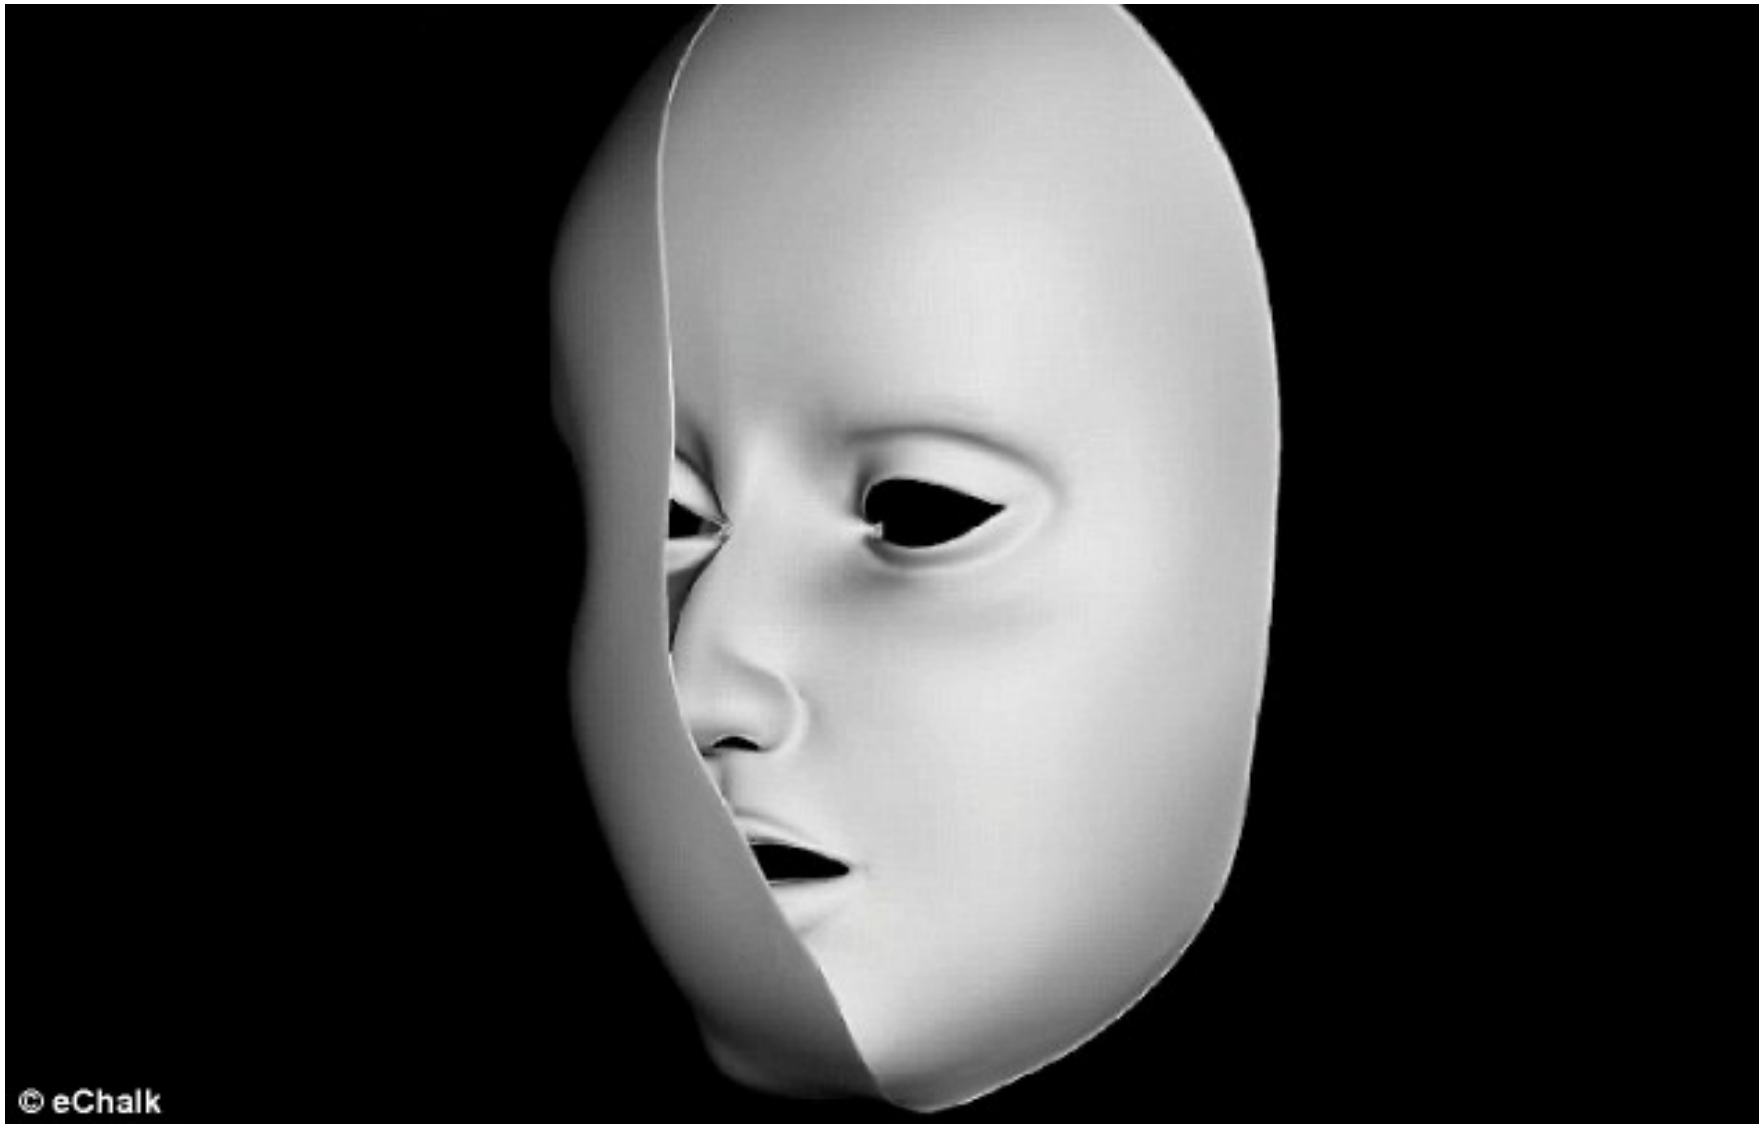

Perception: Rotating cylinders

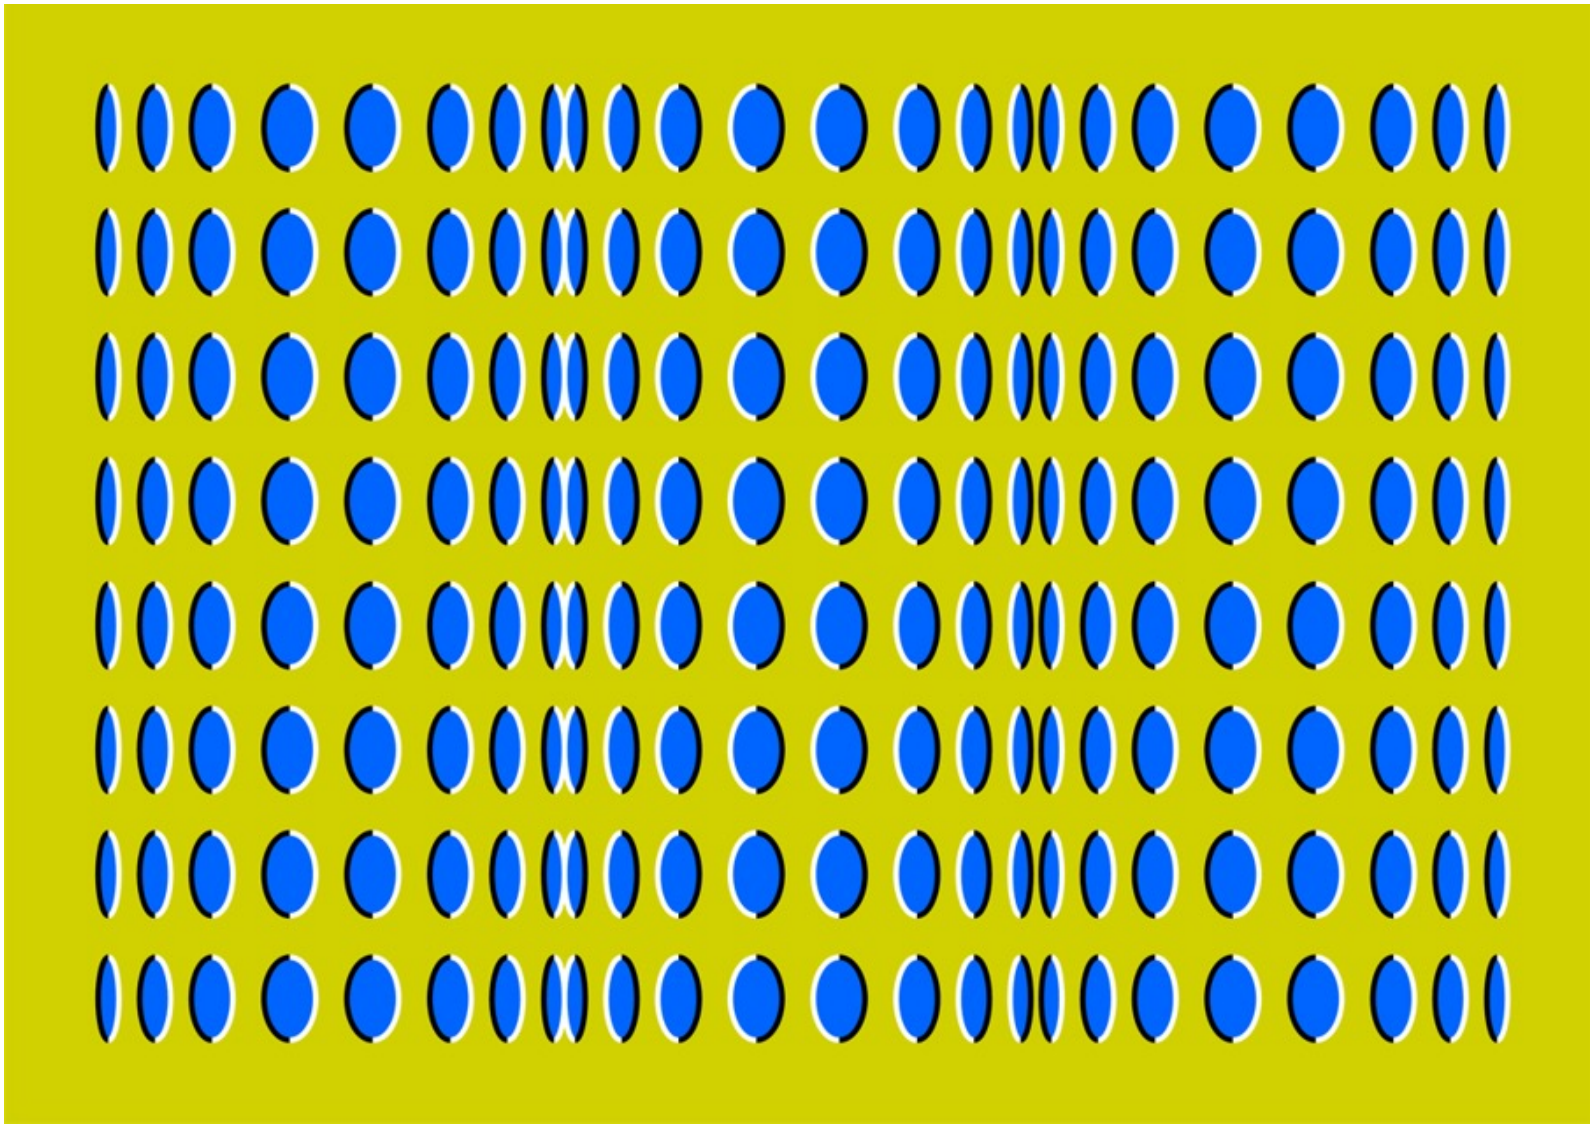

Perception: Chromatic error

All letters are white

Perception: Colors

Lower part looks bluish but all pixels are red

Origins of computer vision

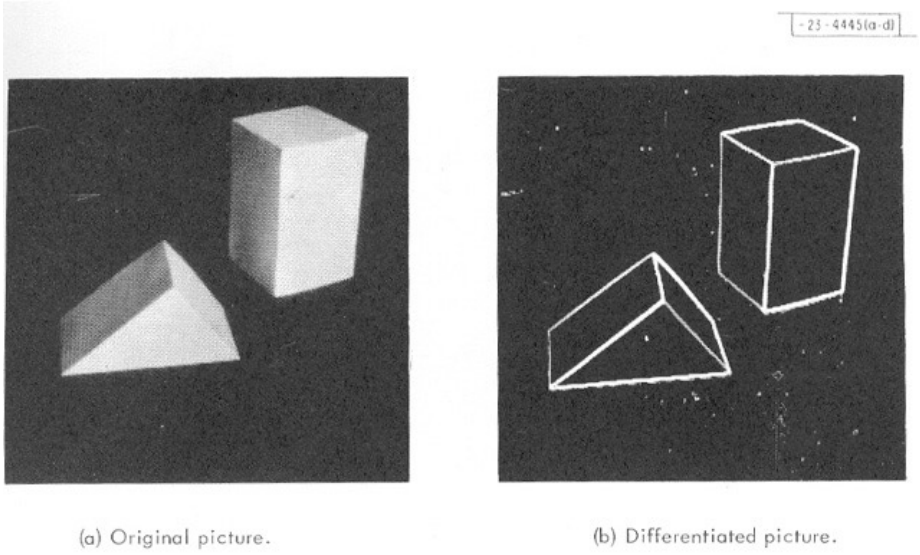

(a) Original picture. (b) Differentiated picture.

(c) Line drawing. (d) Rotated view.

L. G. Roberts *Machine Perception of Three Dimensional Solids*

Connections to other disciplines

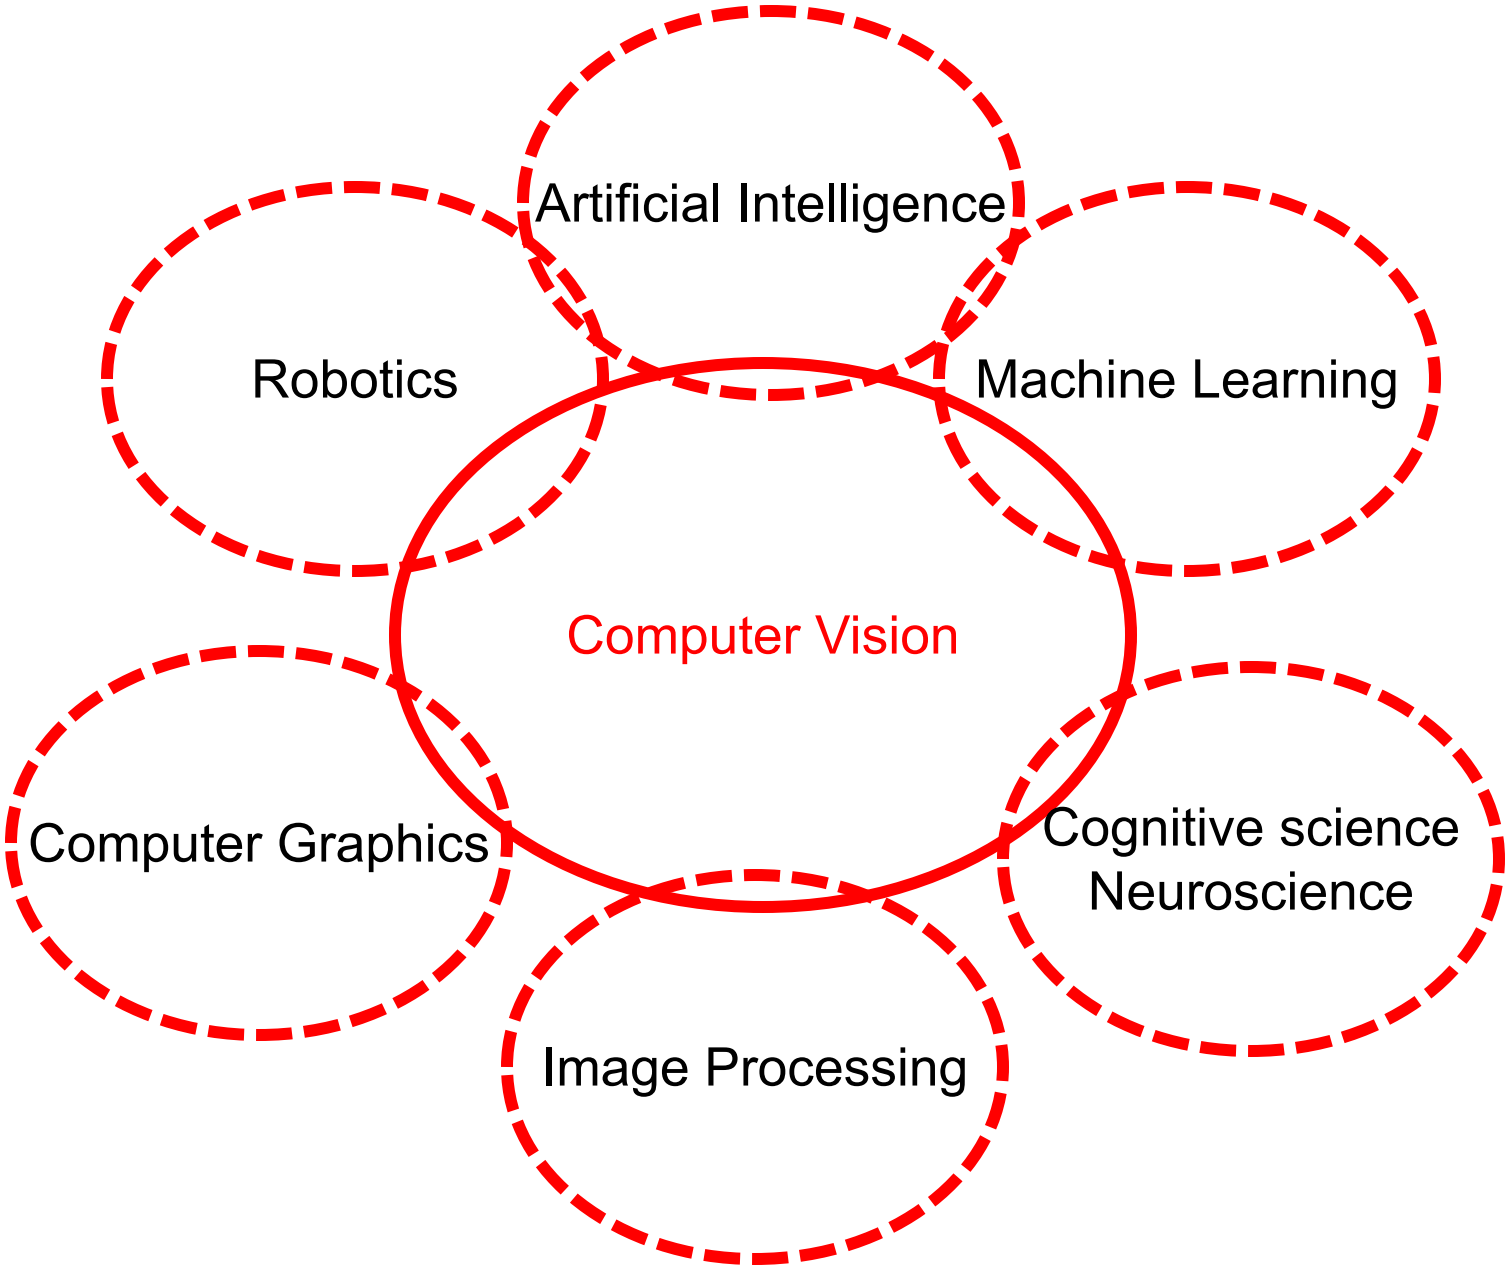

The computer vision industry

- A pre 2015 list of companies here:

  http://www.cs.ubc.ca/spider/lowe/vision.html

- Now there has been an explosion of companies especially in autonomous driving, virtual reality and recognition

Course overview

I. Image formation
   - camera and projection models

II. Image processing
   - filtering, edges, minimization framework

III. Correspondence
   - tracking, optical flow, registration, interest points, fitting

IV. 3D motion and geometry
   - epipolar geometry, (multiview) stereo and structure from motion

V. Numerical methods
   - Bayesian framework, advanced optimization techniques

VI. Recognition and segmentation

Early vision

- Basic image formation and processing

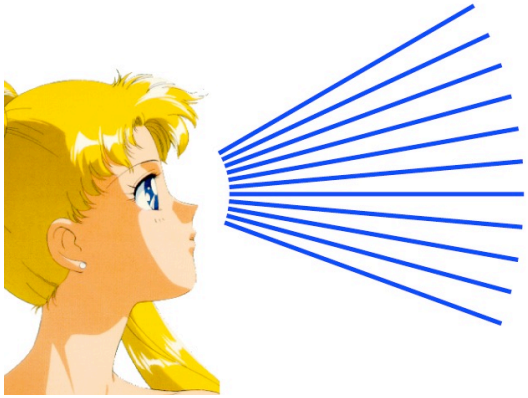

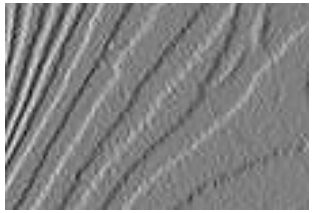

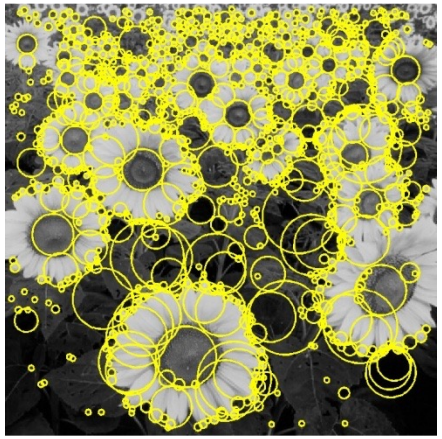

Correspondence and 3D

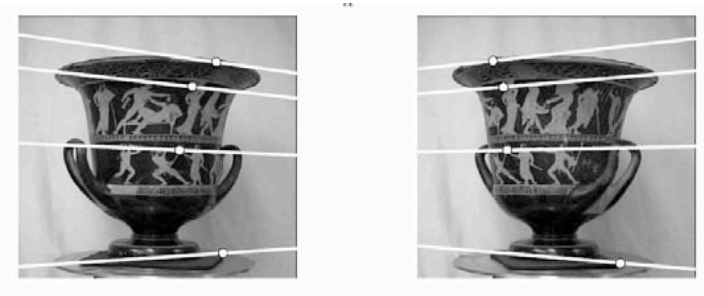

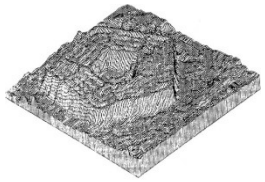

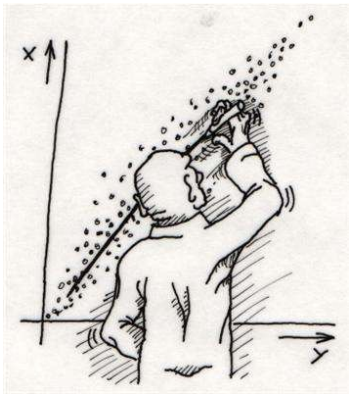

Драконъ, видимый подъ различными углами зрѣнія
По гравюрѣ на мѣди изъ „Oculus artificialis telediopricus“ Цана. 1702 года.

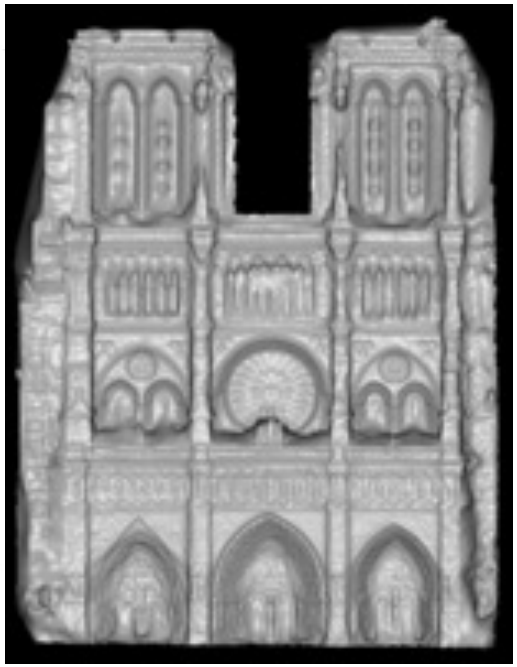

Numerical methods

- Bayesian framework, segmentation