

Image segmentation



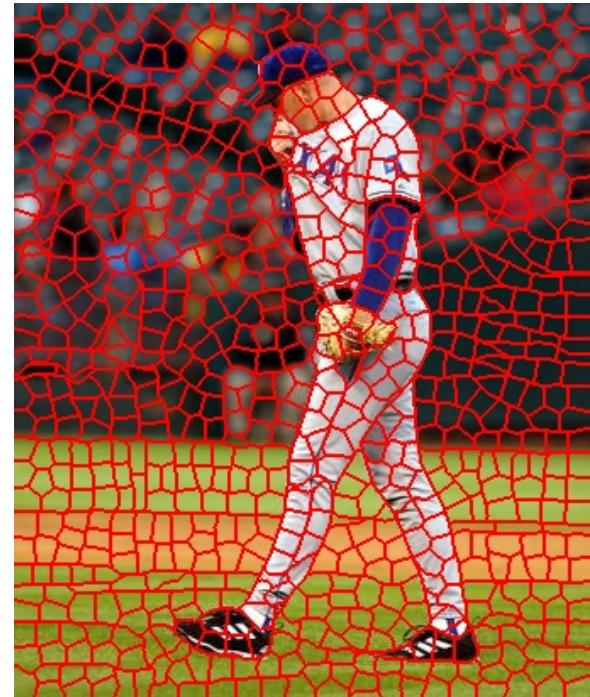
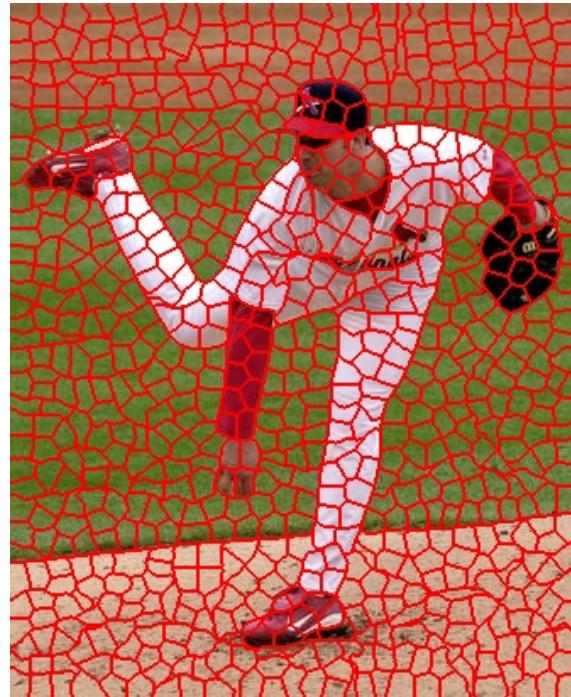
The goals of segmentation

- Obtain primitives for other tasks
 - e.g., feature support
- Perceptual organization, recognition
- Graphics, image manipulation

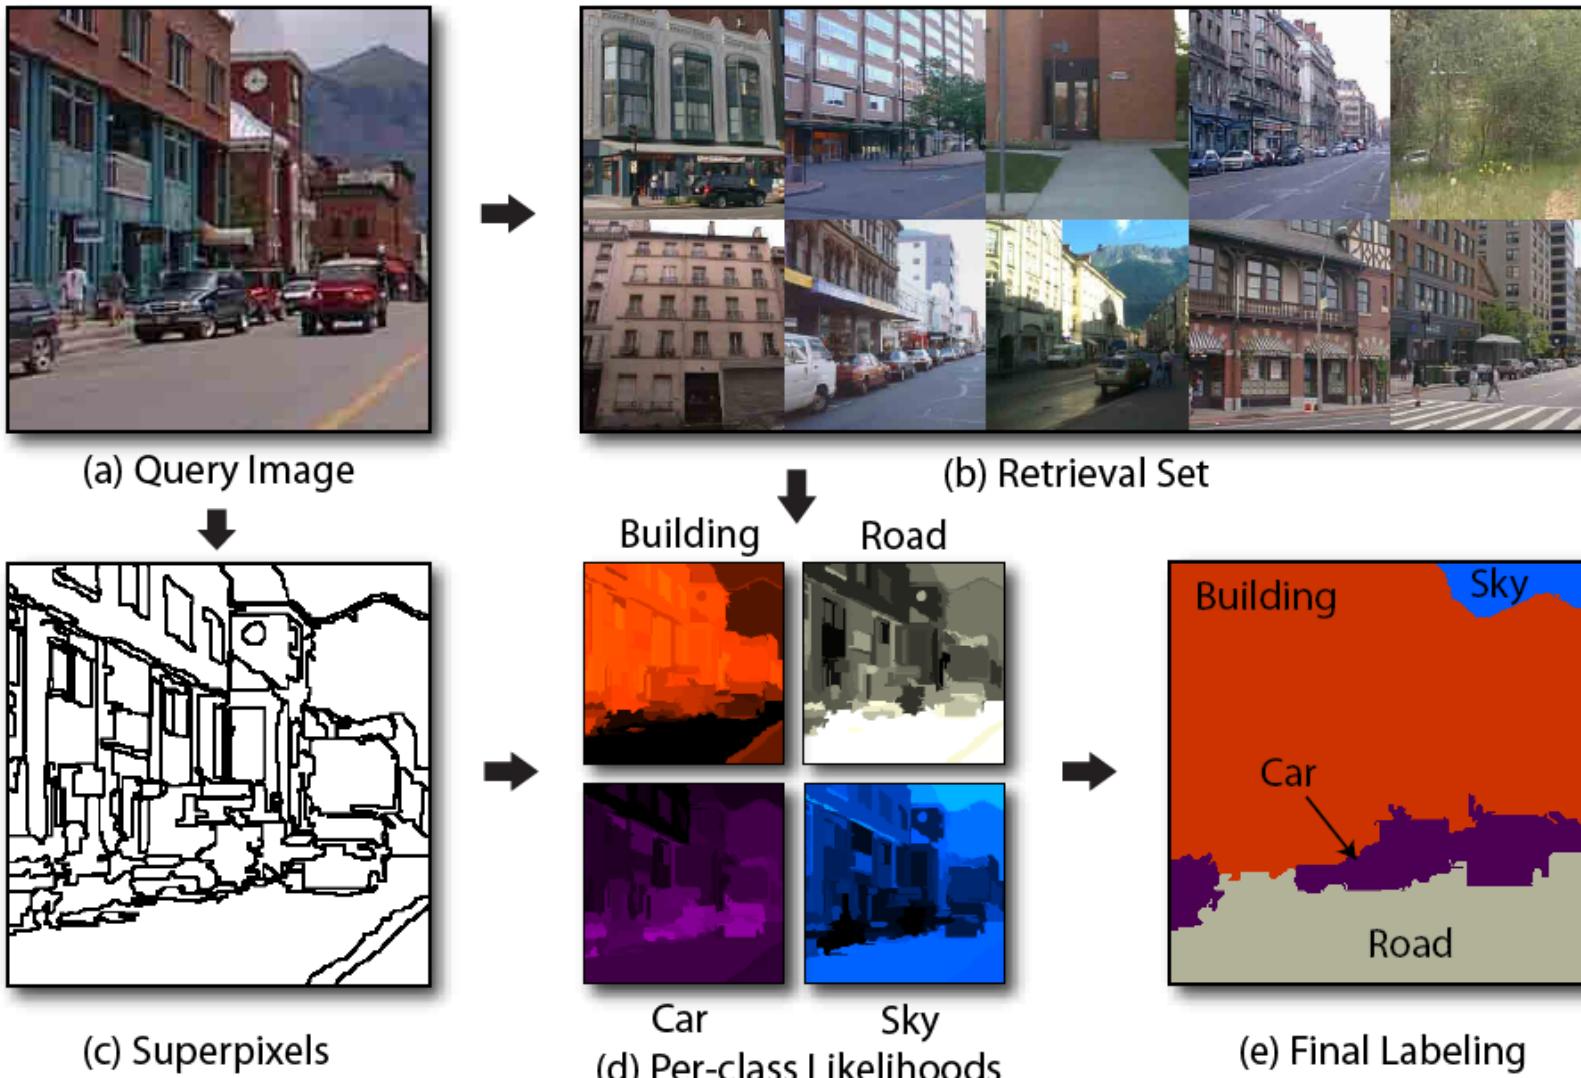
The goals of segmentation

- Group together similar-looking pixels for efficiency of further processing
 - “Bottom-up” process
 - Unsupervised

“superpixels”



Segments as primitives for recognition

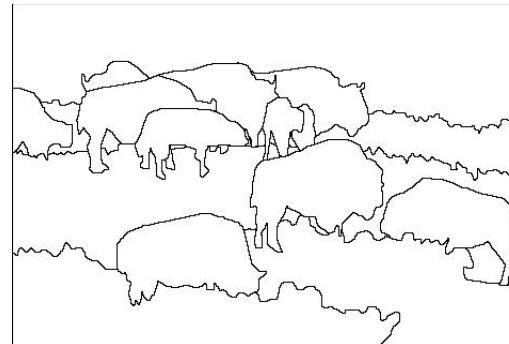


The goals of segmentation

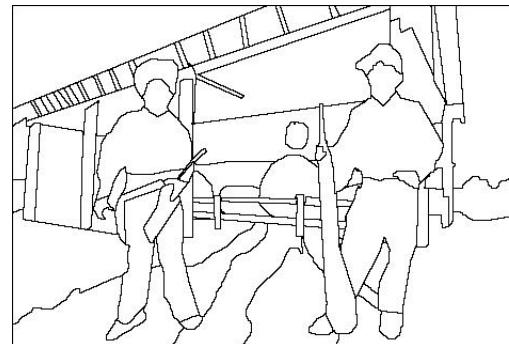
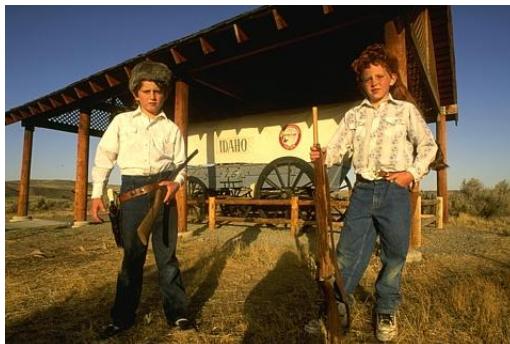
- Separate image into coherent “objects”
 - “Bottom-up” or “top-down” process?
 - Supervised or unsupervised?



image



human segmentation

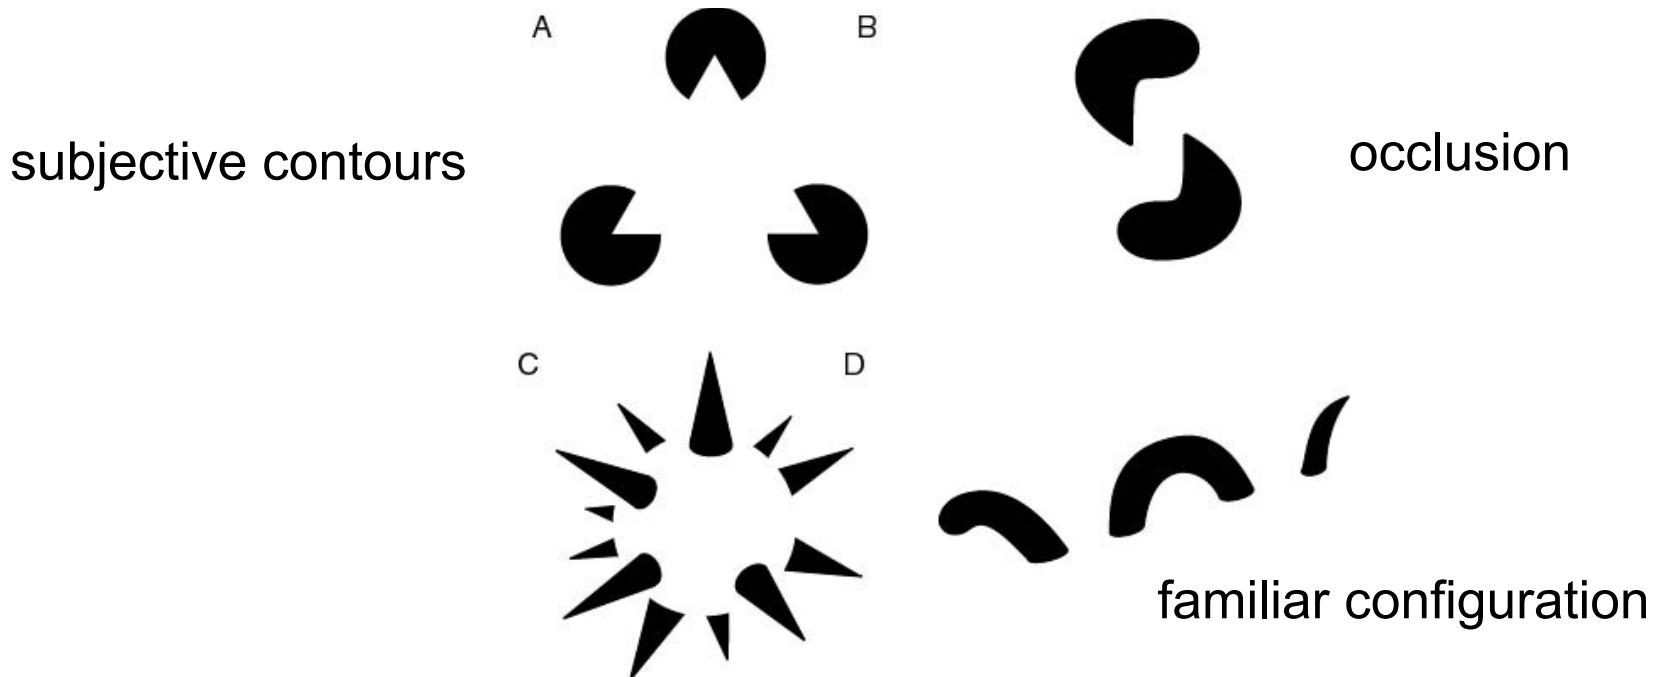


Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

Inspiration from psychology

- The Gestalt school: Grouping is key to visual perception
 - “The whole is greater than the sum of its parts”



Emergence



http://en.wikipedia.org/wiki/Gestalt_psychology

The goals of segmentation

- Interactive segmentation for graphics

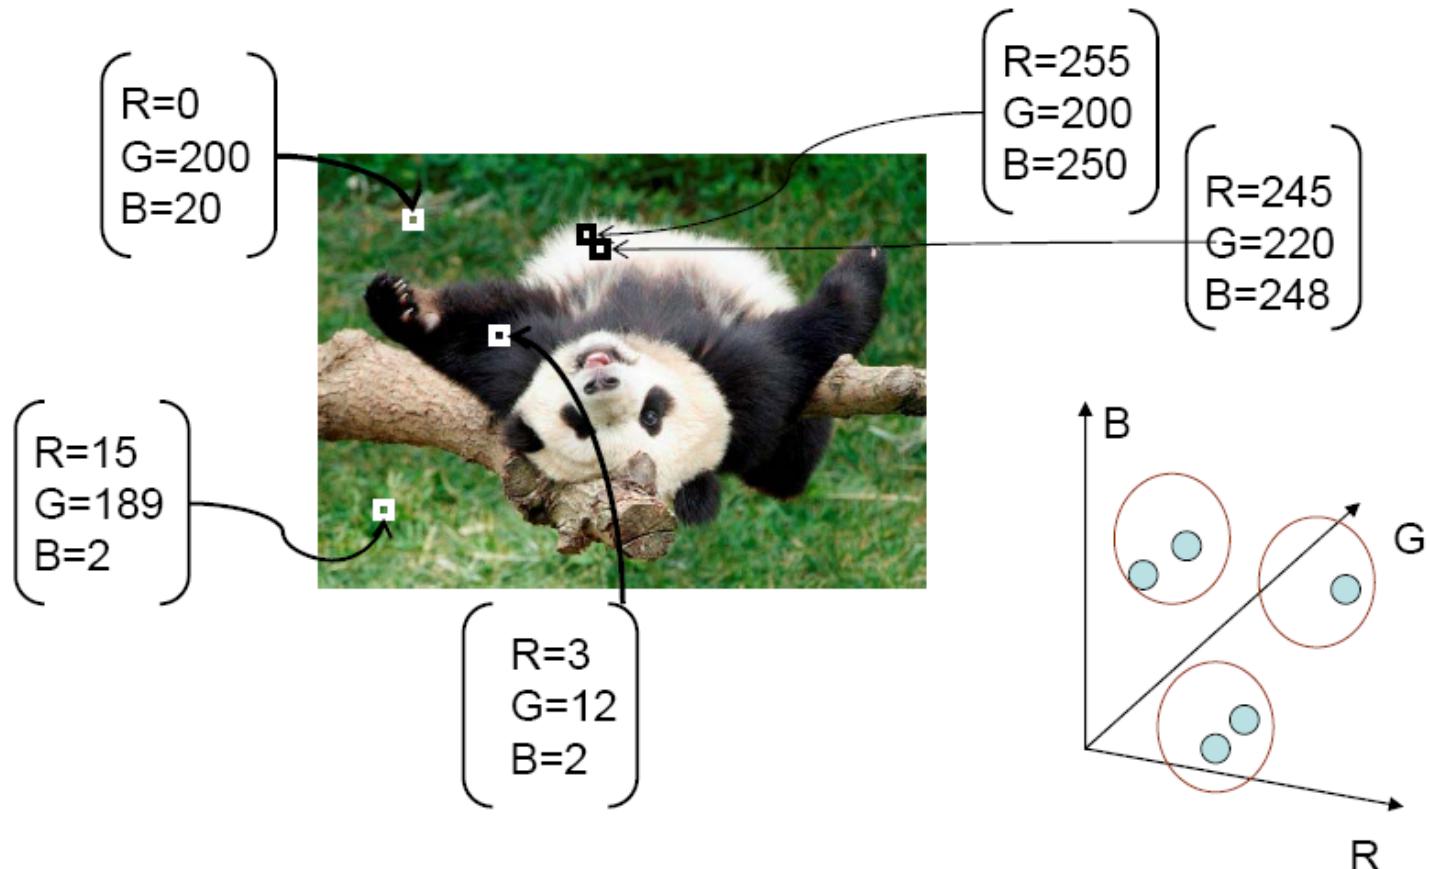


Approaches to segmentation

- Segmentation as clustering
- Segmentation as graph partitioning
- Segmentation as labeling

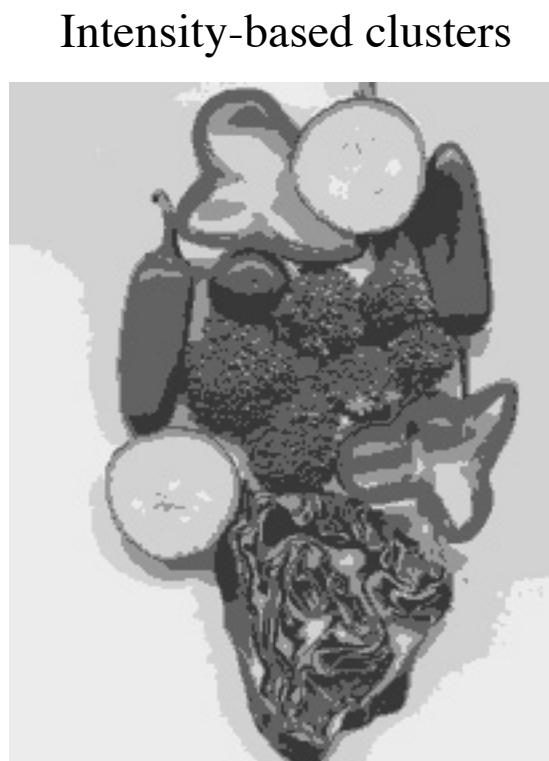
Segmentation as clustering

- Cluster similar pixels (features) together



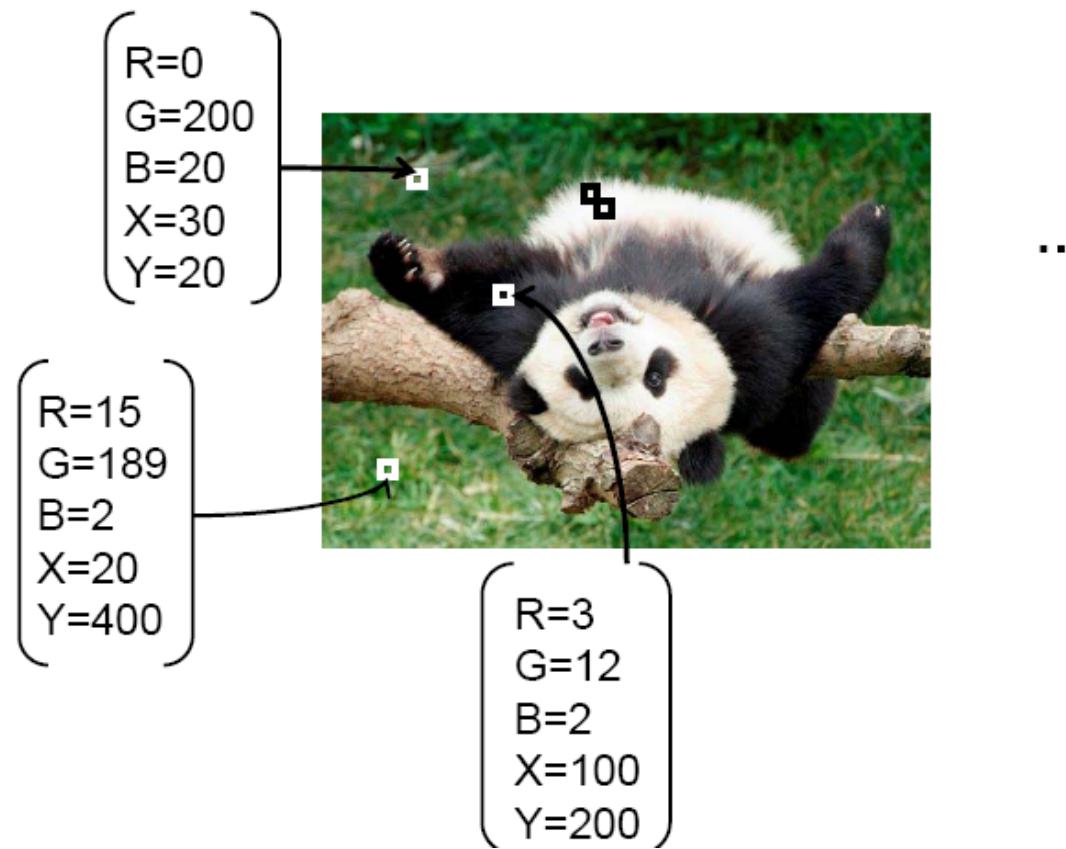
Segmentation as clustering

- K-means clustering based on intensity or color is essentially vector quantization of the image attributes
 - Clusters don't have to be spatially coherent



Segmentation as clustering

- Cluster similar pixels (features) together



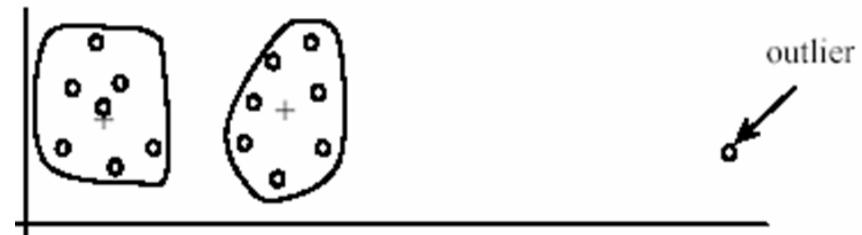
Segmentation as clustering

- Clustering based on (r,g,b,x,y) values enforces more spatial coherence

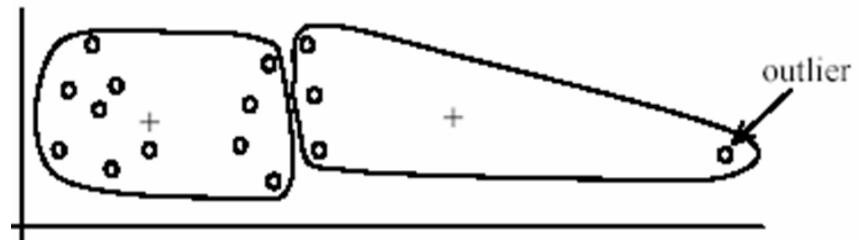


K-Means pros and cons

- Pros
 - Simple and fast
 - Easy to implement
- Cons
 - Need to choose K
 - Sensitive to outliers
- Usage
 - Rarely used for pixel segmentation



(B): Ideal clusters

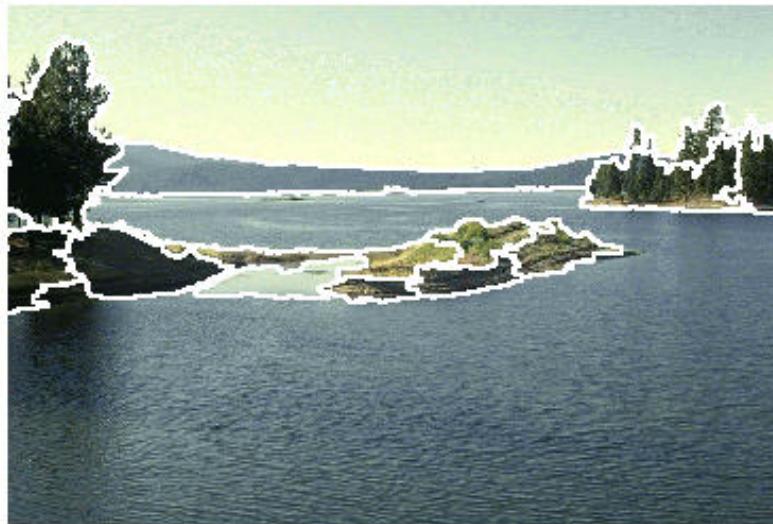


Mean shift segmentation

D. Comaniciu and P. Meer, Mean Shift: A Robust Approach toward Feature Space Analysis, PAMI 2002.

Versatile technique for clustering-based segmentation

Segmented "landscape 1"

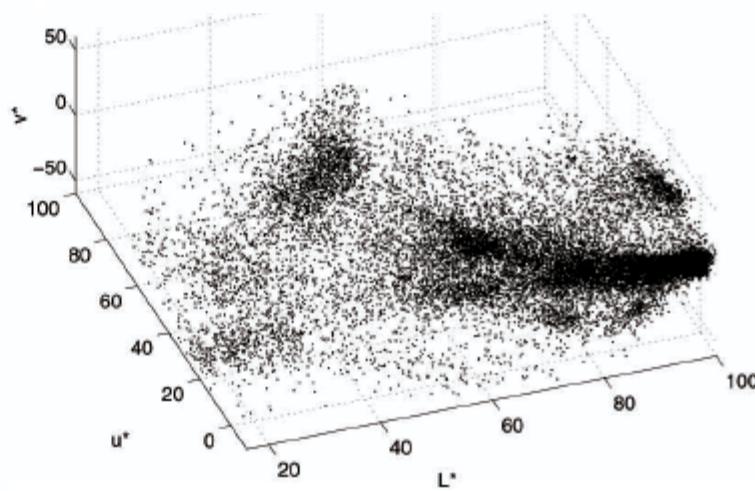
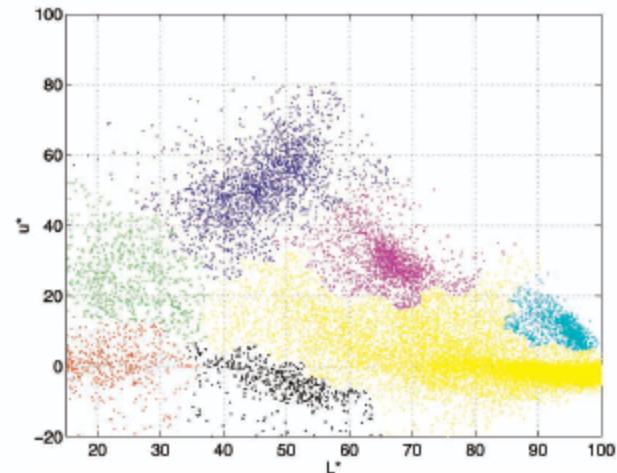
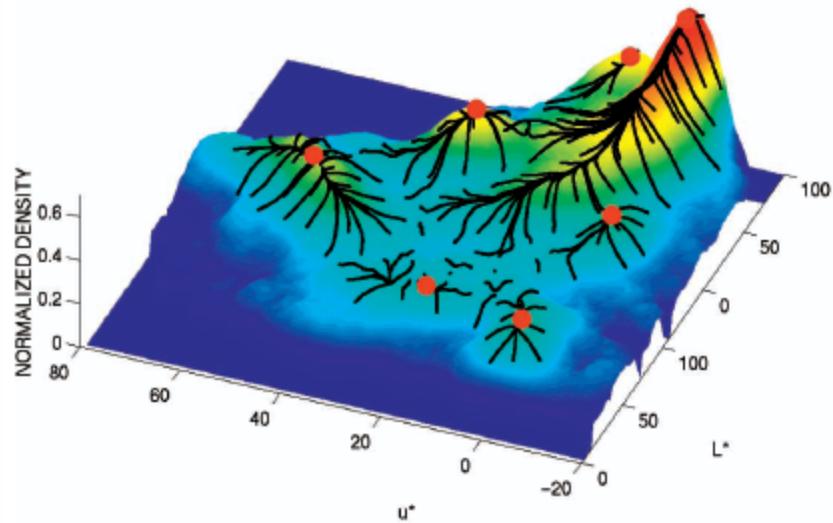


Segmented "landscape 2"

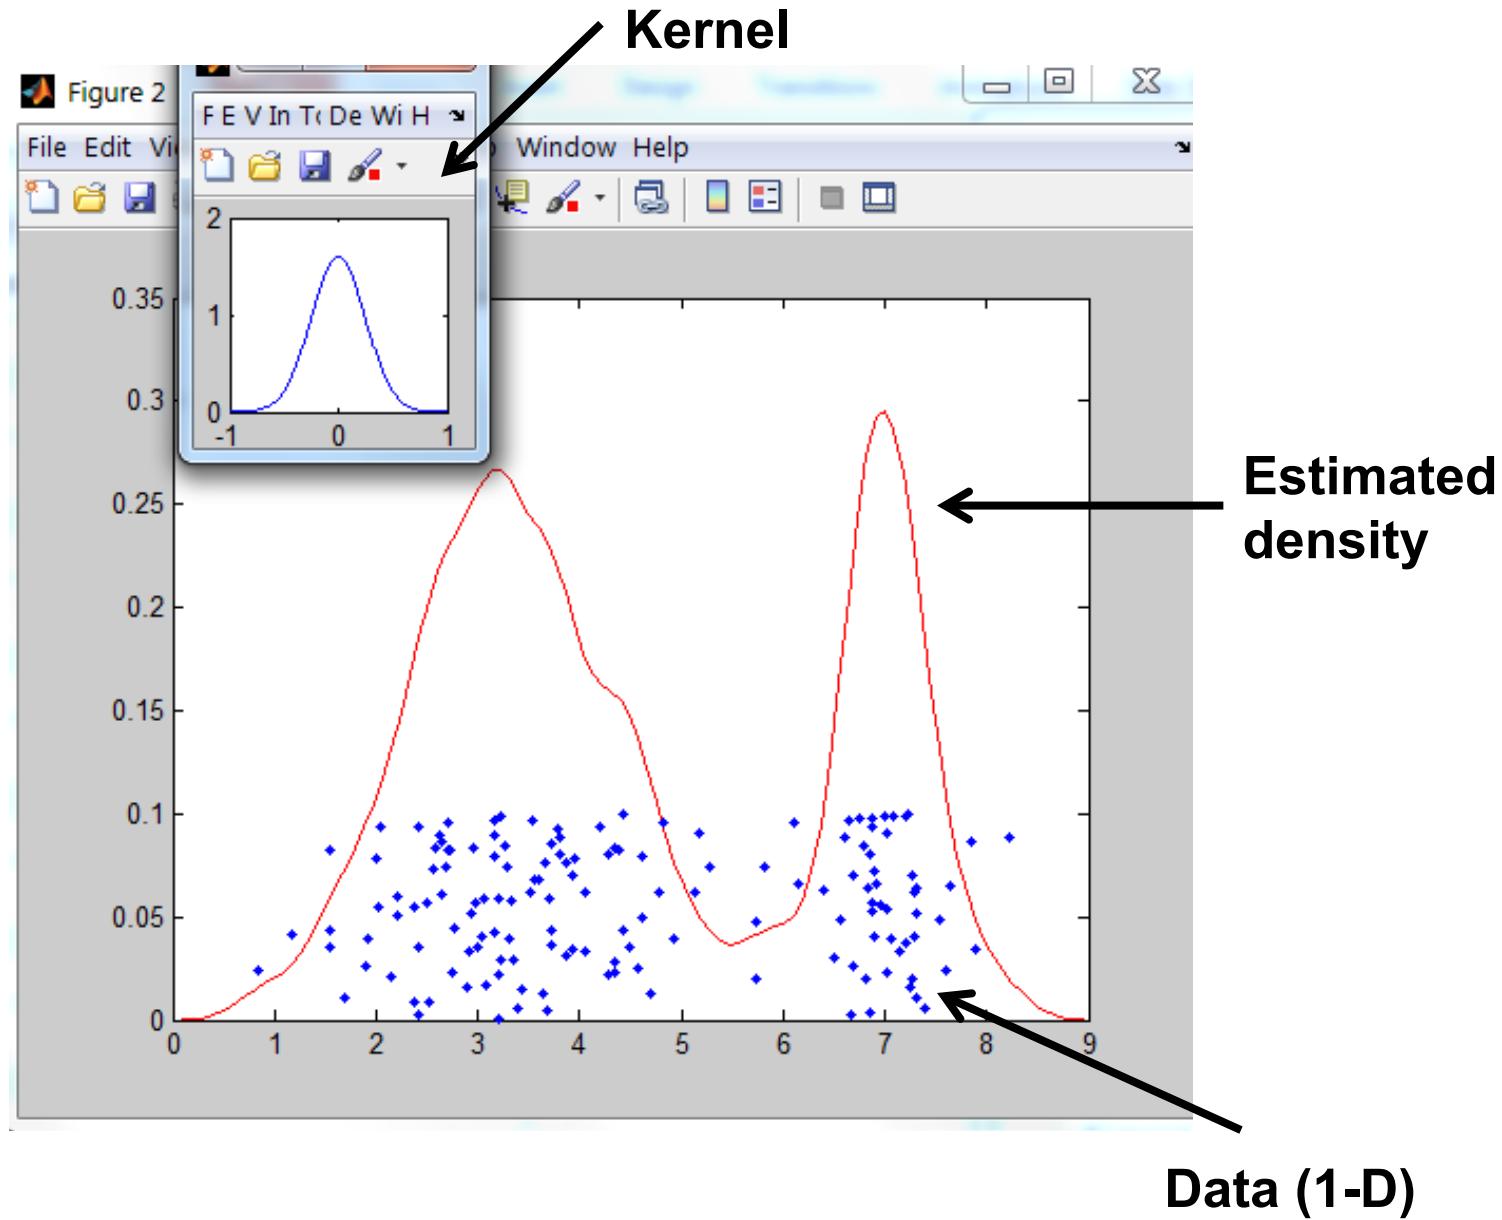


Mean shift algorithm

Try to find *modes* of this non-parametric density



Kernel density estimation



Kernel density estimation

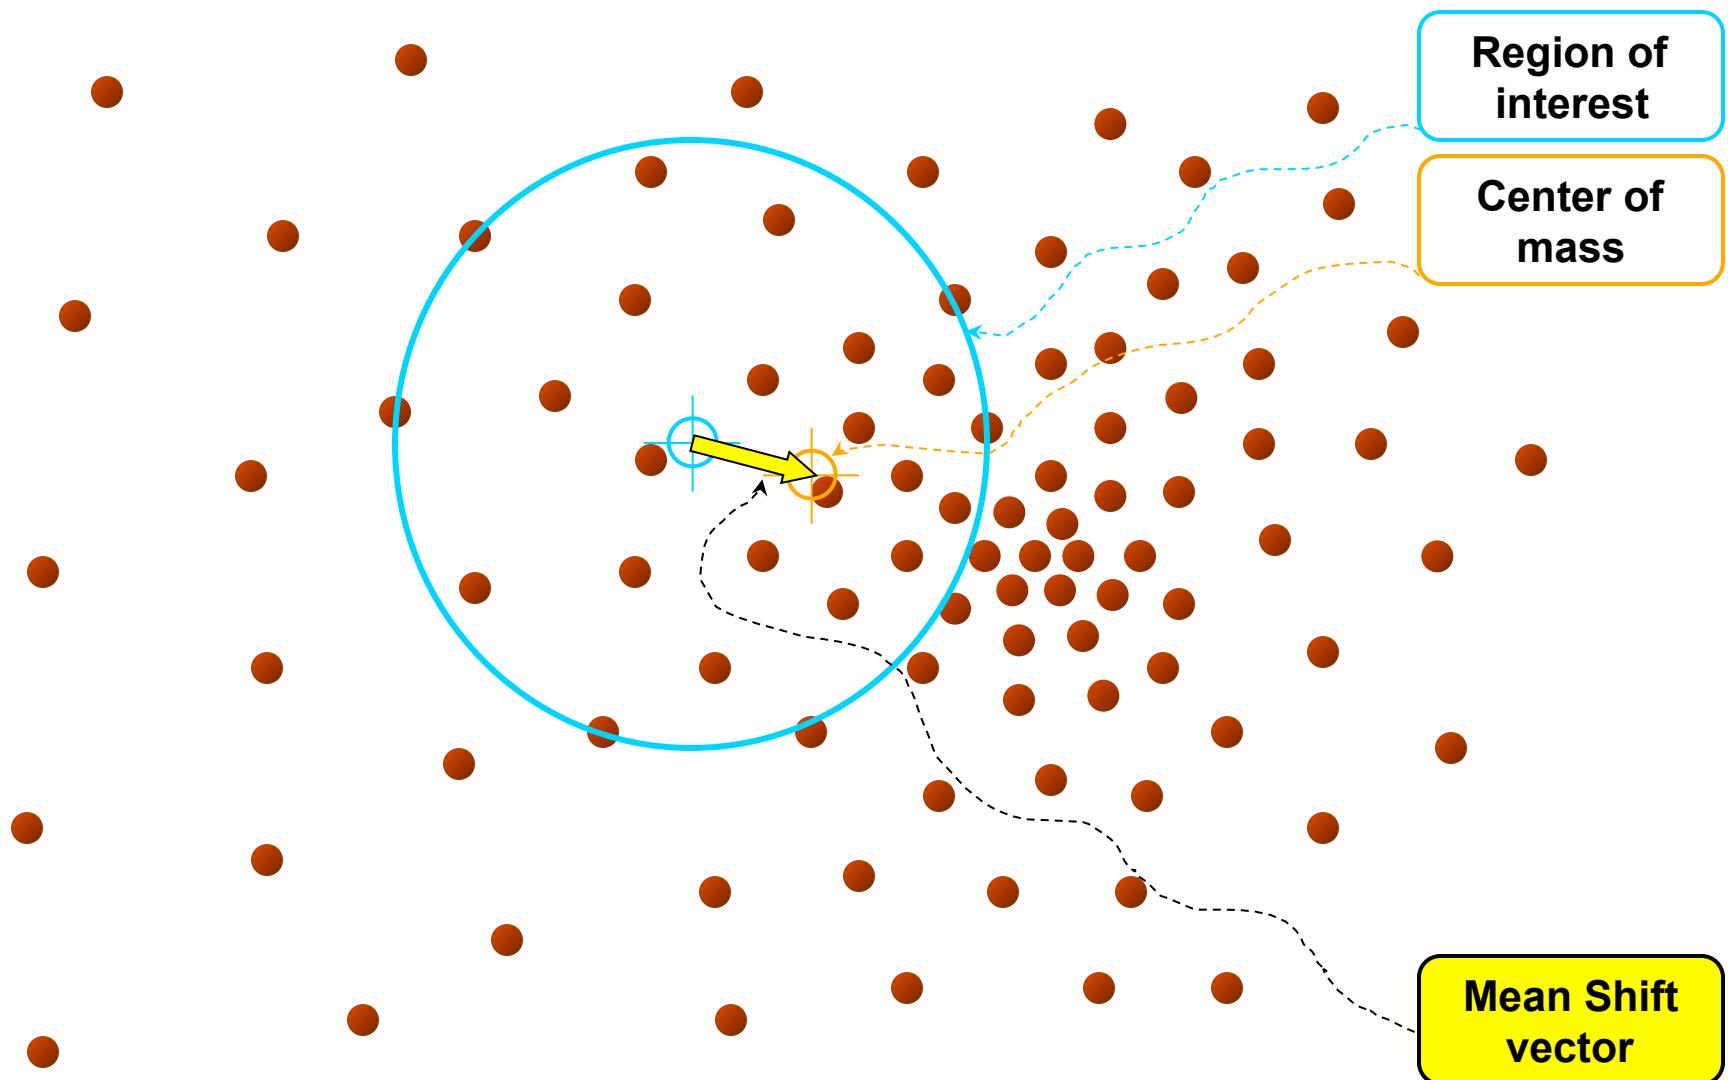
Kernel density estimation function

$$\hat{f}_h(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

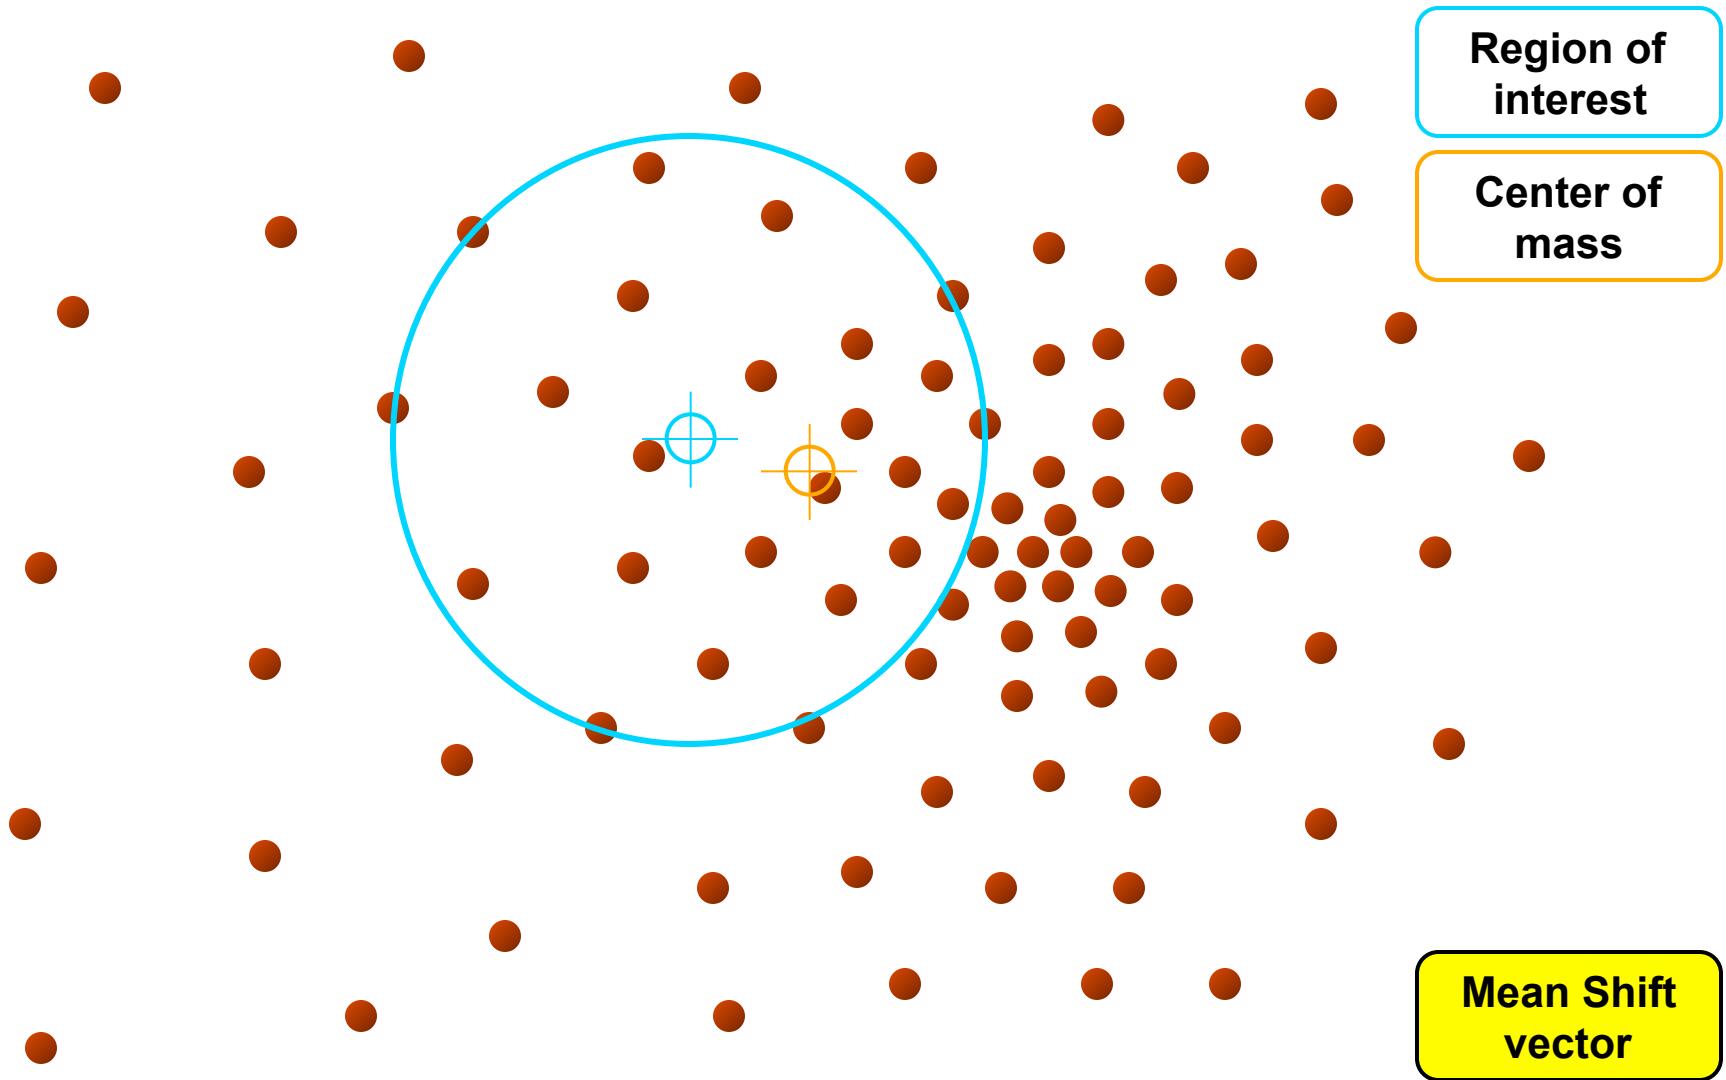
Gaussian kernel

$$K\left(\frac{x - x_i}{h}\right) = \frac{1}{\sqrt{2\pi}} e^{-\frac{(x-x_i)^2}{2h^2}}.$$

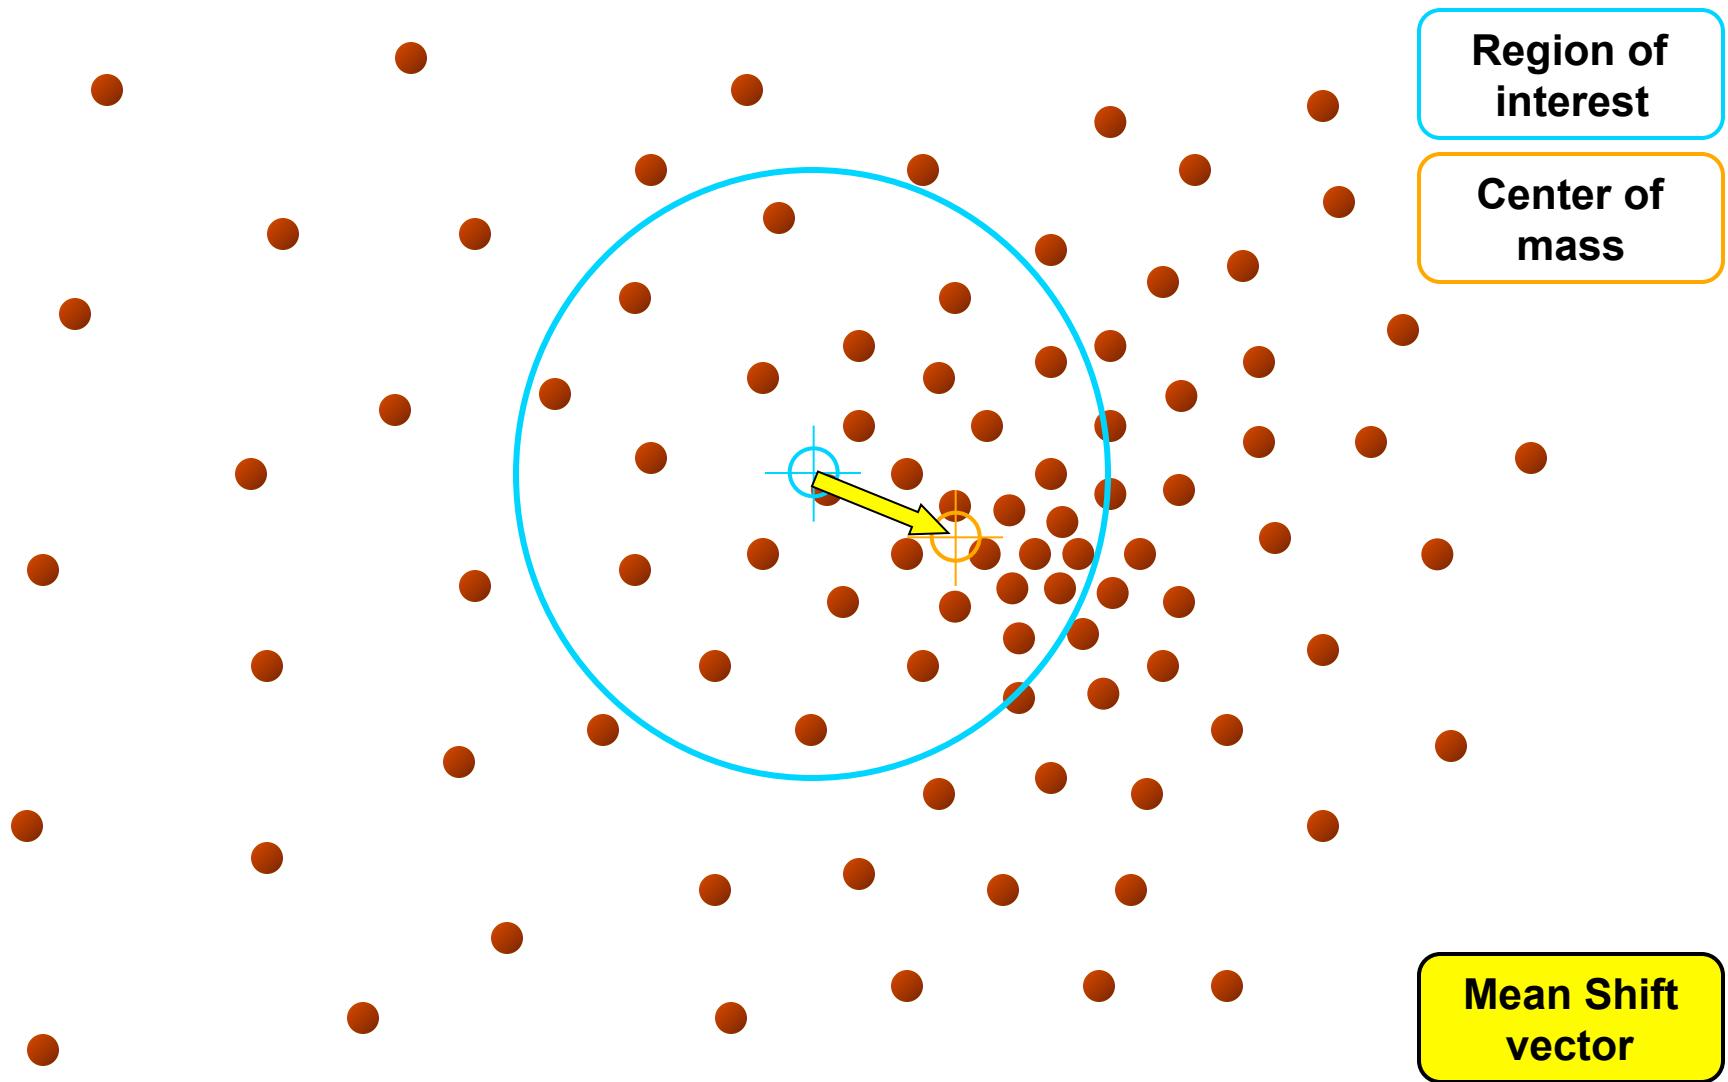
Mean shift



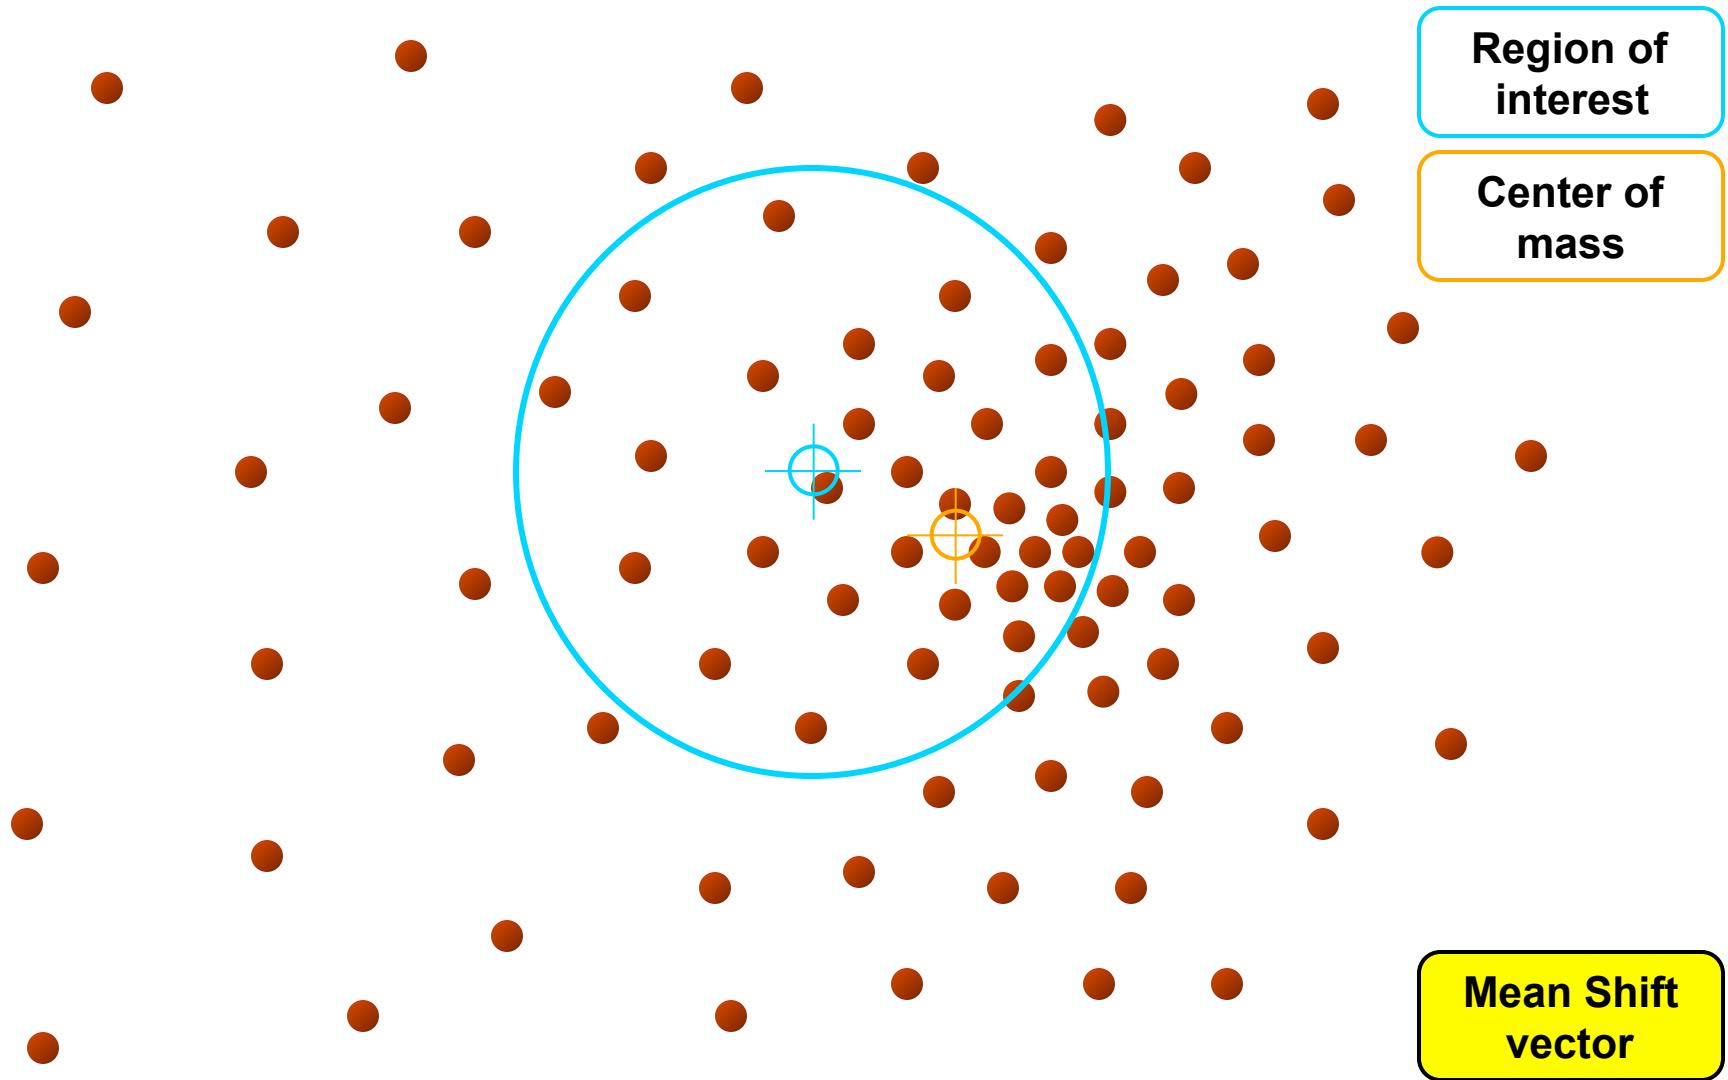
Mean shift



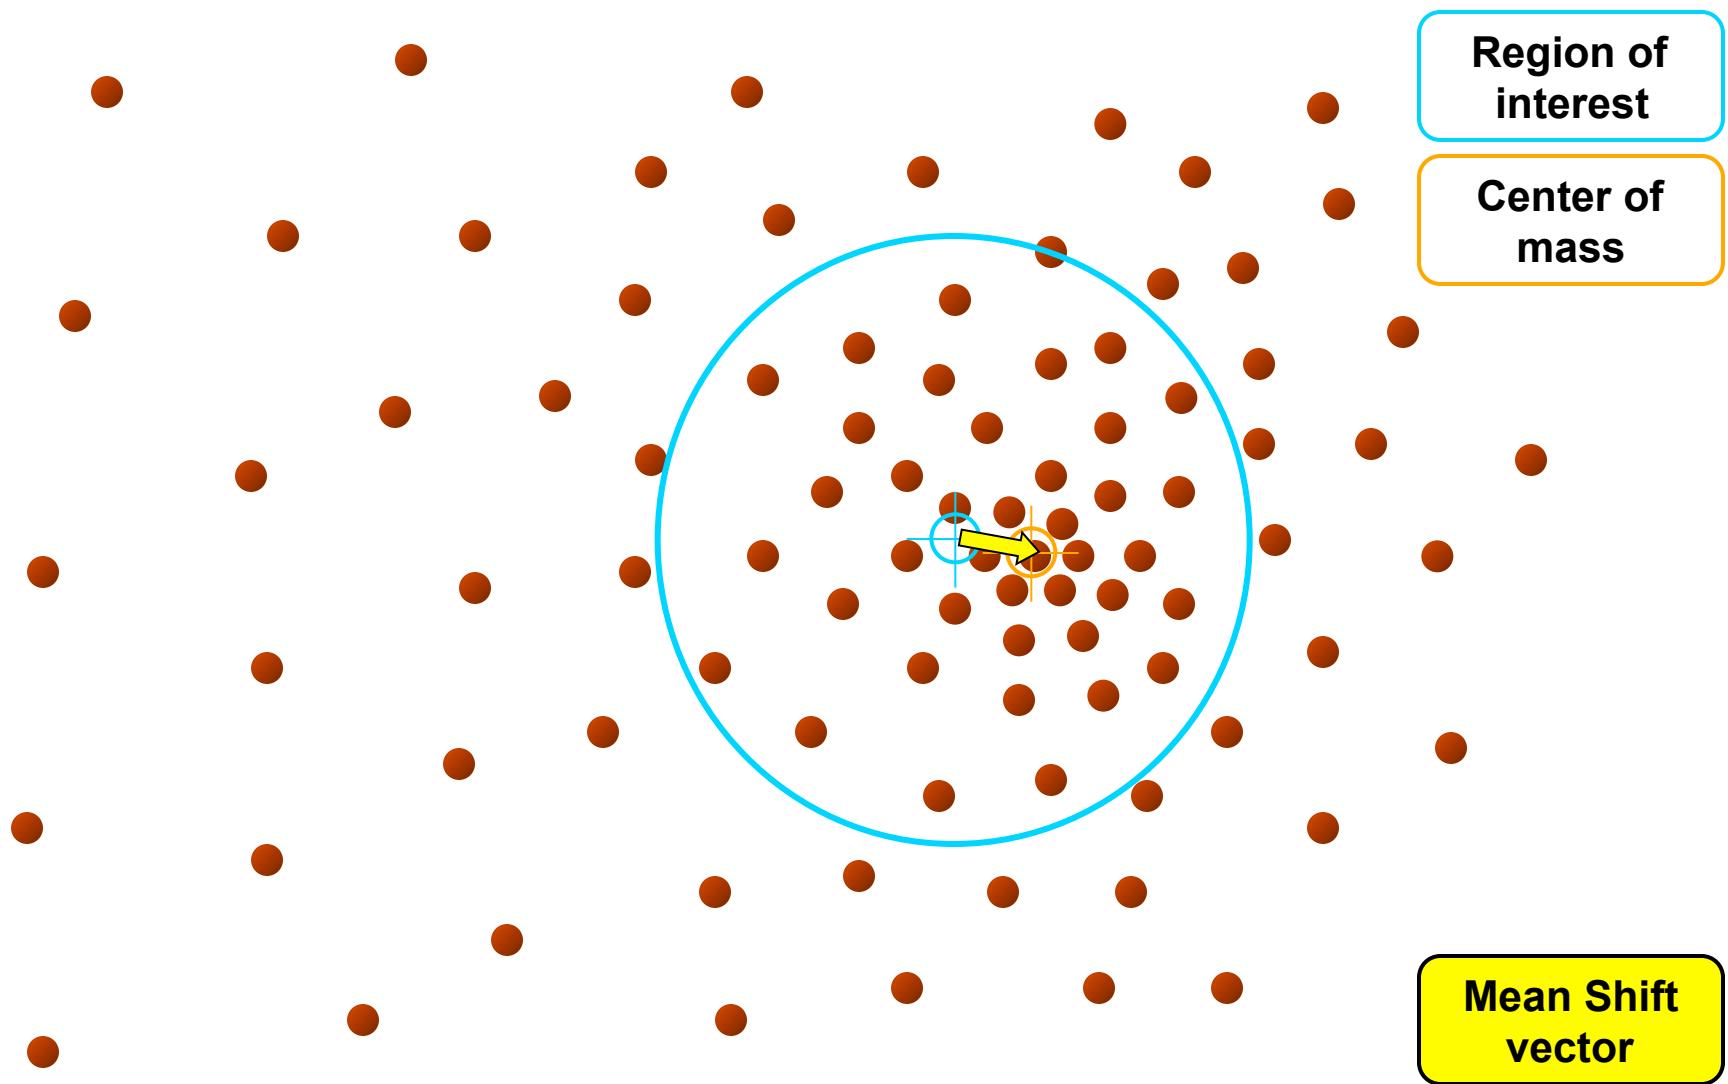
Mean shift



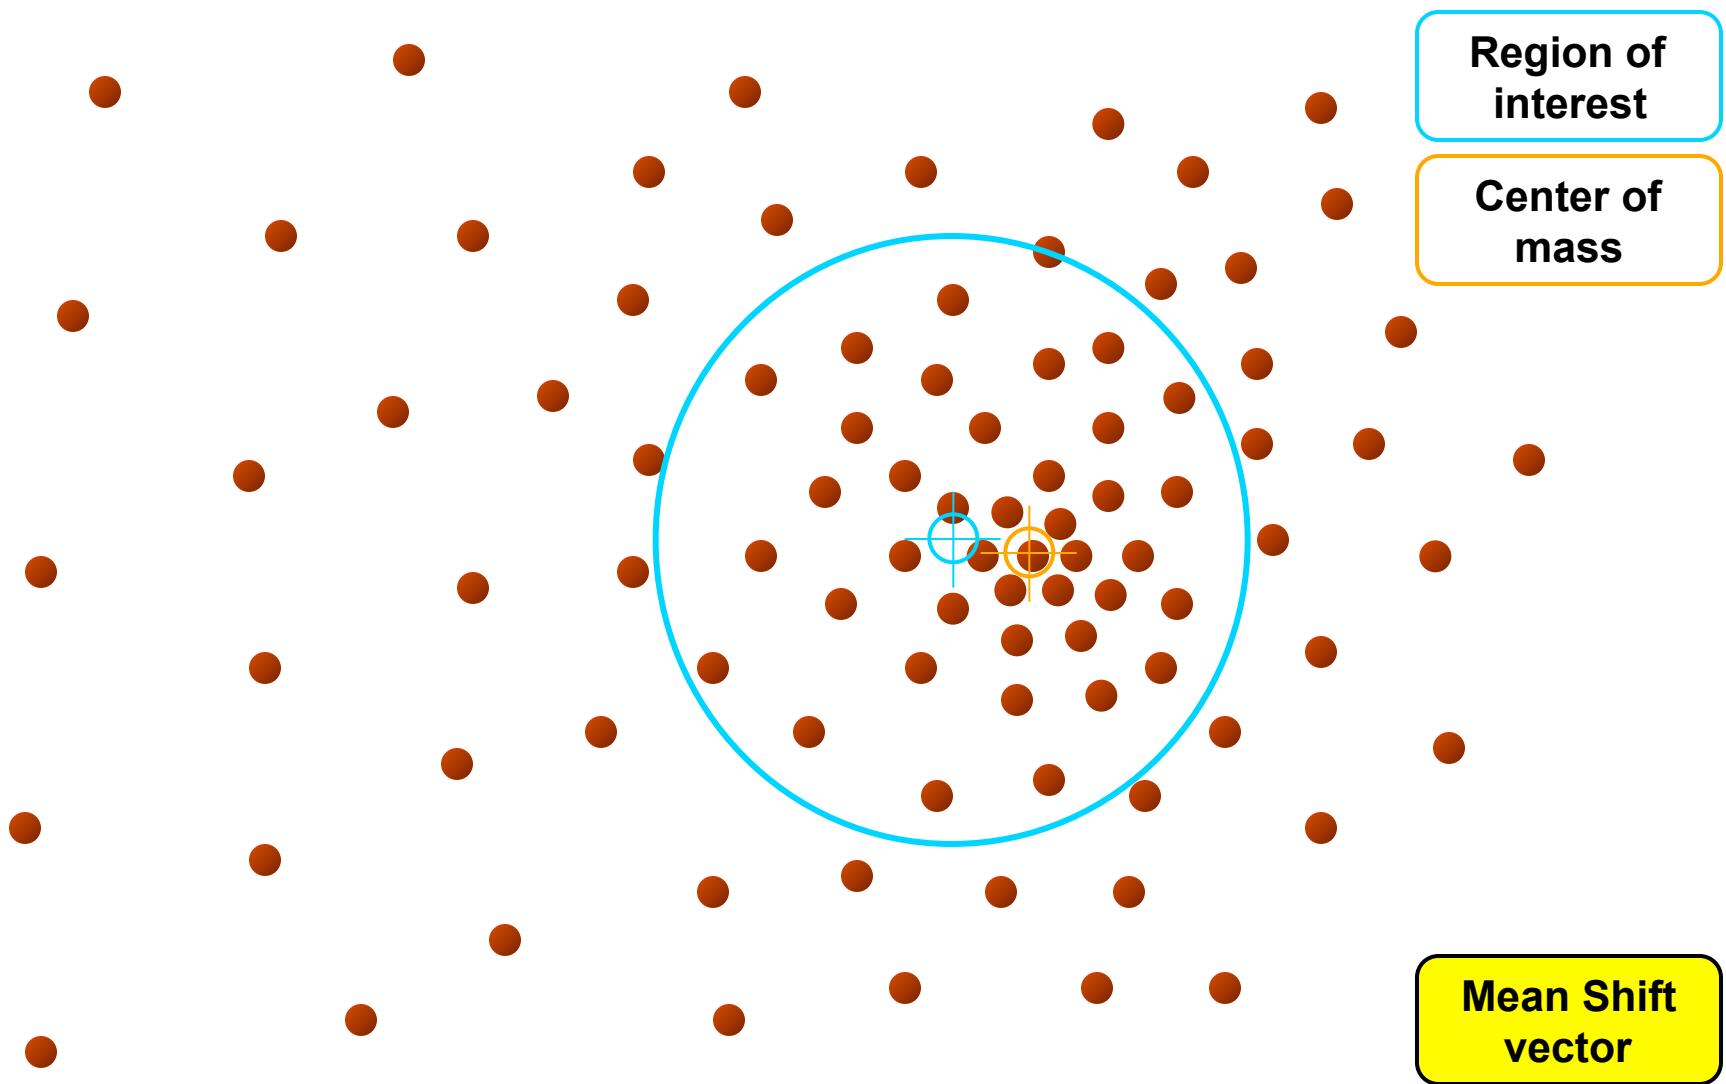
Mean shift



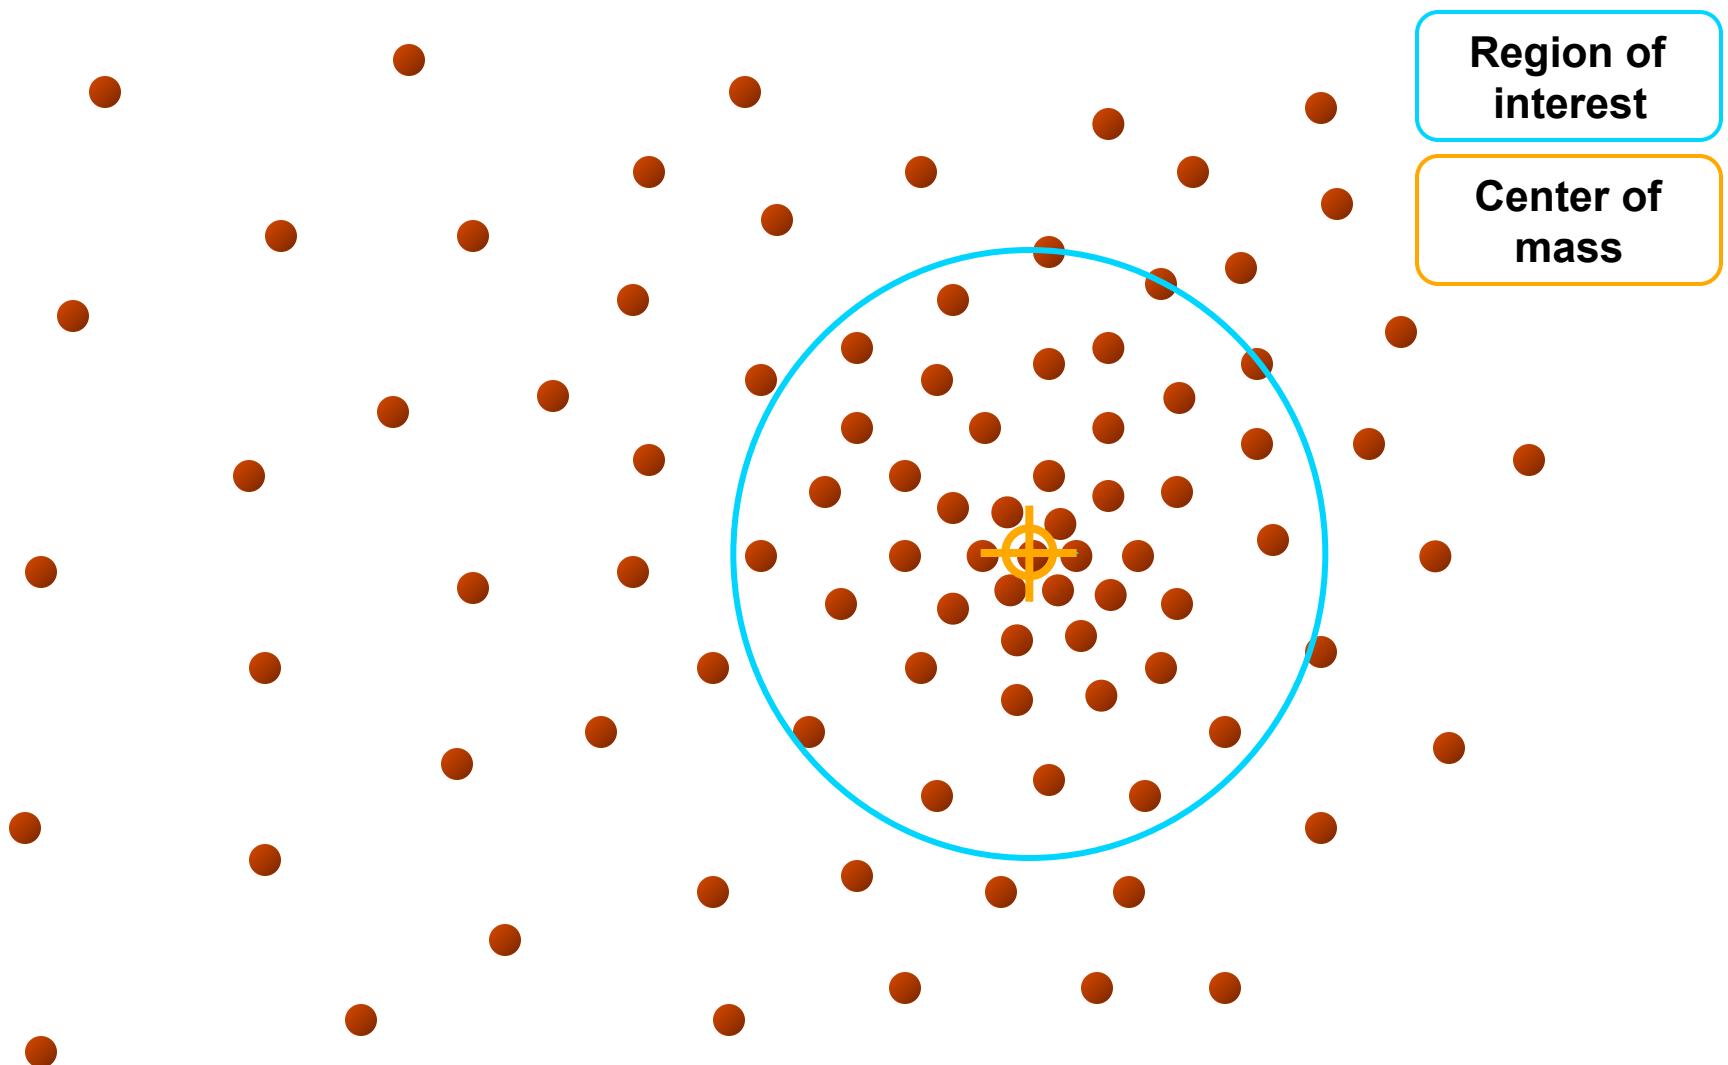
Mean shift



Mean shift



Mean shift



Region of
interest

Center of
mass

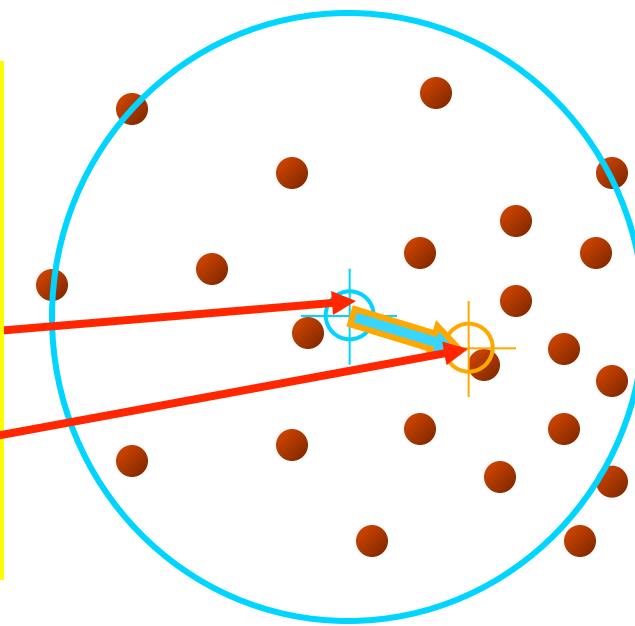
Computing the Mean Shift

Simple Mean Shift procedure:

- Compute mean shift vector
- Translate the Kernel window by

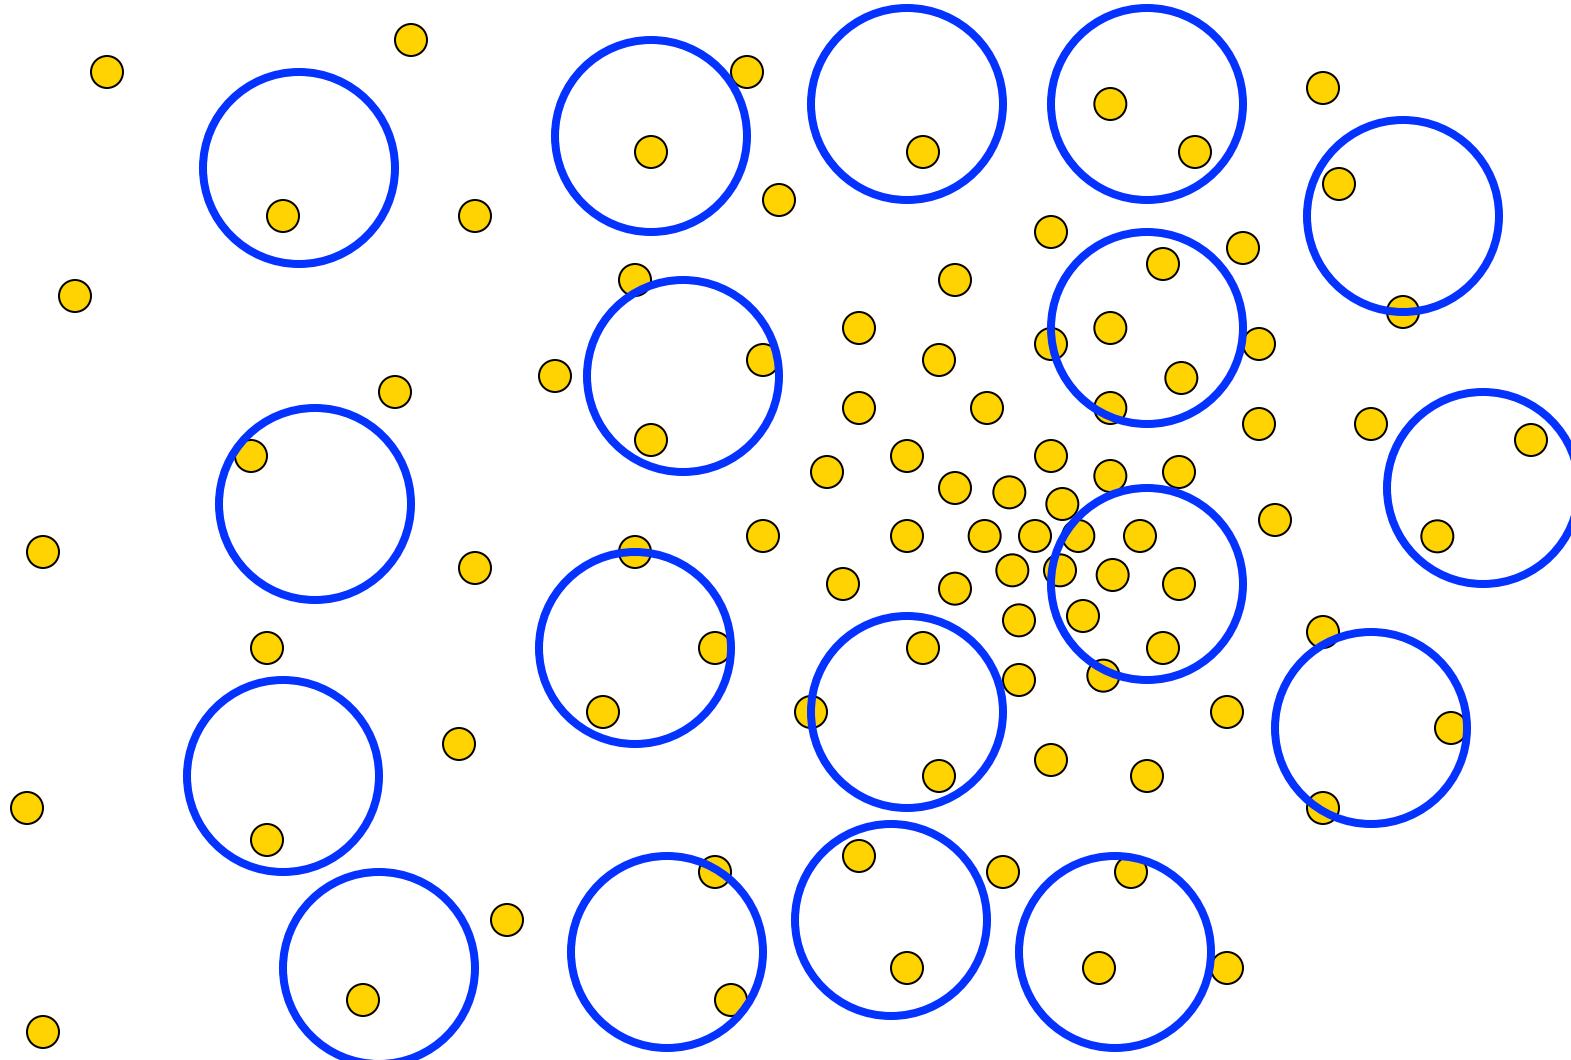
$m(x)$

$$m(x) = \left[\begin{array}{c} \sum_{i=1}^n \mathbf{x}_i g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right) \\ \sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right) \end{array} \right] - \boxed{\mathbf{x}}$$



$$g(x) = -k'(x)$$

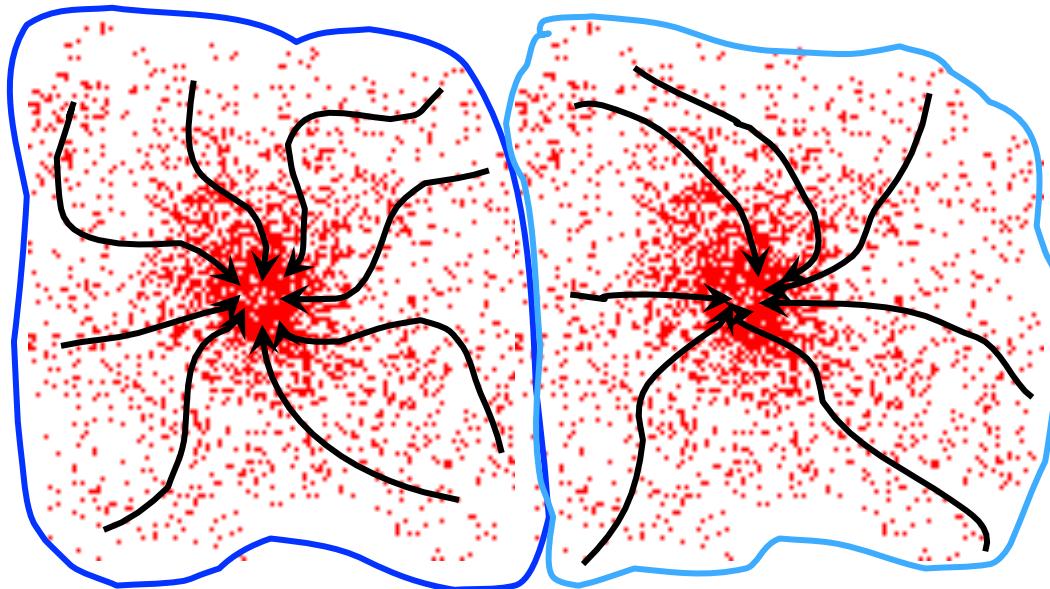
Real Modality Analysis



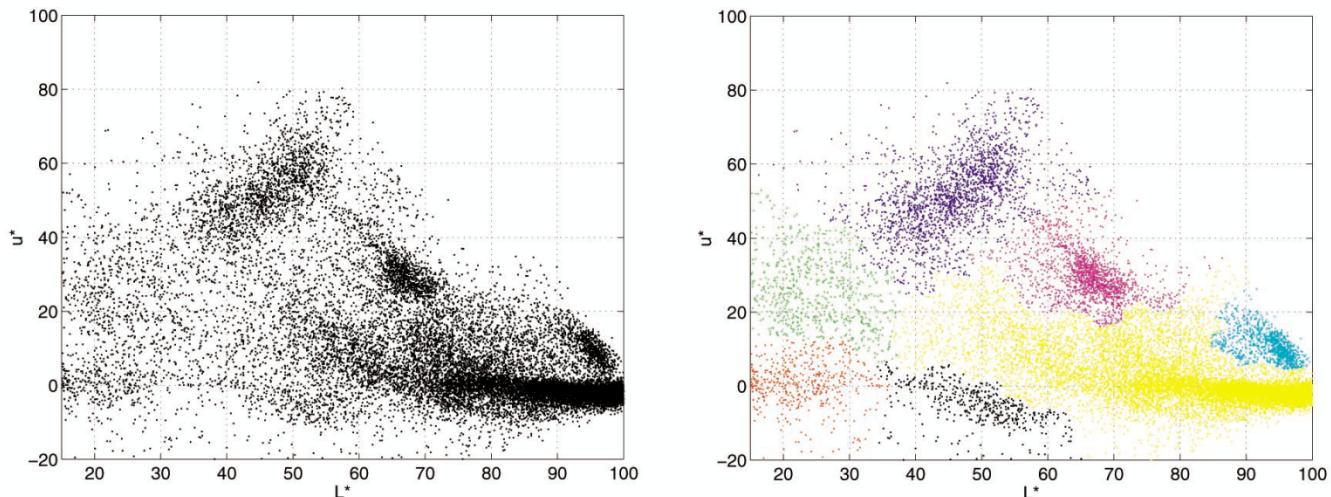
Attraction basin

Attraction basin: the region for which all trajectories lead to the same mode

Cluster: all data points in the attraction basin of a mode

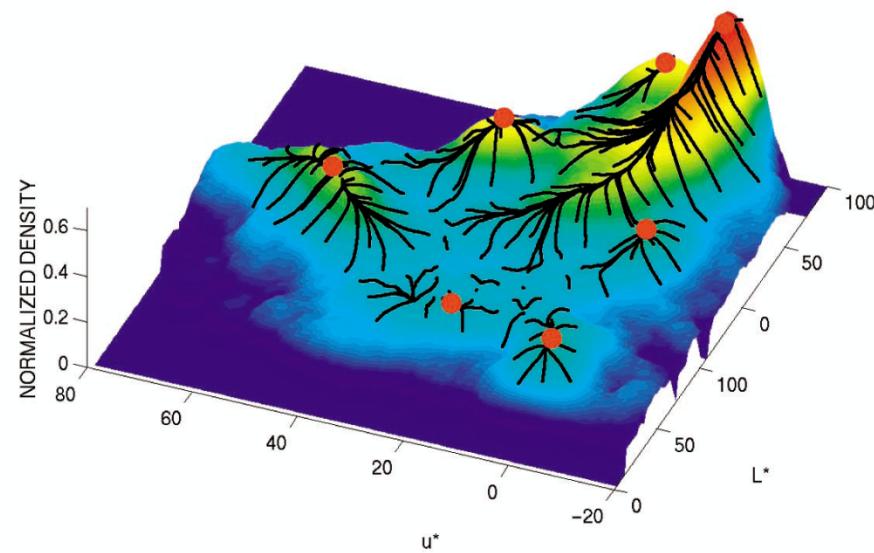


Attraction basin



(a)

(b)



Mean shift clustering

The mean shift algorithm seeks *modes* of the given set of points

1. Choose kernel and bandwidth
2. For each point:
 - a) Center a window on that point
 - b) Compute the mean of the data in the search window
 - c) Center the search window at the new mean location
 - d) Repeat (b,c) until convergence
3. Assign points that lead to nearby modes to the same cluster

Segmentation by Mean Shift

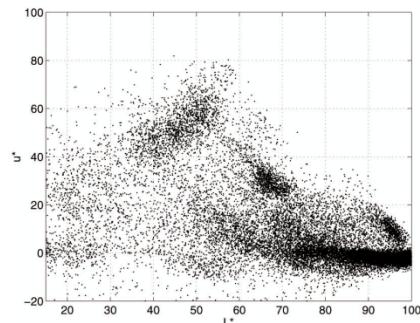
Compute features for each pixel (color, gradients, texture, etc); also store each pixel's position

Set kernel size for features K_f and position K_s

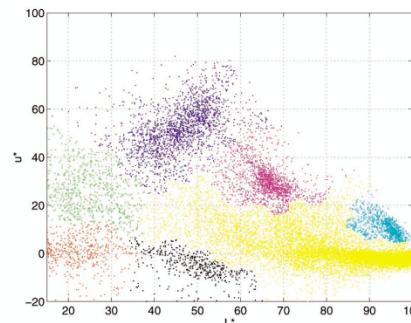
Initialize windows at individual pixel locations

Perform mean shift for each window until convergence

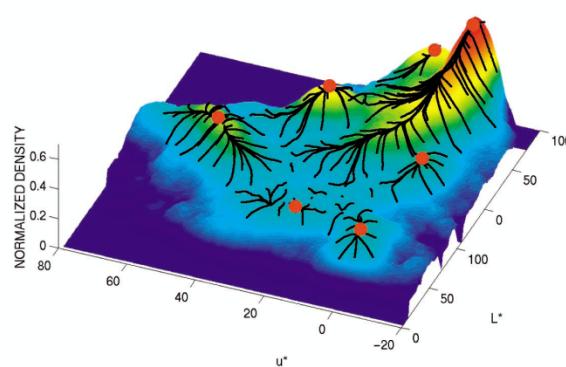
Merge modes that are within width of K_f and K_s



(a)

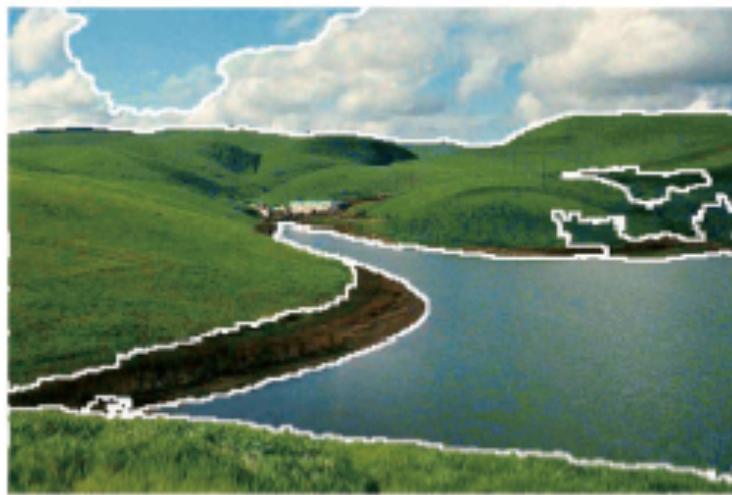
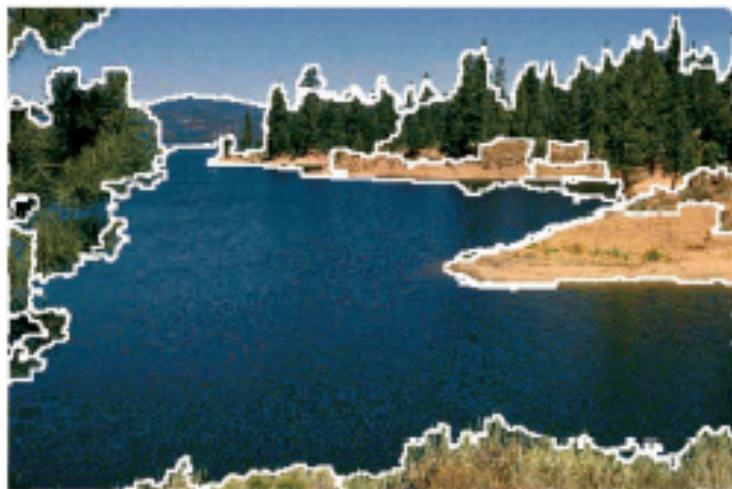


(b)



Mean shift segmentation results





<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>

Mean-shift: other issues

- Speedups
 - Binned estimation – replace points within some “bin” by point at center with mass
 - Fast search of neighbors – e.g., k-d tree or approximate NN
 - Update all windows in each iteration (faster convergence)
- Other tricks
 - Use kNN to determine window sizes adaptively
- Lots of theoretical support

D. Comaniciu and P. Meer, Mean Shift: A Robust Approach toward Feature Space Analysis, PAMI 2002.

Mean shift pros and cons

Pros

- Good general-purpose segmentation
- Flexible in number and shape of regions
- Robust to outliers

Cons

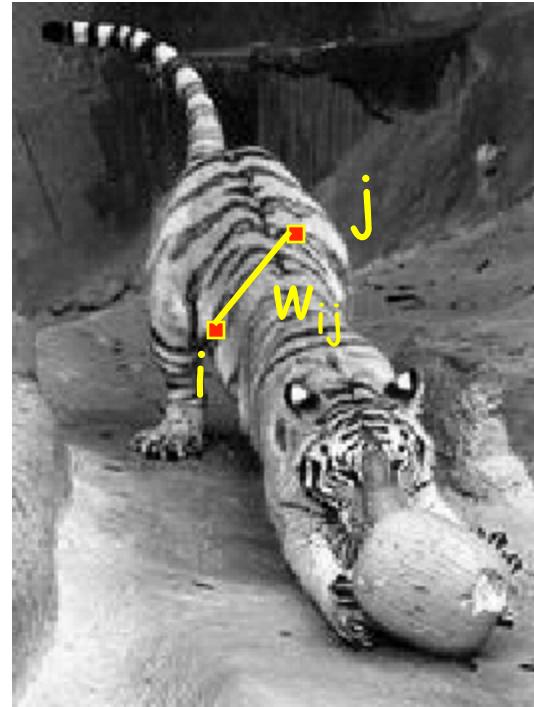
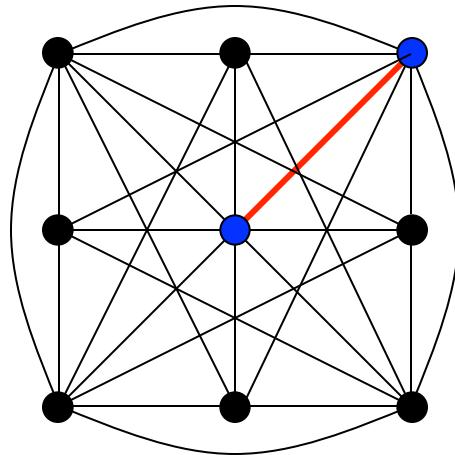
- Have to choose kernel size in advance
- Not suitable for high-dimensional features

When to use it

- Oversegmentation
- Multiple segmentations
- Tracking, clustering, filtering applications
 - D. Comaniciu, V. Ramesh, P. Meer:

[Real-Time Tracking of Non-Rigid Objects using Mean Shift](#),
Best Paper Award, IEEE Conf. Computer Vision and Pattern
Recognition (CVPR'00), Hilton Head Island, South Carolina,
Vol. 2, 142-149, 2000

Segmentation as graph partitioning

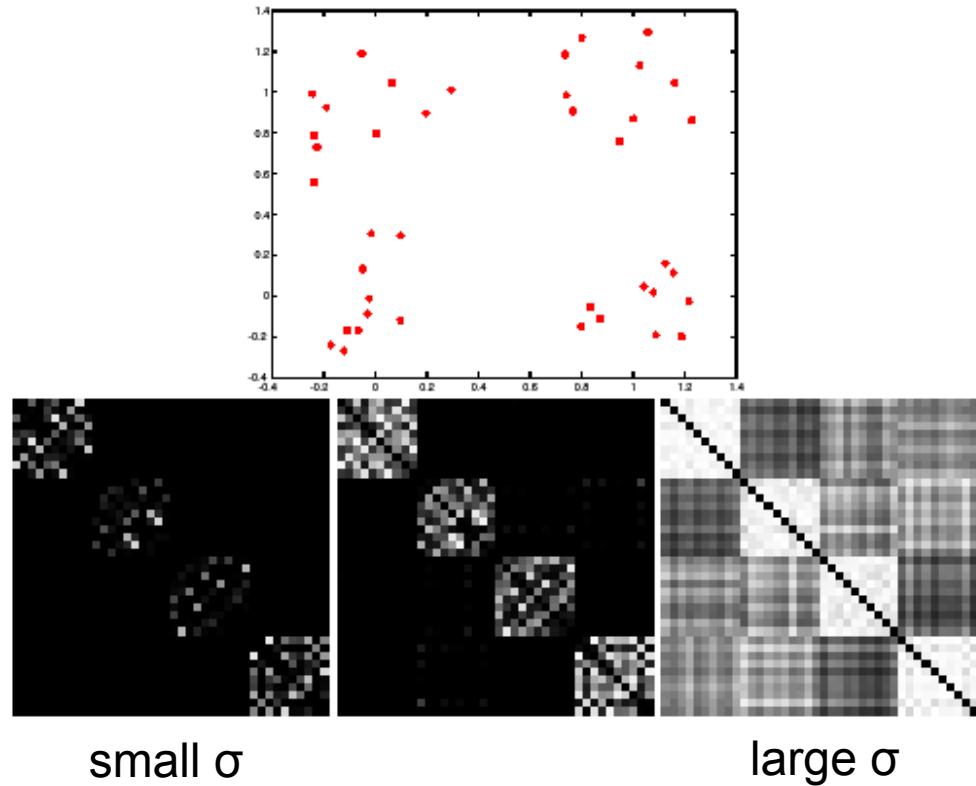
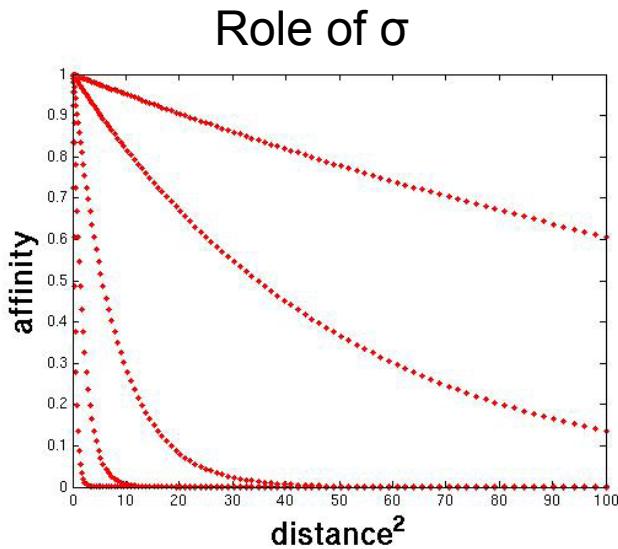


- Node for every pixel
- Edge between every pair of pixels (or every pair of “sufficiently close” pixels)
- Each edge is weighted by the *affinity* or similarity of the two nodes

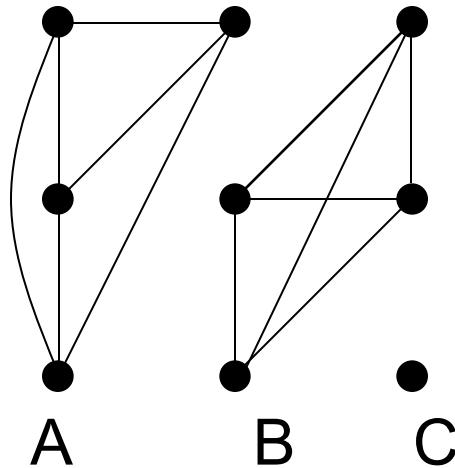
Measuring affinity

- Represent each pixel by a feature vector \mathbf{x} and define an appropriate distance function

$$\text{affinity}(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{1}{2\sigma^2} \text{dist}(\mathbf{x}_i, \mathbf{x}_j)^2\right)$$

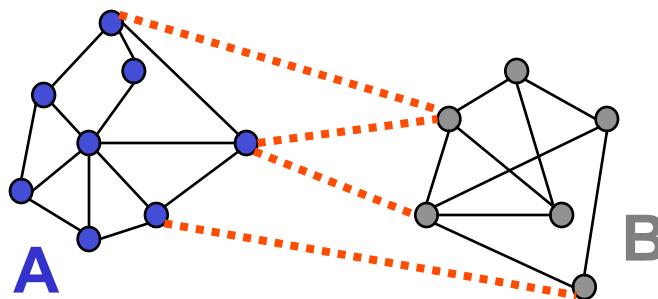


Segmentation as graph partitioning



- Break Graph into Segments
 - Delete links that cross between segments
 - Easiest to break links that have low affinity
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

Graph cut

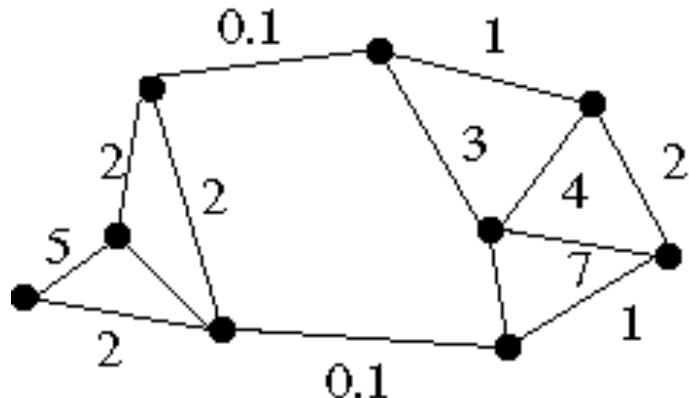


- Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges
- A graph cut gives us a segmentation
 - What is a “good” graph cut and how do we find one?

Minimum cut

- We can do segmentation by finding the *minimum cut* in a graph
 - Efficient algorithms exist for doing this

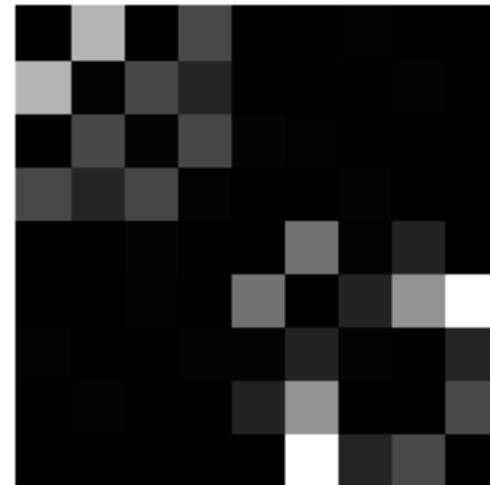
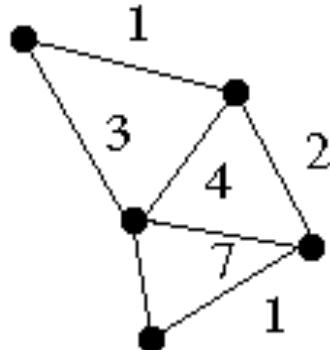
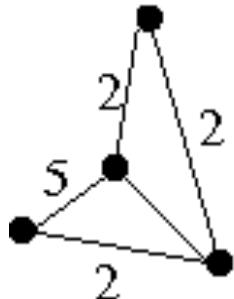
Minimum cut example



Minimum cut

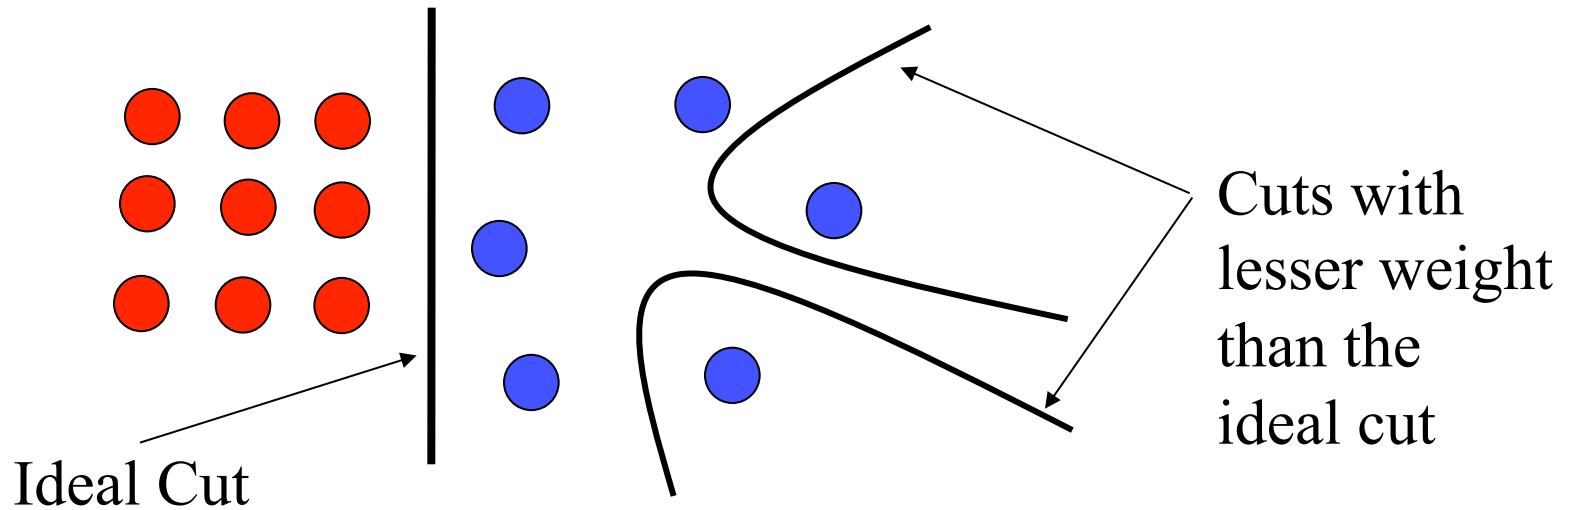
- We can do segmentation by finding the *minimum cut* in a graph
 - Efficient algorithms exist for doing this

Minimum cut example



Normalized cut

- Drawback: minimum cut tends to cut off very small, isolated components



Normalized cut

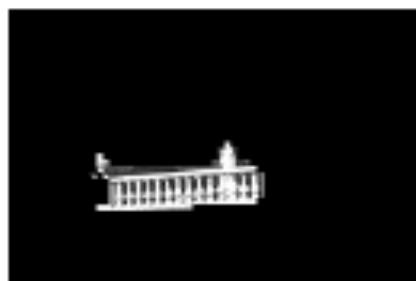
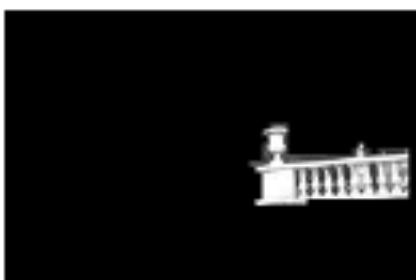
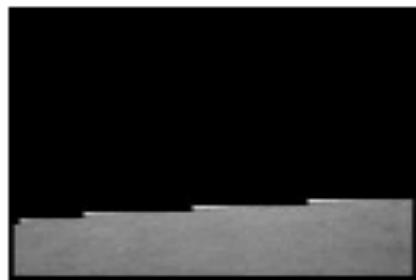
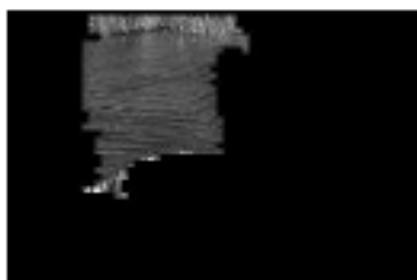
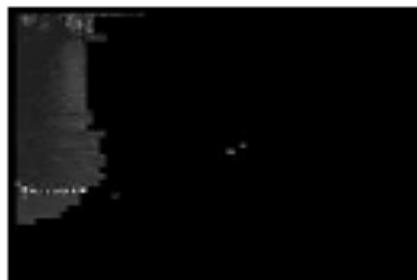
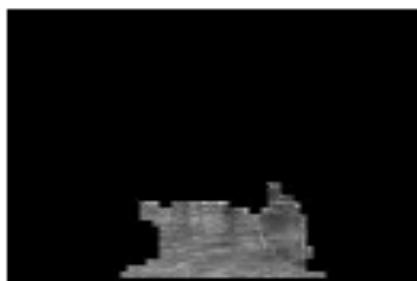
- Drawback: minimum cut tends to cut off very small, isolated components
- This can be fixed by normalizing the cut by the weight of all the edges incident to the segment
- The *normalized cut* cost is:

$$\frac{w(A, B)}{w(A, V)} + \frac{w(A, B)}{w(B, V)}$$

$w(A, B)$ = sum of weights of all edges between A and B

- Finding the globally optimal cut is NP-complete, but a relaxed version can be solved using a generalized eigenvalue problem

Example result

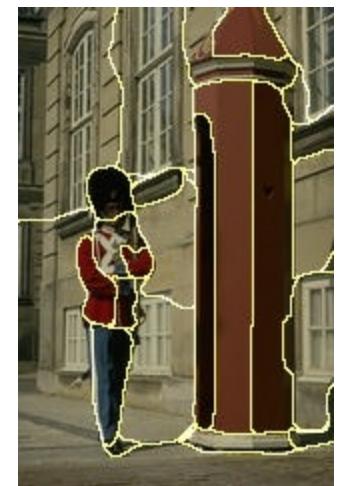
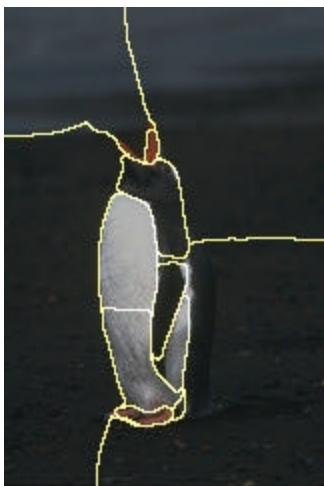
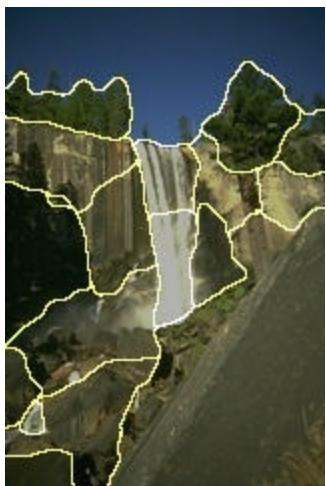


Results: Berkeley Segmentation Engine



<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Results: Berkeley Segmentation Engine



<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Normalized cuts: Pro and con

- Pro
 - Generic framework, can be used with many different features and affinity formulations
- Con
 - High storage requirement and time complexity: involves solving a generalized eigenvalue problem of size $n \times n$, where n is the number of pixels

Efficient graph-based segmentation



- Runs in time nearly linear in the number of edges
- Easy to control coarseness of segmentations
- Results can be unstable

P. Felzenszwalb and D. Huttenlocher,
[Efficient Graph-Based Image Segmentation](#), IJCV 2004

Segmentation as labeling

- Suppose we want to segment an image into foreground and background
 - Binary labeling problem



Segmentation as labeling

- Suppose we want to segment an image into foreground and background
 - Binary labeling problem



User sketches out a few strokes on foreground and background...

How do we label the rest of the pixels?

Binary segmentation as energy minimization

- Define a labeling L as an assignment of each pixel with a 0-1 label (background or foreground)
- Find the labeling L that minimizes

$$E(L) = E_d(L) + \lambda E_s(L)$$

data term

smoothness term



How similar is
each labeled pixel
to the foreground
or background?

Encourage spatially
coherent segments

$$E(L) = E_d(L) + \lambda E_s(L)$$



$$E_d(L) = \sum_{(x,y)} C(x,y, L(x,y))$$

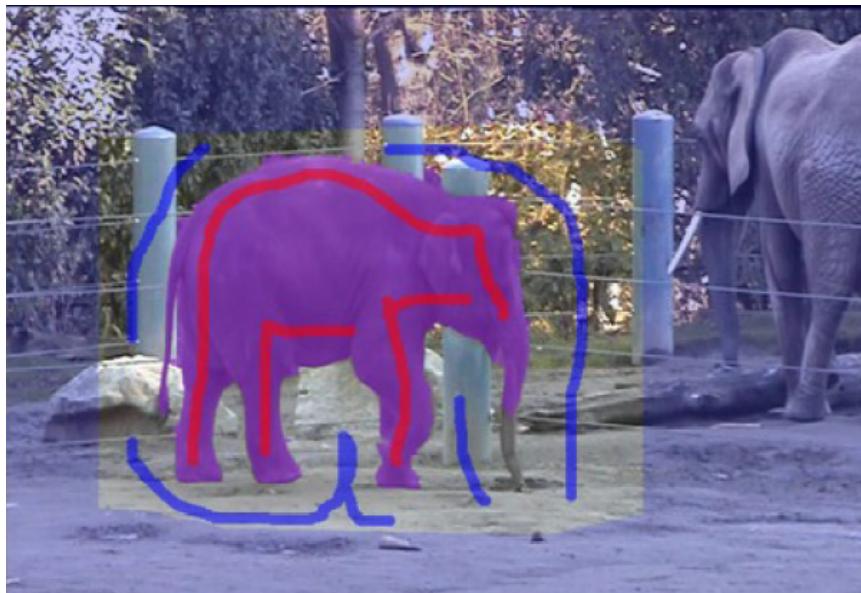
$$\tilde{L}(x, y)$$

$$C(x,y, L(x,y)) = \begin{cases} \infty & \text{if } L(x,y) \neq \tilde{L}(x,y) \\ C'(x,y, L(x,y)) & \text{otherwise} \end{cases}$$

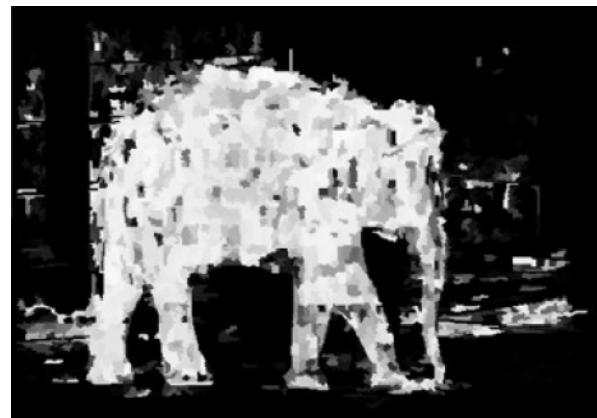
$C'(x,y,0)$: “distance” from pixel to background
 $C'(x,y,1)$: “distance” from pixel to foreground

} computed by creating a color model from user-labeled pixels

$$E(L) = E_d(L) + \lambda E_s(L)$$



$$C'(x, y, 0)$$



$$C'(x, y, 1)$$

$$E(L) = E_d(L) + \lambda E_s(L)$$

- Neighboring pixels should generally have the same labels
 - Unless the pixels have very different intensities



$$E_s(L) = \sum_{\text{neighbors } (p,q)} w_{pq} |L(p) - L(q)|$$

w_{pq} : similarity in intensity of p and q

$$w_{pq} = 0.1$$
$$w_{pq} = 10.0$$

Binary segmentation as energy minimization

$$E(L) = E_d(L) + \lambda E_s(L)$$

- For this problem, we can efficiently find the global minimum using the max flow / min cut algorithm

Y. Boykov and M.-P. Jolly,

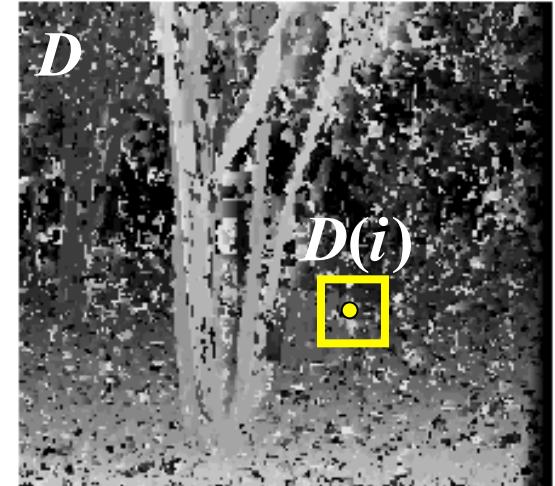
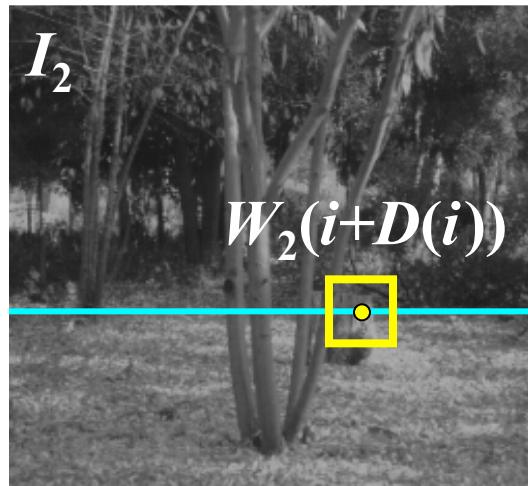
[Interactive Graph Cuts for Optimal Boundary and Region Segmentation of Objects in N-D Images](#), ICCV 2001

GrabCut



C. Rother, V. Kolmogorov, and A. Blake,
[“GrabCut” — Interactive Foreground Extraction using Iterated Graph Cuts,](#)
SIGGRAPH 2004

Recall: Stereo as energy minimization



$$E(D) = \underbrace{\sum_i (W_1(i) - W_2(i + D(i)))^2}_{\text{data term}} + \lambda \underbrace{\sum_{\text{neighbors } i,j} \rho(D(i) - D(j))}_{\text{smoothness term}}$$

- Energy functions of this form can be minimized using *graph cuts*

Y. Boykov, O. Veksler, and R. Zabih,
[Fast Approximate Energy Minimization via Graph Cuts](#), PAMI 2001

Further reading

- Nicely written mean-shift explanation (with math)

<http://saravananthirumuruganathan.wordpress.com/2010/04/01/introduction-to-mean-shift-algorithm/>

- Includes .m code for mean-shift clustering

- Mean-shift paper by Comaniciu and Meer

<http://www.caip.rutgers.edu/~comanici/Papers/MsRobustApproach.pdf>

- Adaptive mean shift in higher dimensions

<http://mis.hevra.haifa.ac.il/~ishimshoni/papers/chap9.pdf>

- Contours to regions (watershed): Arbelaez et al. 2009

http://www.eecs.berkeley.edu/~arbelaez/publications/Arbelaez_Maire_Fowlkes_Malik_CVPR2009.pdf